

Numerical Weather Prediction

Facing the Future with C++

T. Quintino, F. Rathgeber, W. Deconinck,

B. Raoult, M. Fuentes

ECMWF

C++Now 2015, May 13 2015



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

1

Overview

- ECMWF
- Numerical Weather Prediction
- Challenges Ahead
- Learning from the Past
- Preparing the Future



ECMWF



ECMWF EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

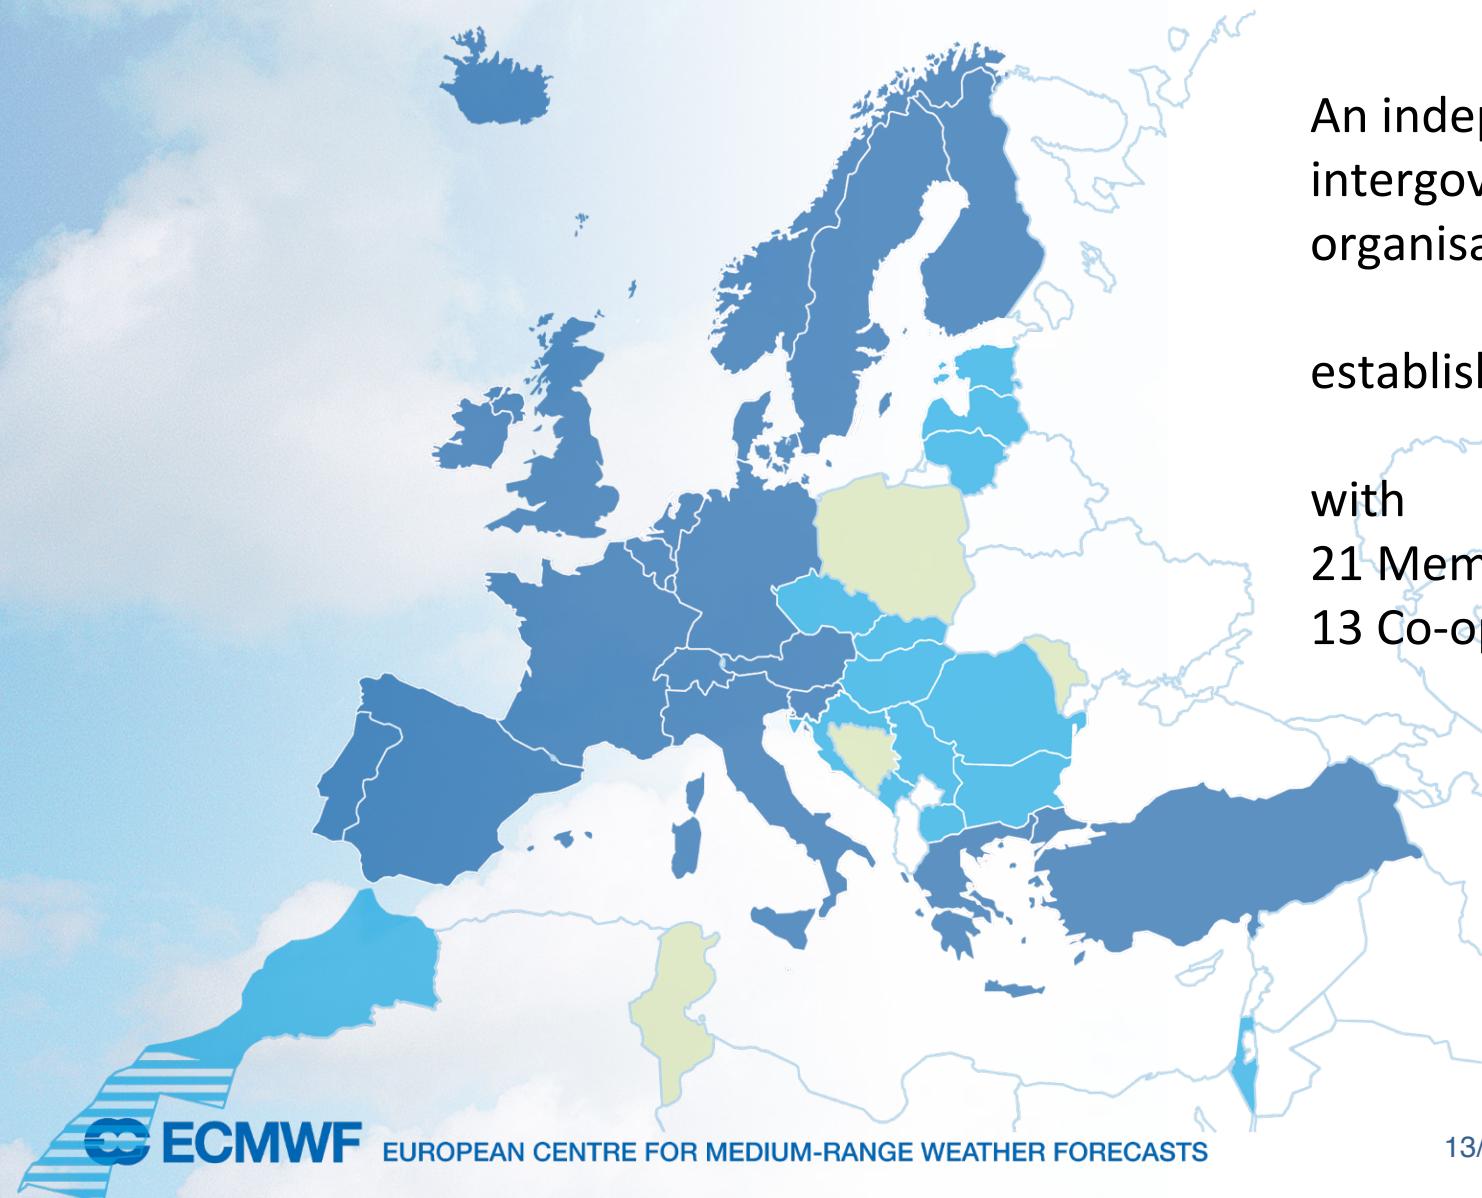
3

ECMWF

■ Member States

■ Co-operating States

■ Under negotiation



An independent
intergovernmental
organisation

established in 1975

with
21 Member States
13 Co-operating States

Who are we and what do we do?

Weather Forecasts

We produce
global weather forecasts

Medium-Range

Up to 15 days ahead.
Our products also include
monthly and **seasonal** forecasts
and we collect and store
meteorological data.

What do we have to achieve this?

People

About 260 staff,
specialists and contractors

Equipment

State-of-the-art supercomputers
and data handling systems

Budget

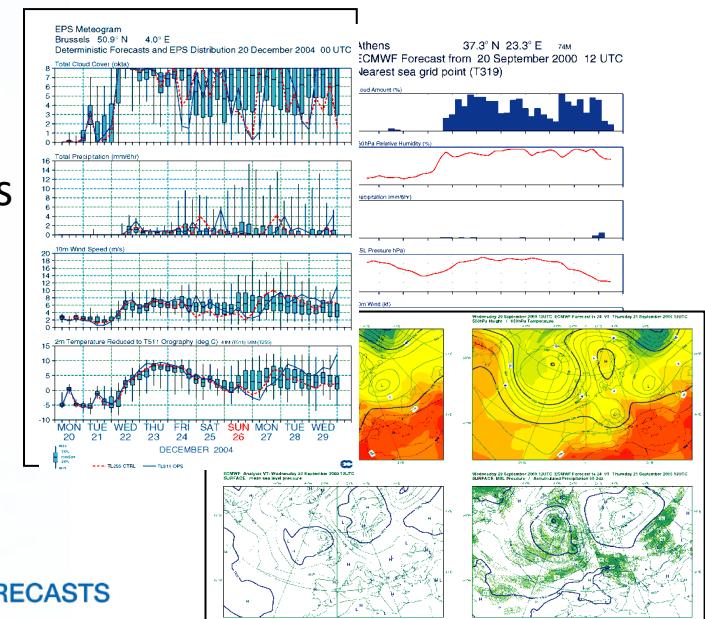
£50 million per year



Reading, United Kingdom



ECMWF EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS



Global observation system



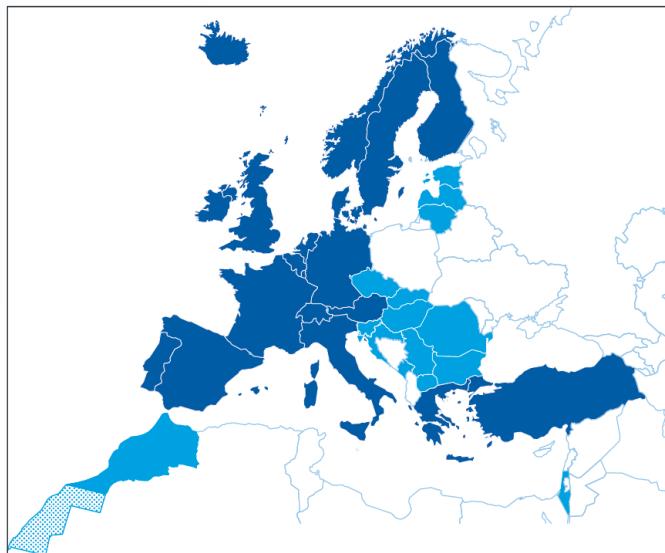
Global numerical weather forecasts



Users



National weather services



Numerical Weather Prediction

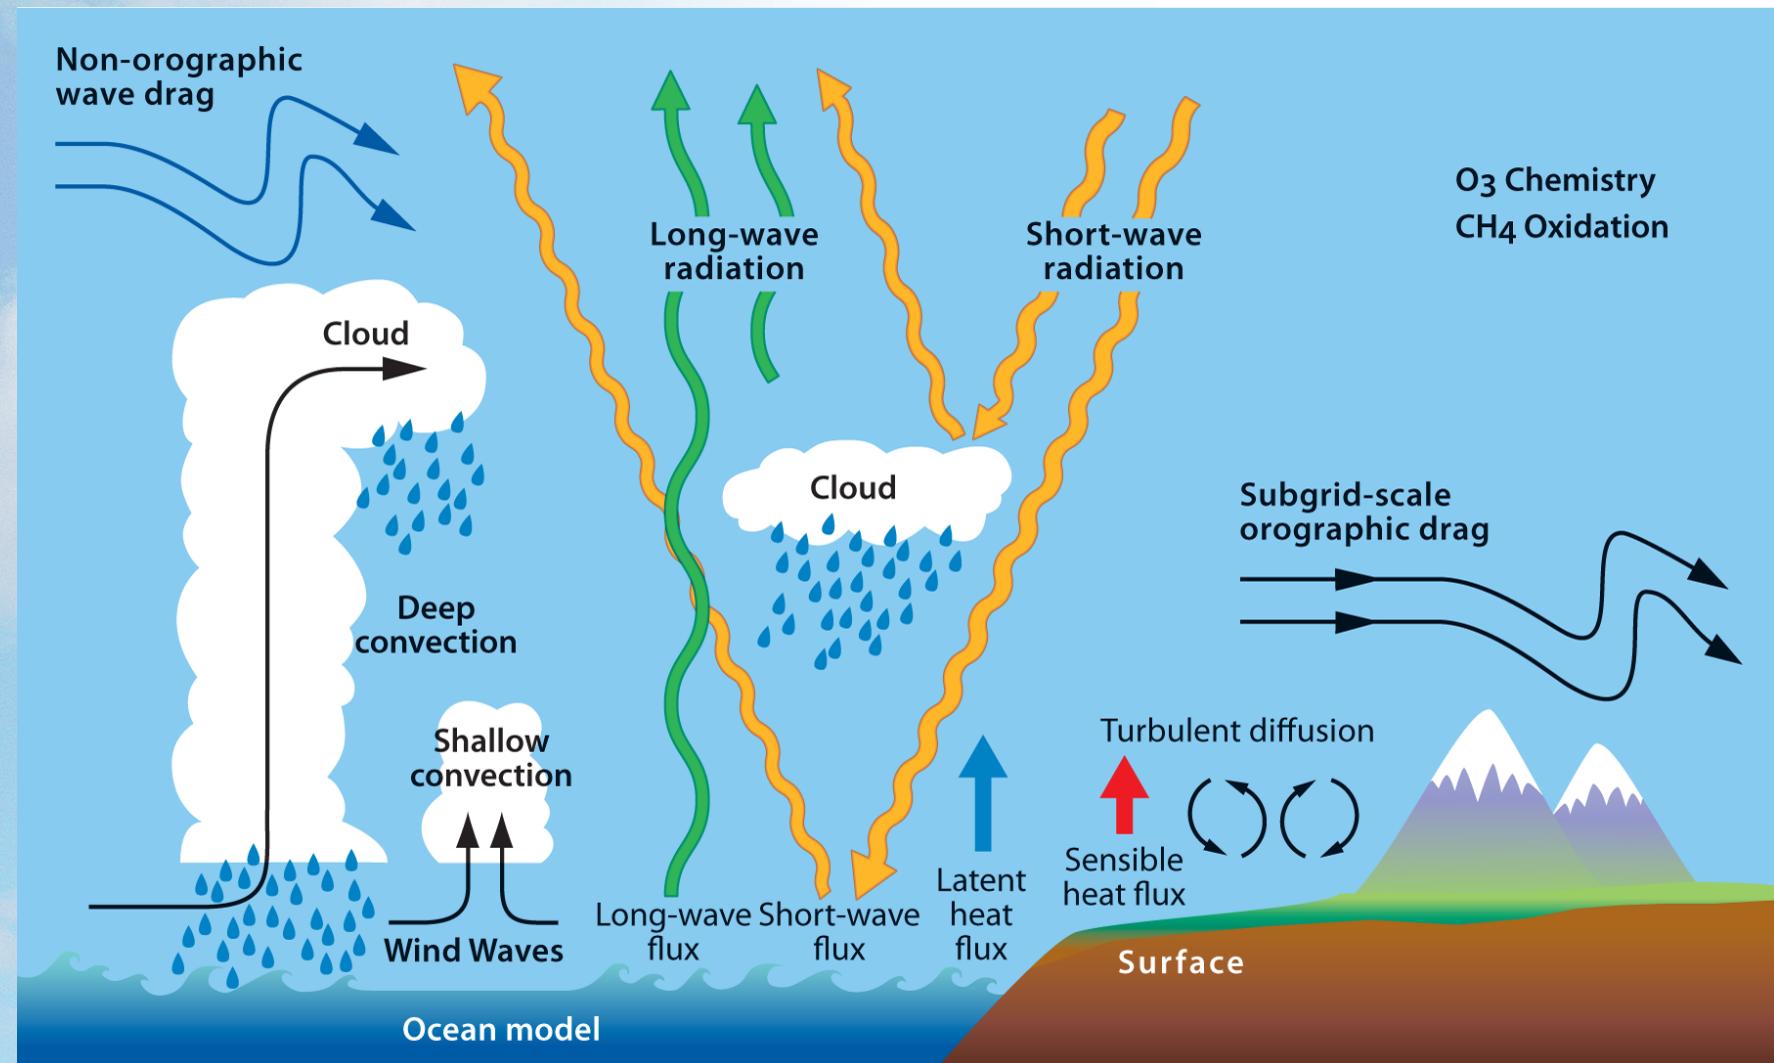


EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

7

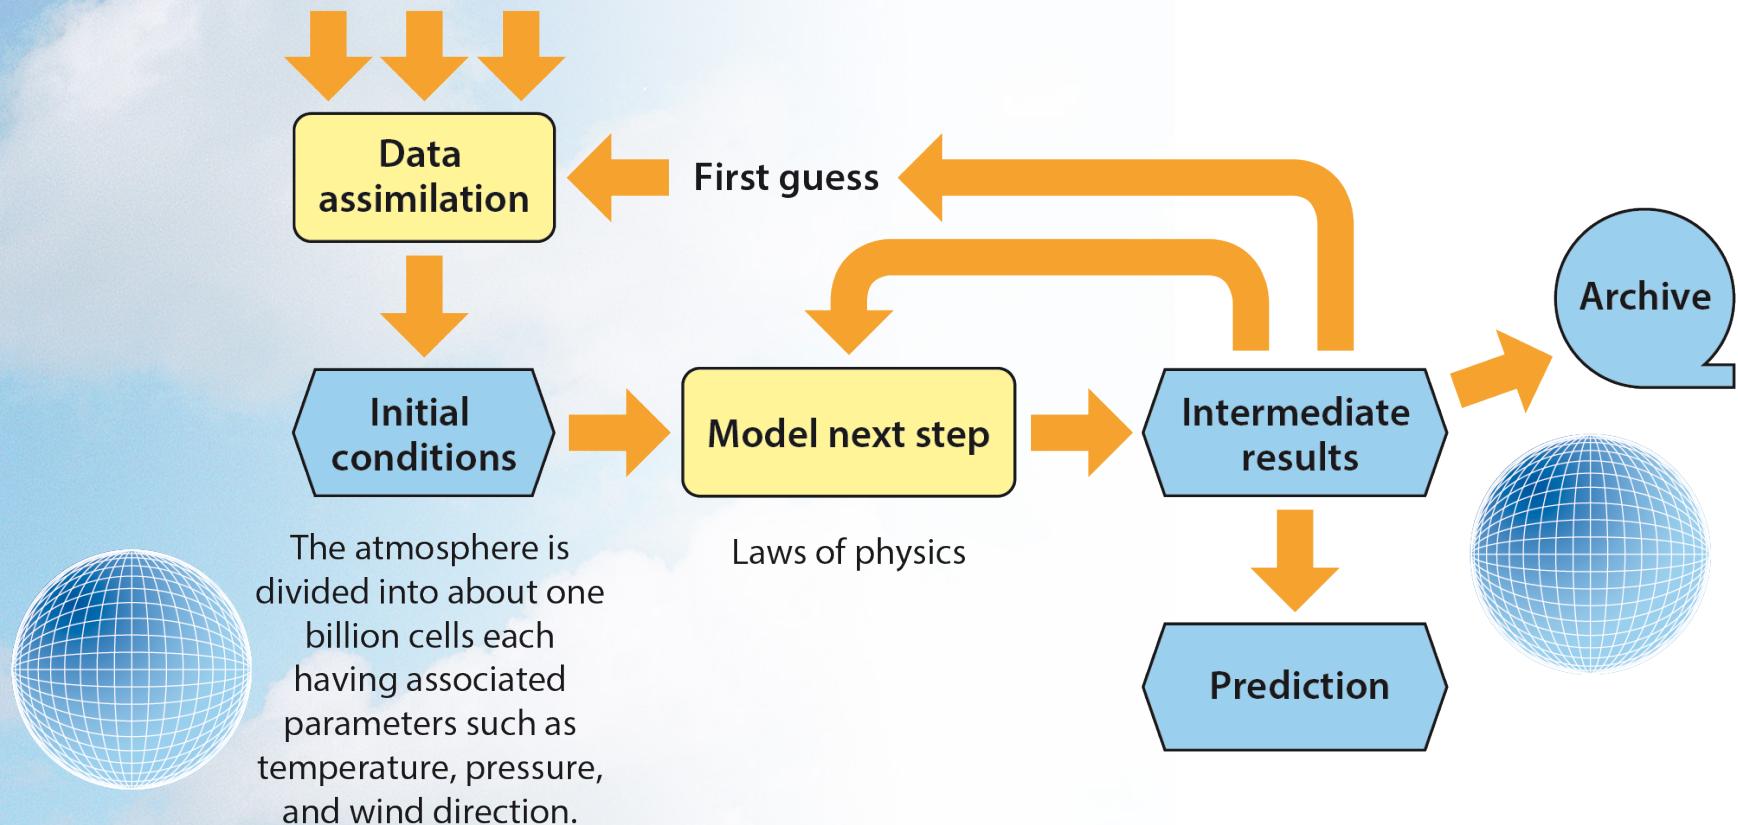
Model

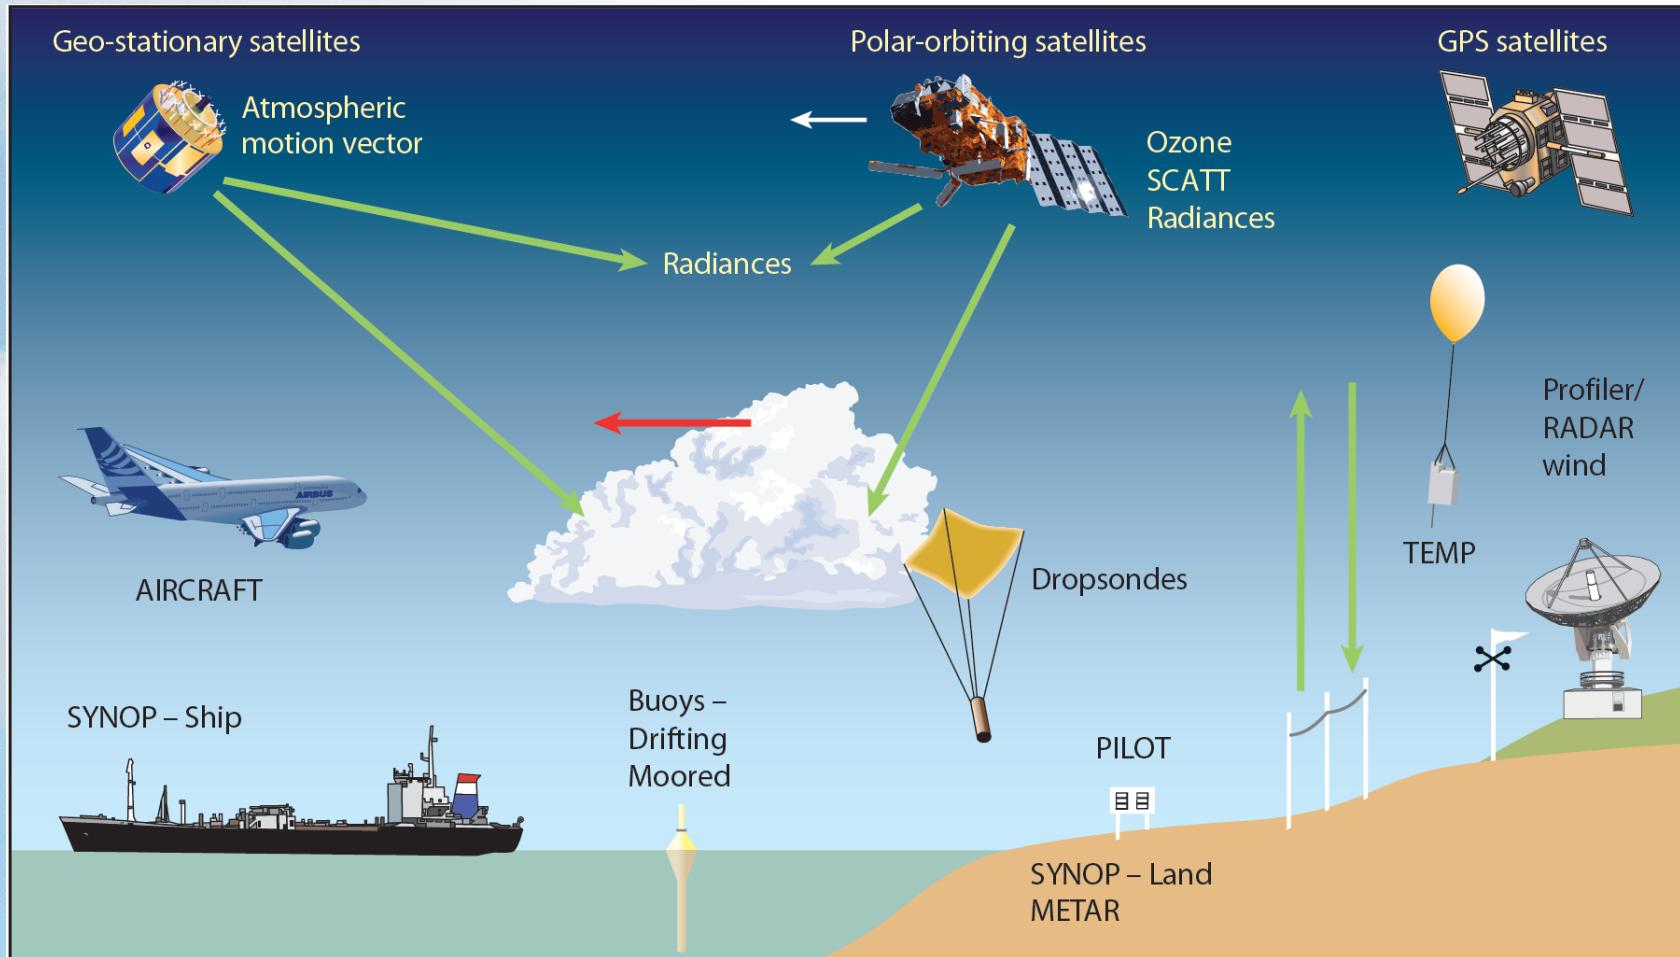


A basic description of our models



Approximately 20 million observations



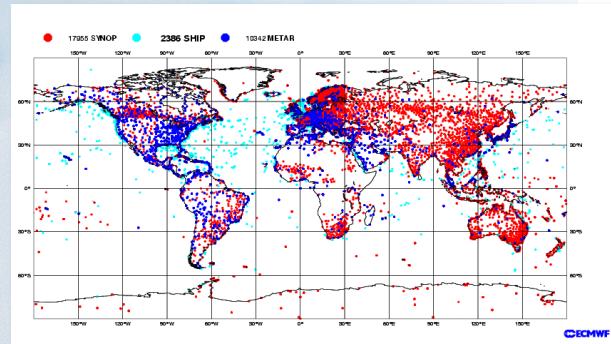


Data sources for the ECMWF Meteorological Operational System

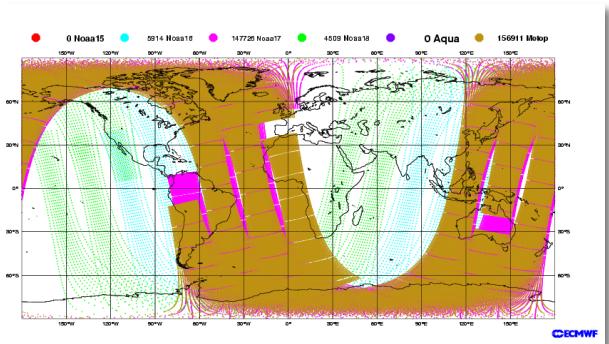
Major assimilated datasets

Receive **300 million** observations
from **130 sources** daily.

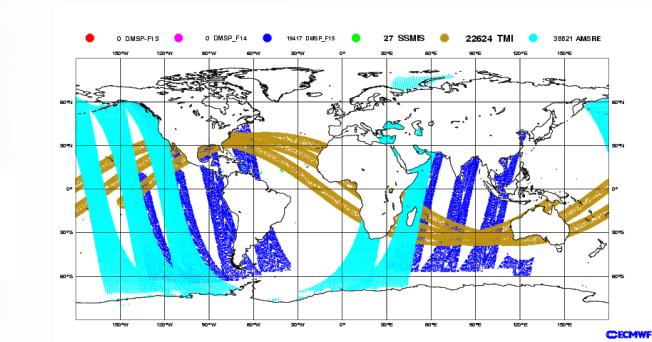
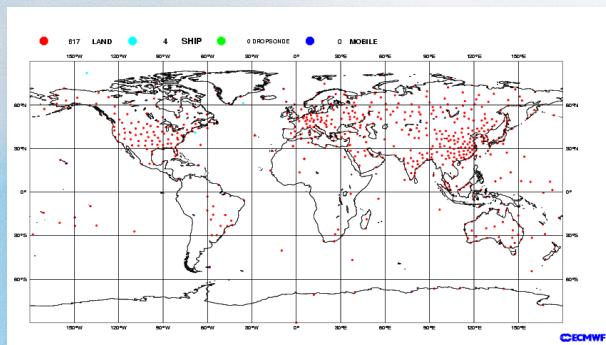
Surface stations



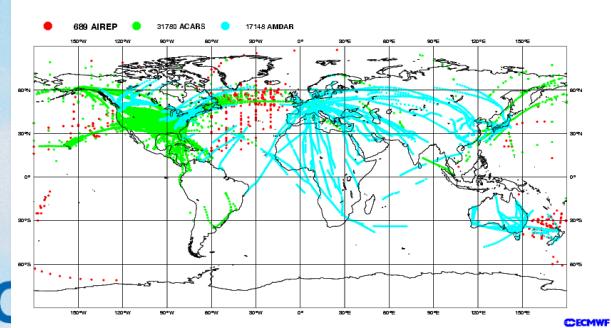
Polar,
infrared



Radiosonde balloons

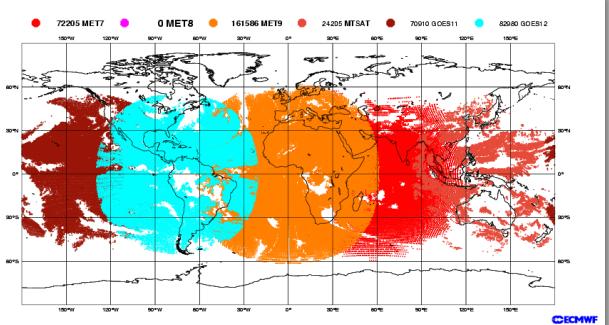


Aircraft

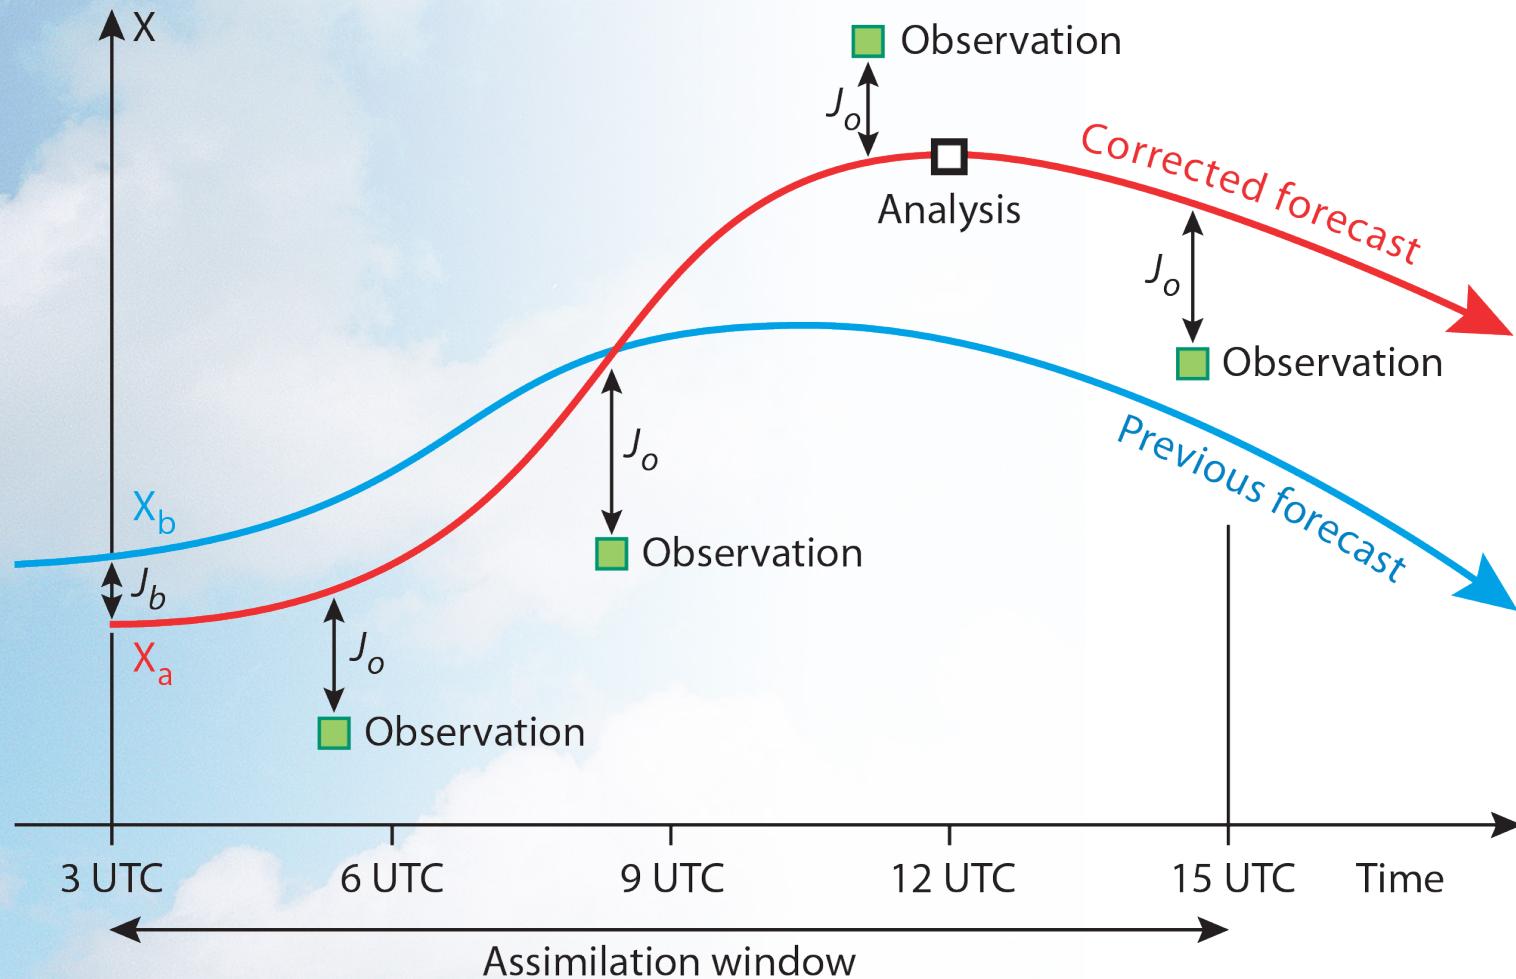


Geostationary, IR

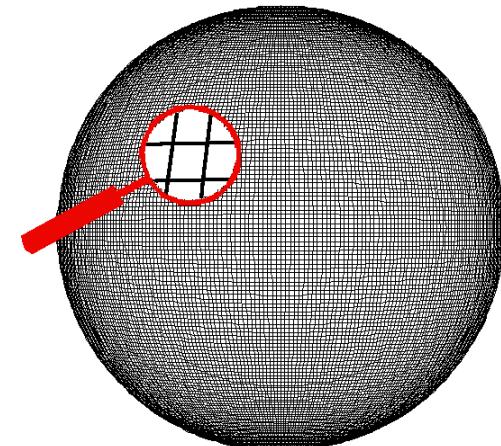
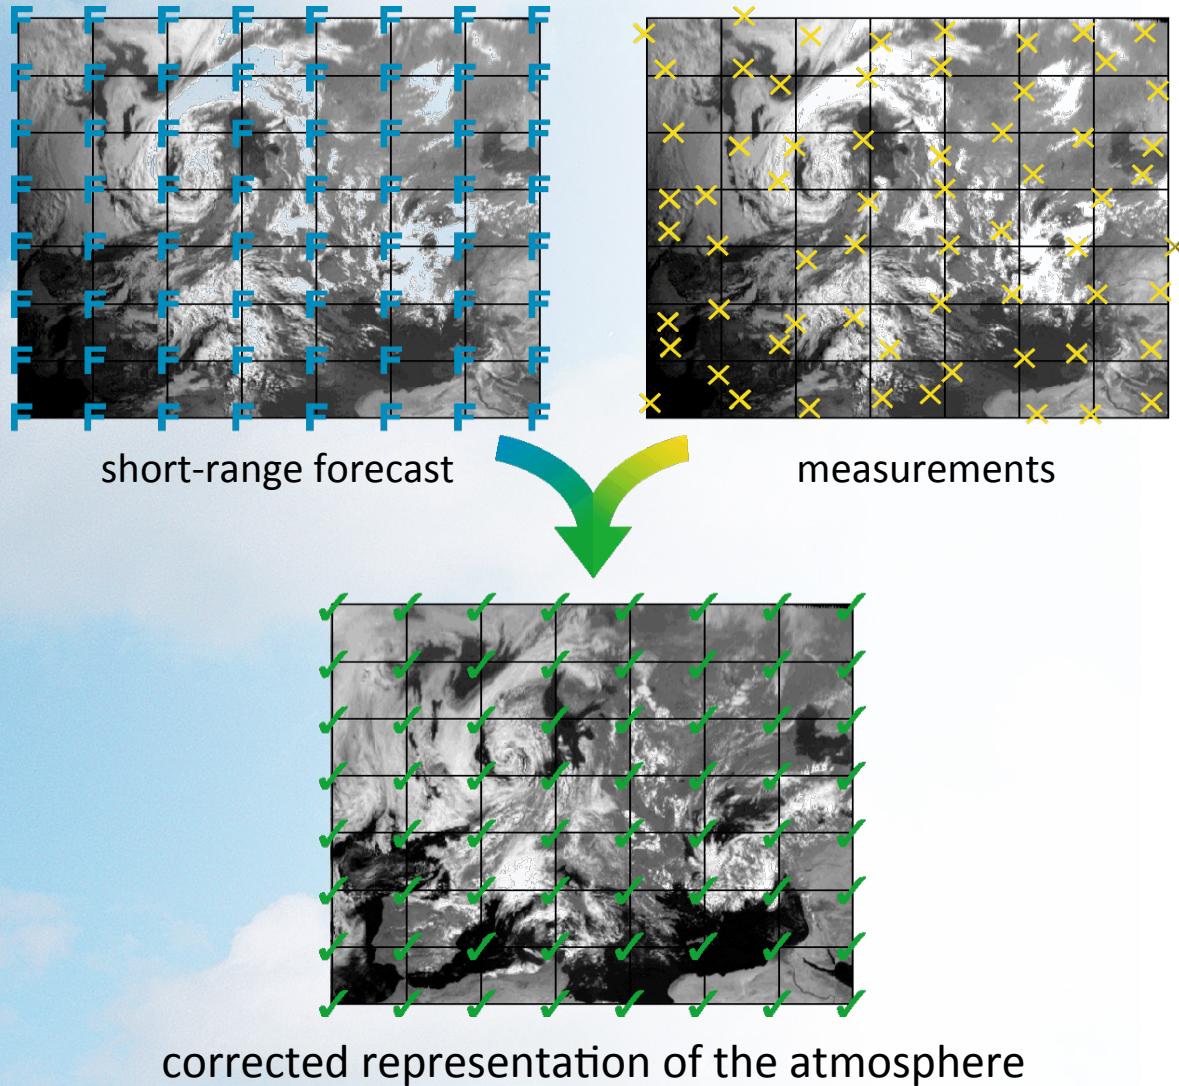
RANGE WEATHER FORECAST



4D-Var Assimilation



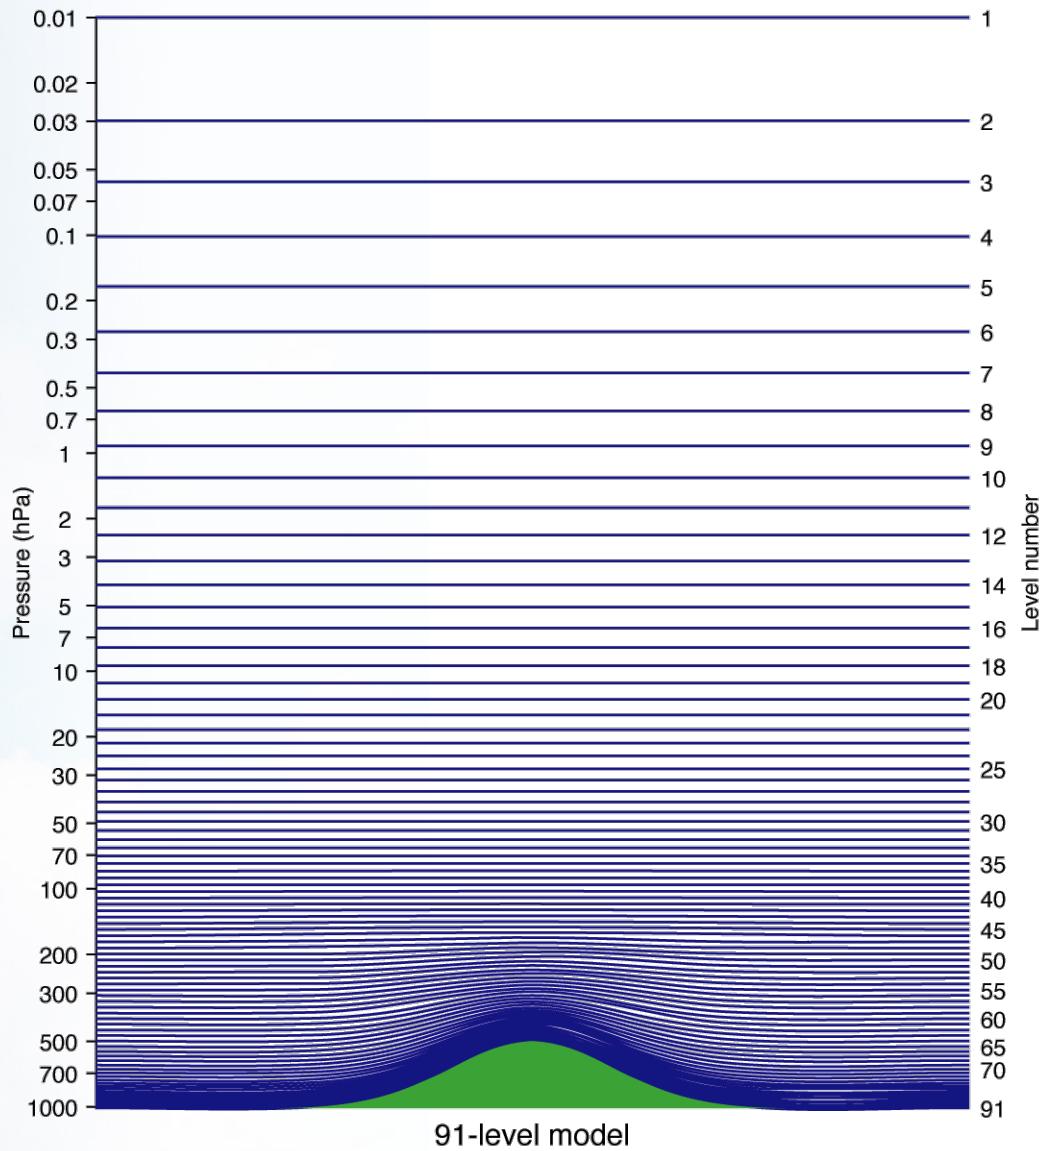
Data Assimilation



Model Levels

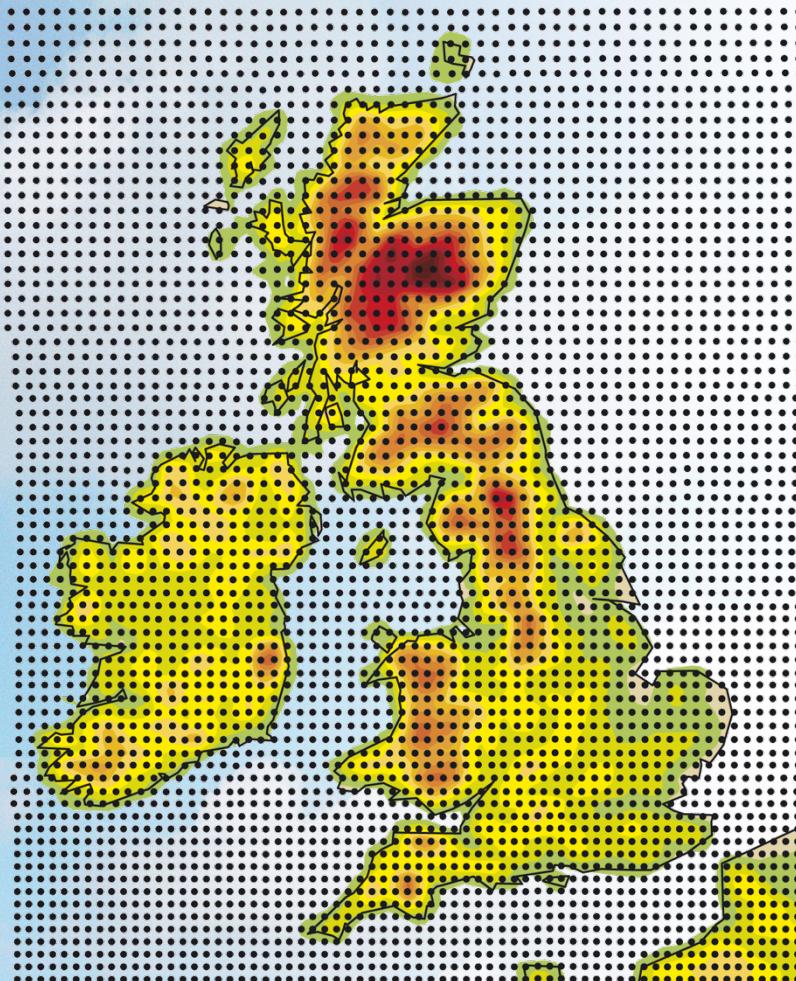
Current operational:
137 levels

Alt. $\approx 40.000\text{m}$

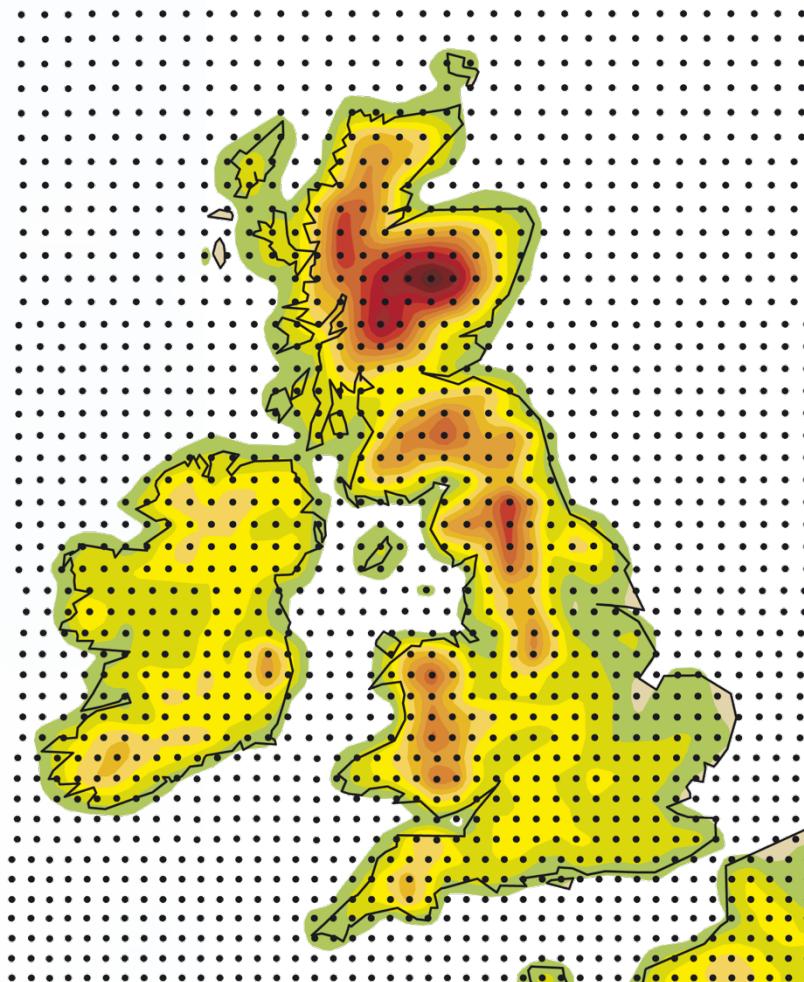


Operational model grid

High Resolution (16km)

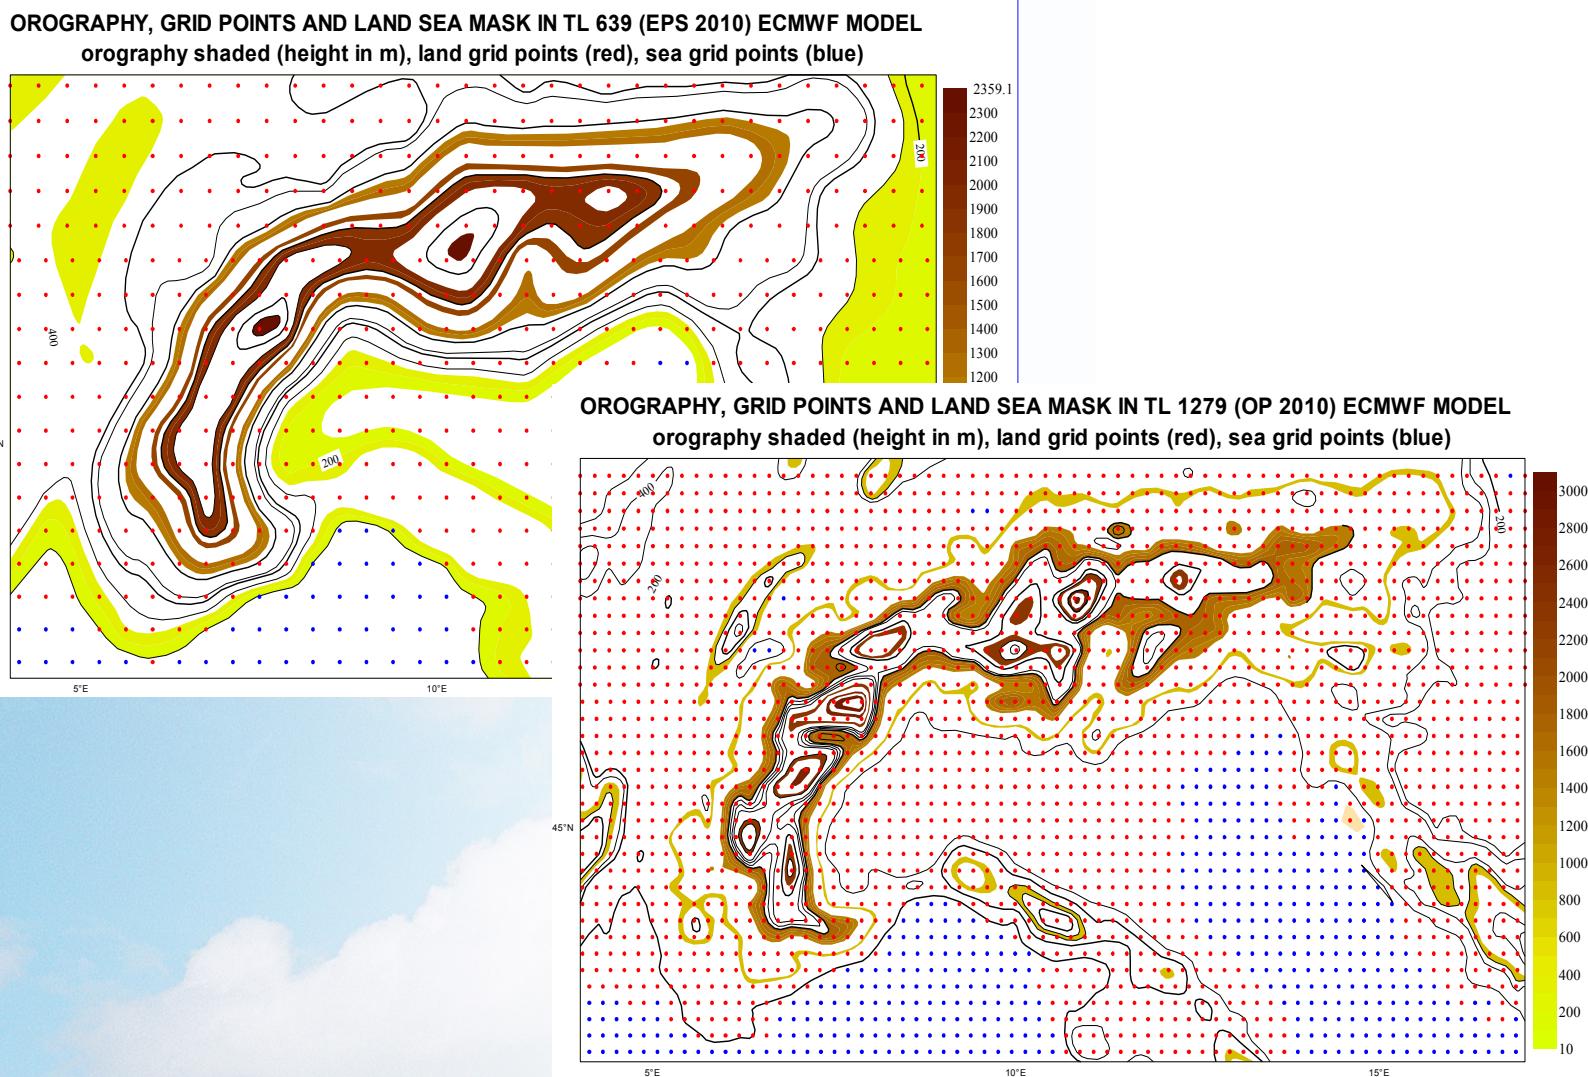


ENS (32Km)



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

Grid Effect



Meteorological Fields

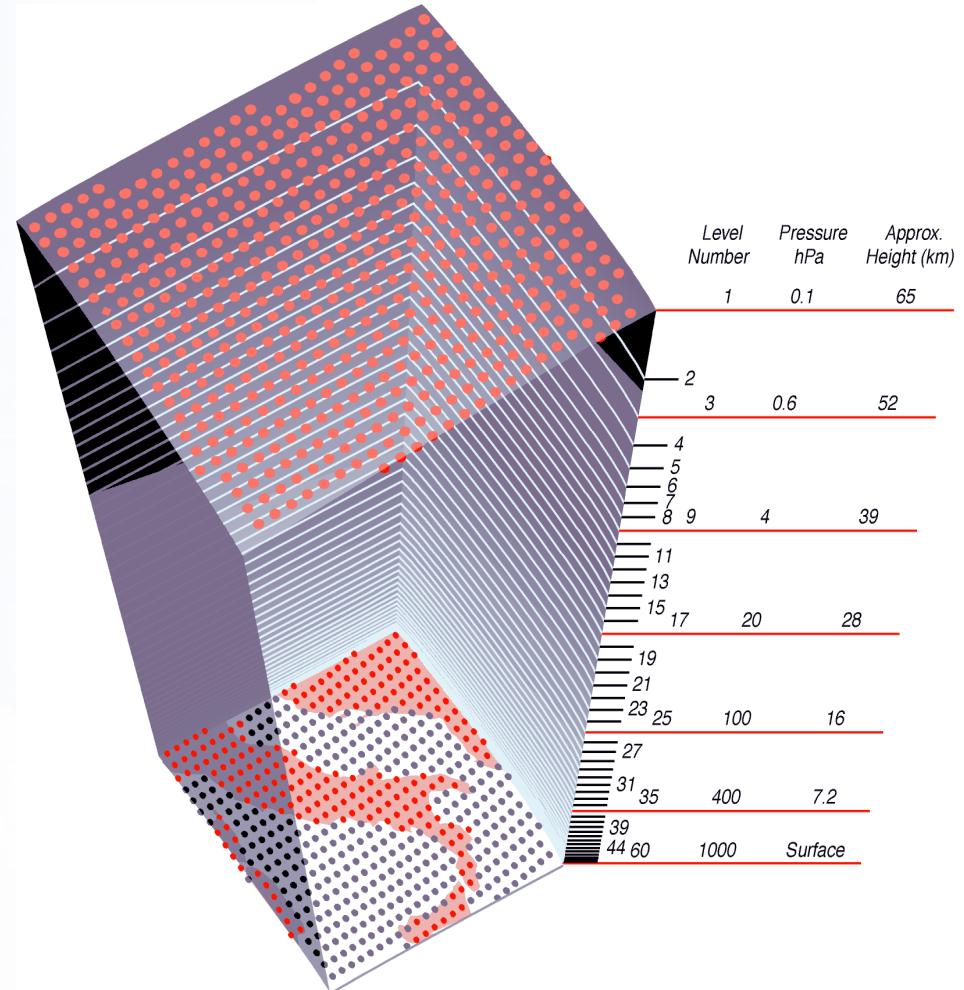
Fields:

2M grid points x 137 levels
287M values, **per variable**

Operational models produce:

- 13 millions fields daily
- Totalling 14 TB/day

Time critical window: 1h !!!



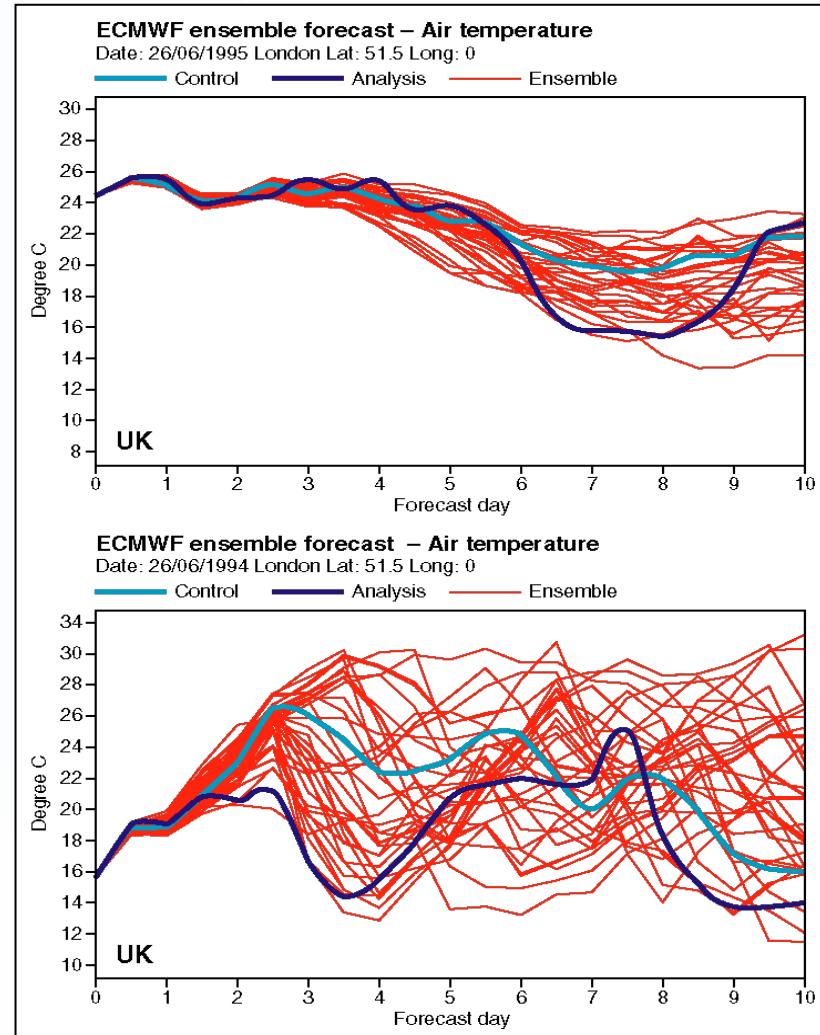
Chaos and weather prediction

The atmosphere is a chaotic system

- Small errors can grow to have major impact (butterfly effect)
- Can never perfectly measure the state of the atmosphere
- Limits weather prediction to a week or so ahead

Ensemble Forecasts

- Parallel set of forecasts from very slightly different initial conditions and model formulation
- Assess uncertainty of today's forecast



Why run an ensemble prediction system?

Basic idea

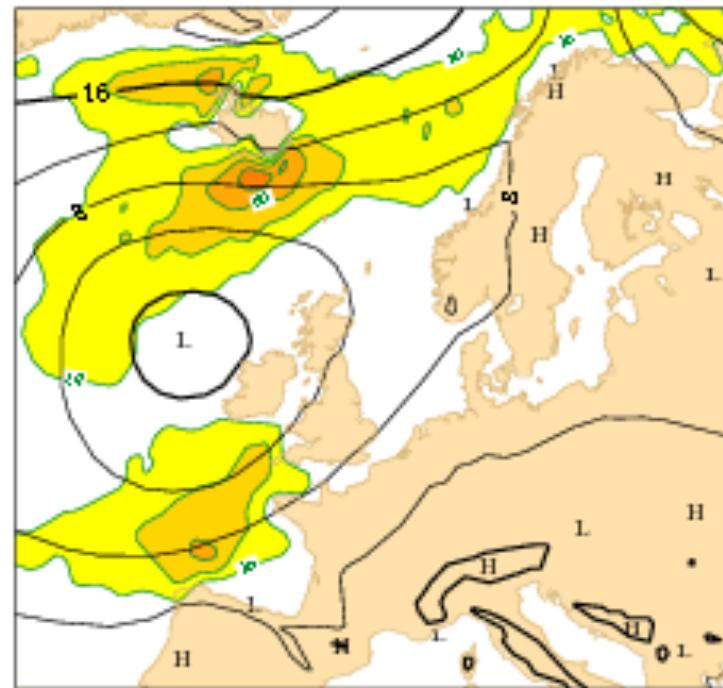
- Taking account of uncertainty
- Forecasting forecast skill

Forecasting benefits

- Assess uncertainty of today's forecast
- Provide alternative forecast scenarios
- Highlight the predictable (large-scale) component and the risk for a less likely but significant (small-scale) event

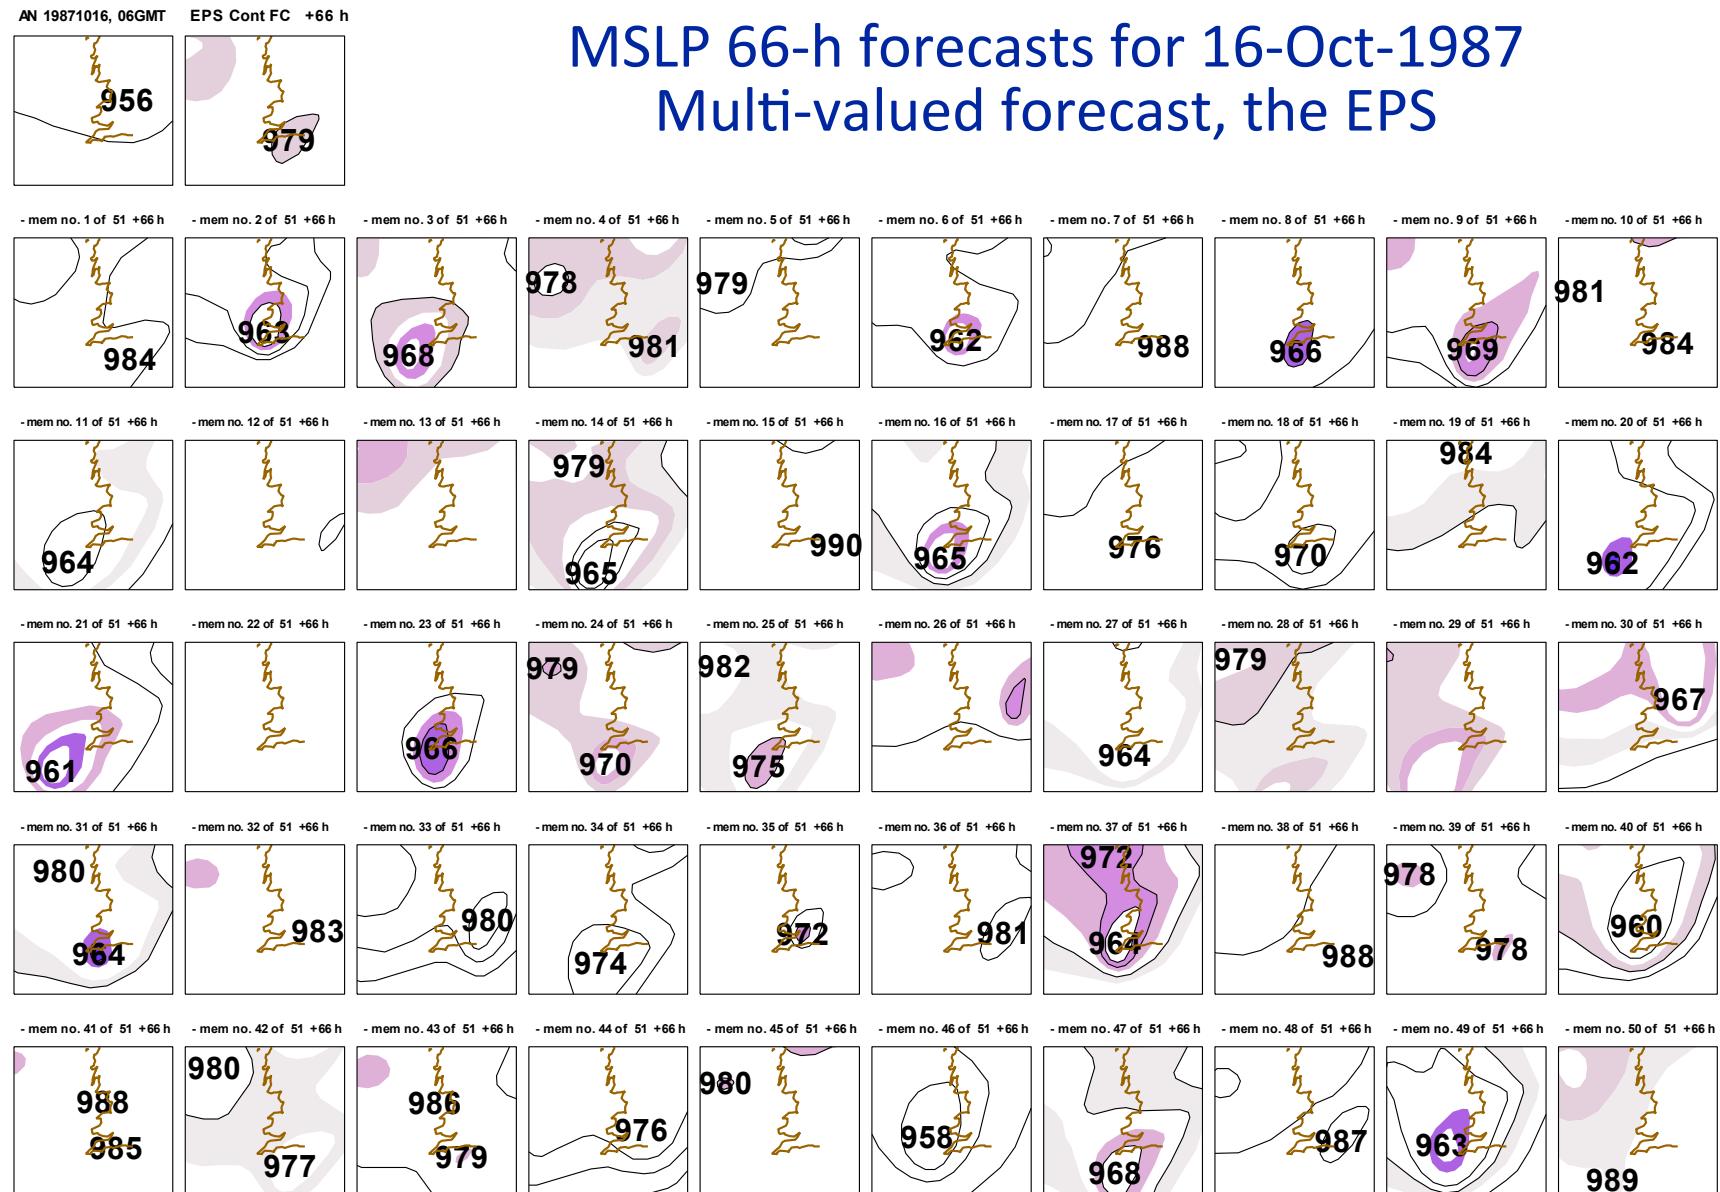
Continuing challenges

- Forecasting extreme events
- Extending the forecast range



MSLP 66-h forecasts for 16-Oct-1987

Multi-valued forecast, the EPS



Product Generation

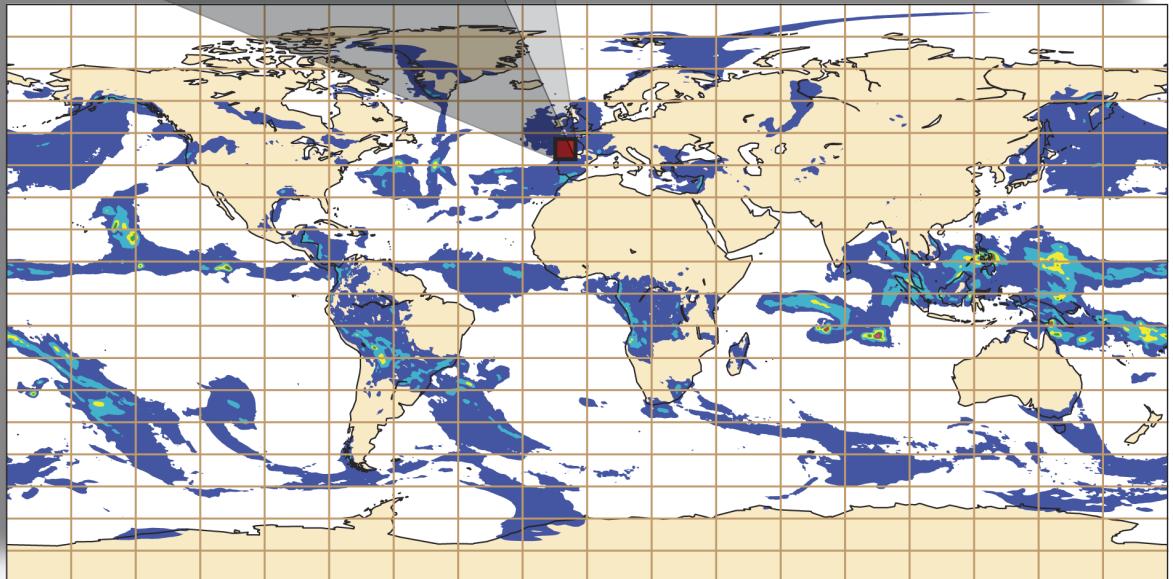
77 million products disseminated ever day, totalling 6 TB.

Interpolate output fields into user required grids

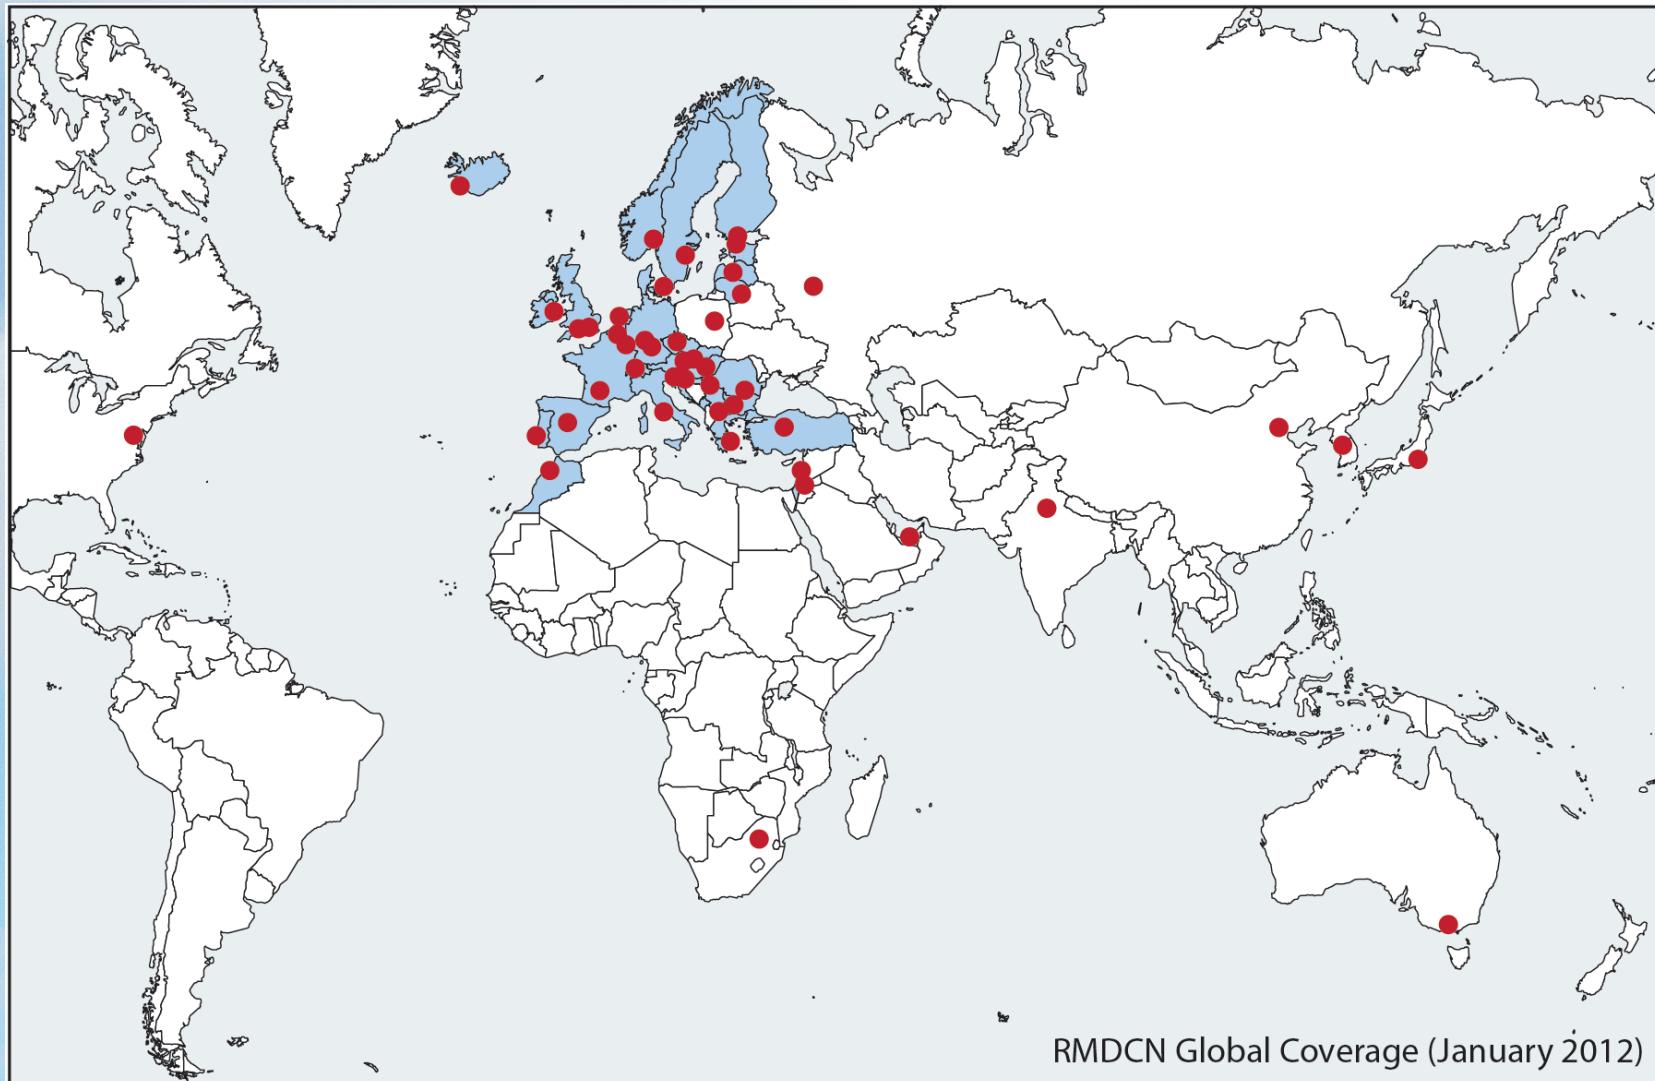
Product generation is also subject to a dissemination schedule...

Time critical window: 1h !!!

7.2	9.9	3.6	0.4	8.3	0.2	0.5	0.1	9.1	6.7
0.3	8.8	1.8	0.5	0.3	0.1	2.7	0.1	7.9	6.9
7.1	9.2	3.6	0.4	8.3	0.2	6.5	3.3	5.5	5.3
2.2	1.1	1.7	0.7	3.5	2.4	0.8	1.9	9.0	6.7
5.1	0.9	1.9	8.9	5.9	0.4	1.5	2.0	7.7	0.7
6.2	0.4	1.4	9.8	9.9	7.7	0.9	3.2	7.2	4.8
8.1	1.4	4.4	0.4	0.3	7.2	3.5	3.4	1.1	9.7
7.0	3.6	4.9	0.7	6.8	1.2	0.1	2.2	6.6	6.0
0.2	7.7	3.6	3.1	8.6	0.5	9.5	0.8	5.6	5.0
3.2	7.2	3.1	0.4	0.9	0.3	0.7	0.4	0.2	0.0



RMDCN Connections





Products also served via web
visualisation services
(demo)

Challenges Ahead

Scalable Algorithms

Platform Uncertainty

Big Data



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

24

New Algorithms Challenge

To perform better forecasts:

- Assimilate better & more observations
- Keep increasing model resolution
- Add modeling of new physical phenomena
- Add more ensemble members

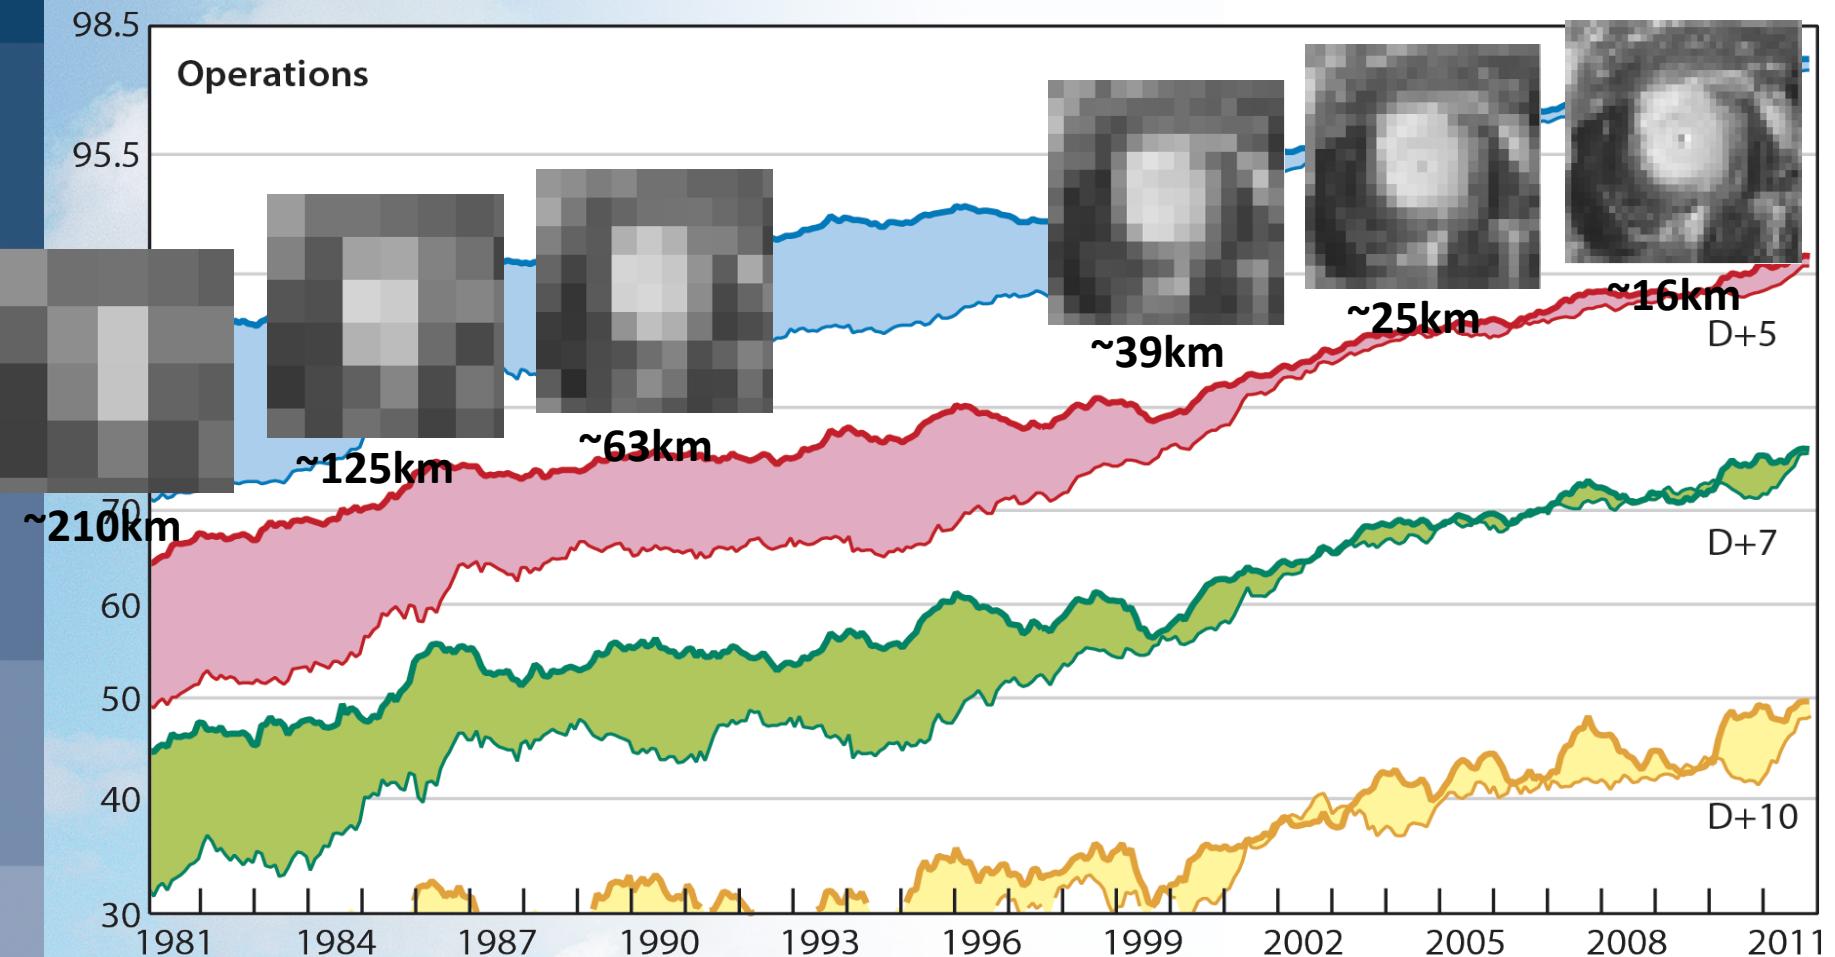
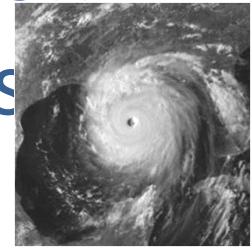
More computations

All run within a critical 1h time window

Evolution of ECMWF scores comparison northern and southern hemispheres

Anomaly correlation of 500 hPa height forecasts

— Northern hemisphere — Southern hemisphere



Benefits of High Resolution



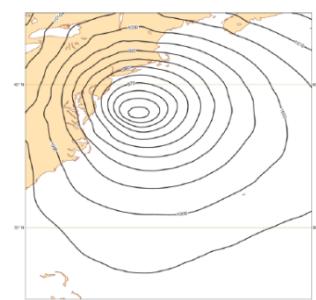
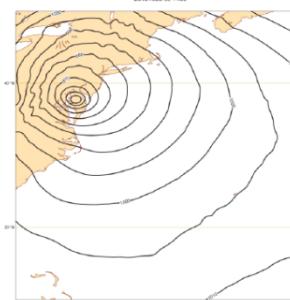
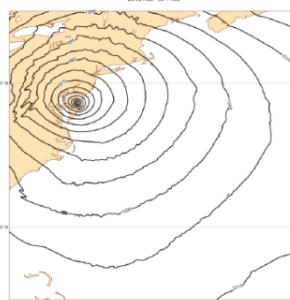
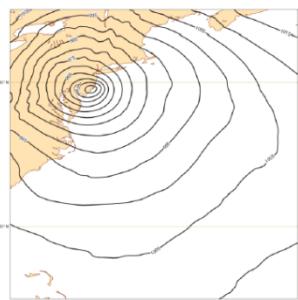
Mean sea-level pressure

AN 30 Oct

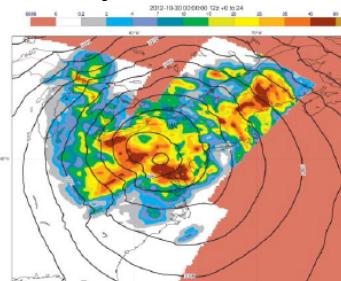
5d FC T3999

5d FC T1279

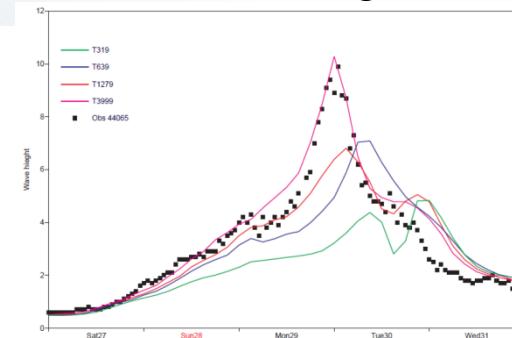
5d FC T639



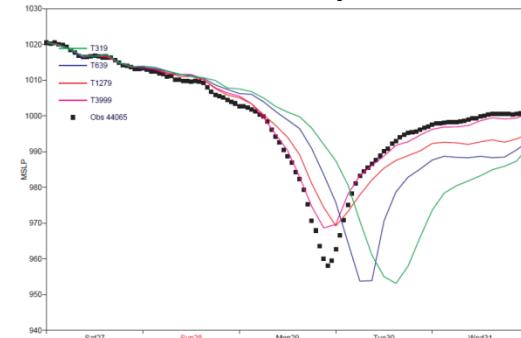
Precipitation: NEXRAD 27 Oct



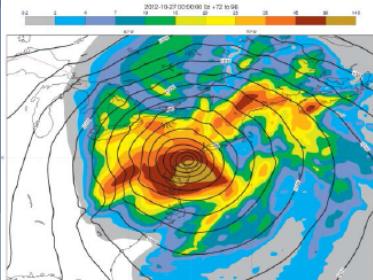
3d FC: Wave height



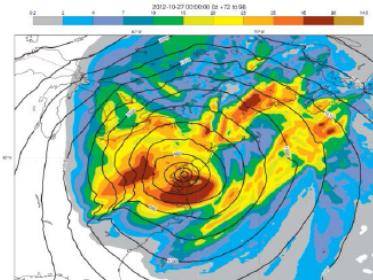
Mean sea-level pressure



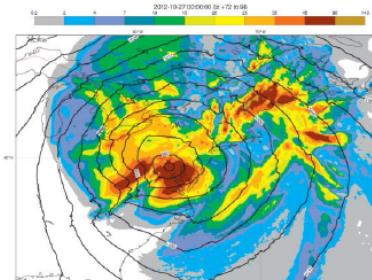
4d FC T639



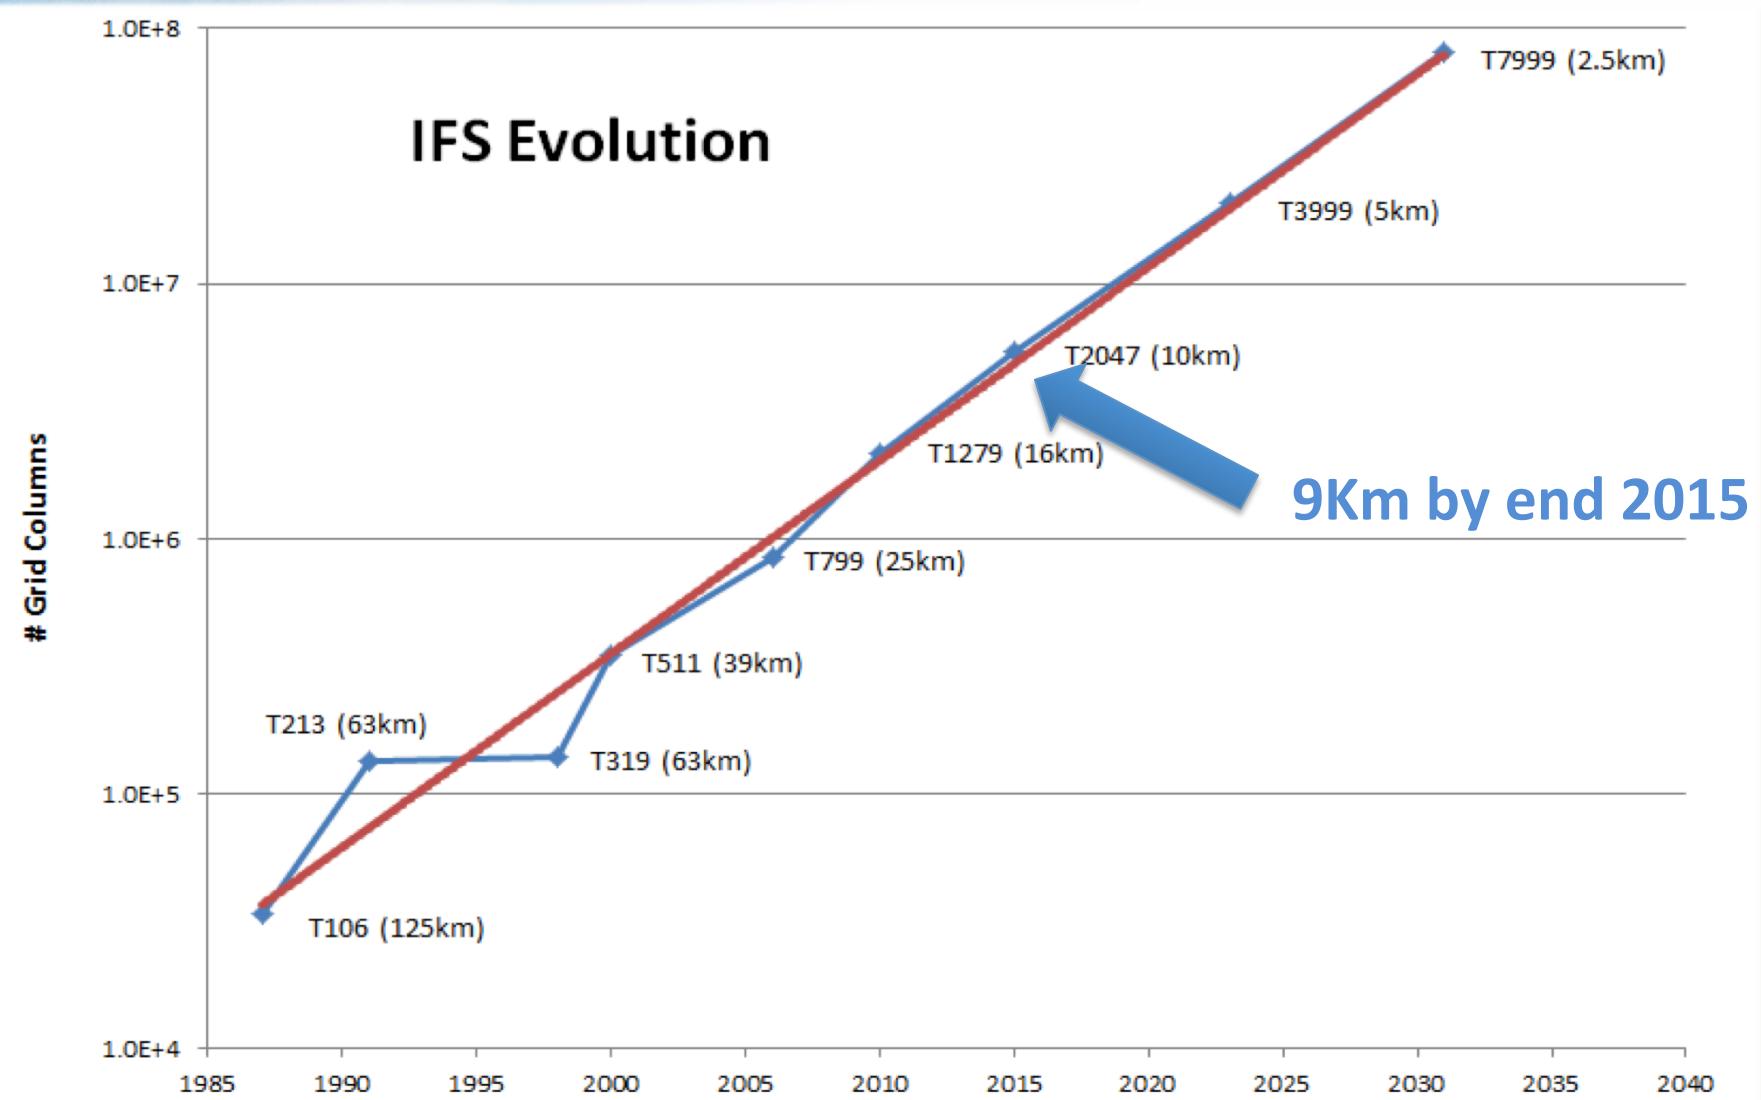
4d FC T1279



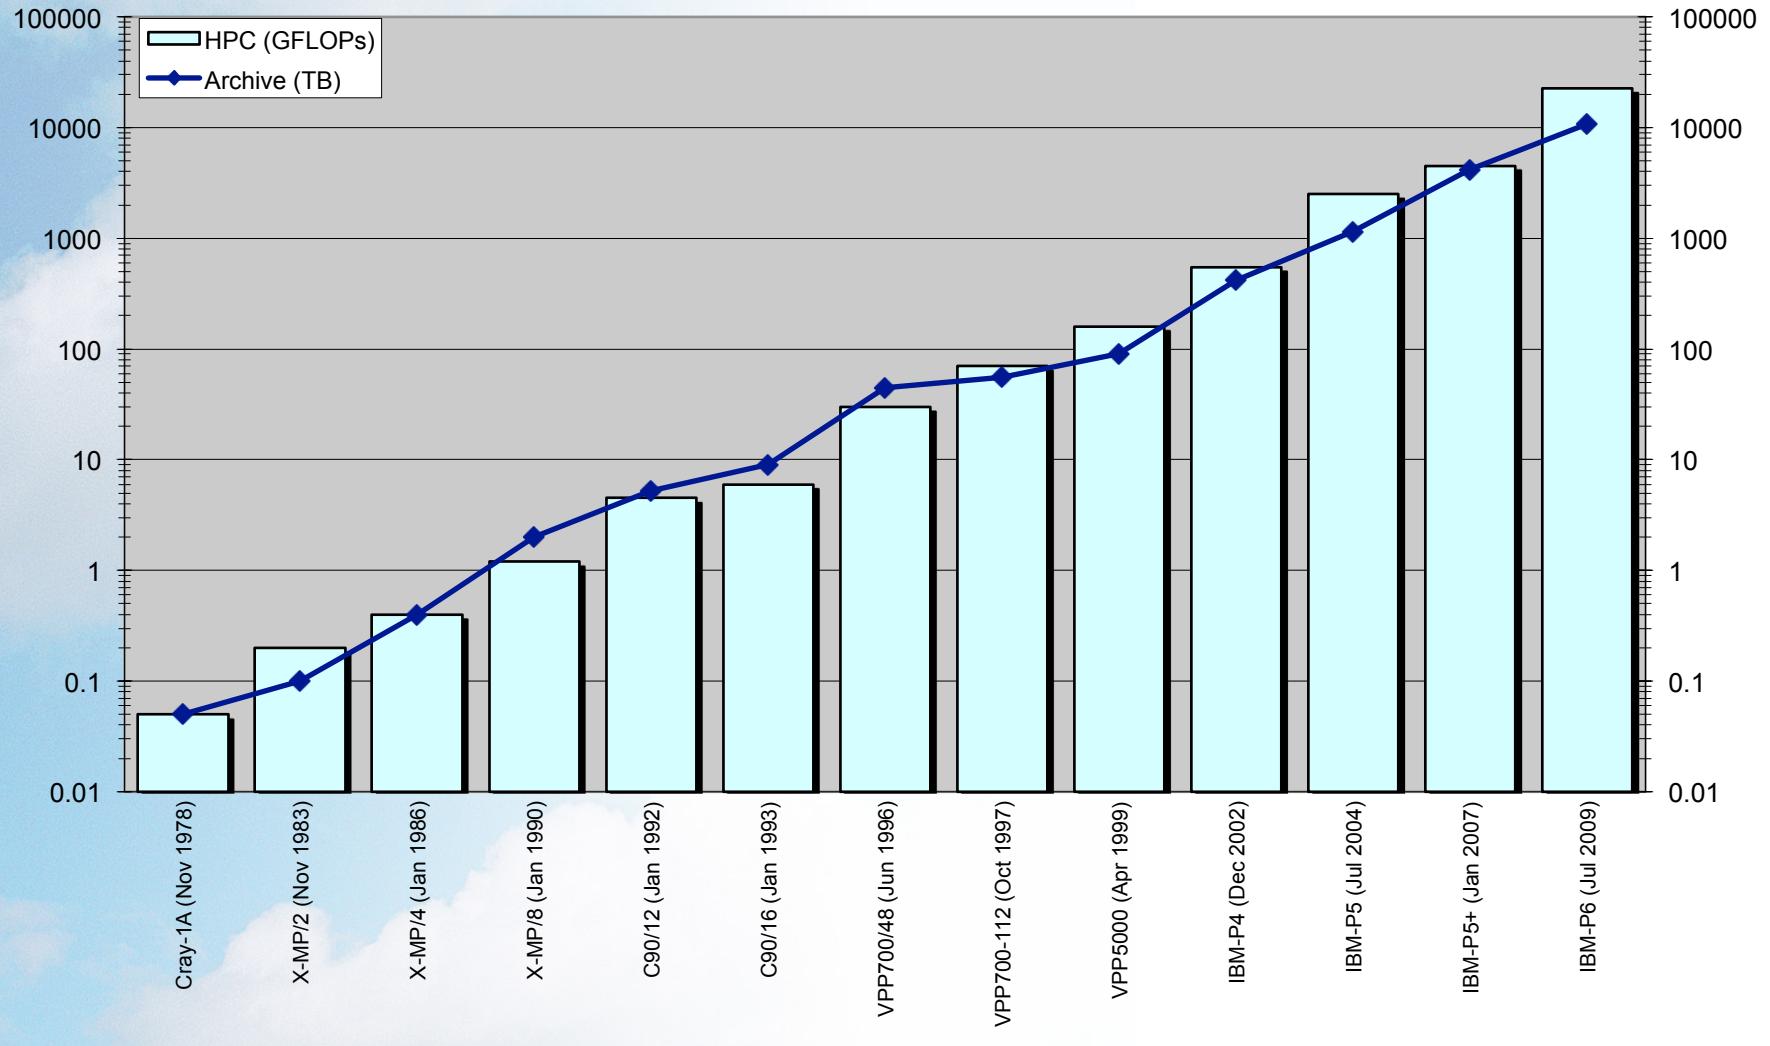
4d FC T3999



Planned Resolution Upgrades



Archive size vs. Supercomputer power

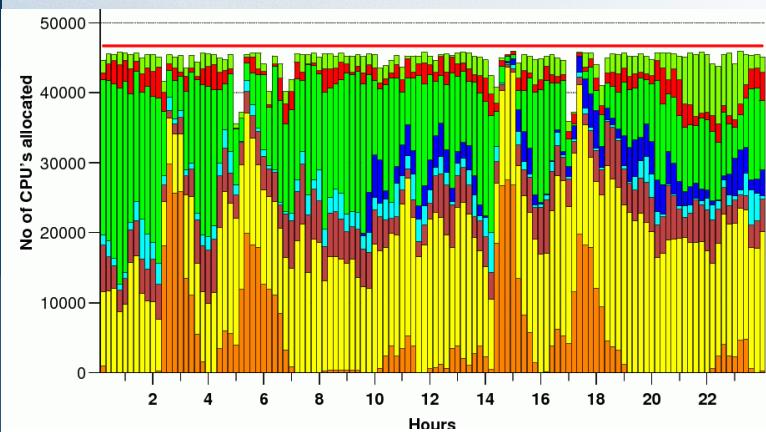


Resolution Upgrades

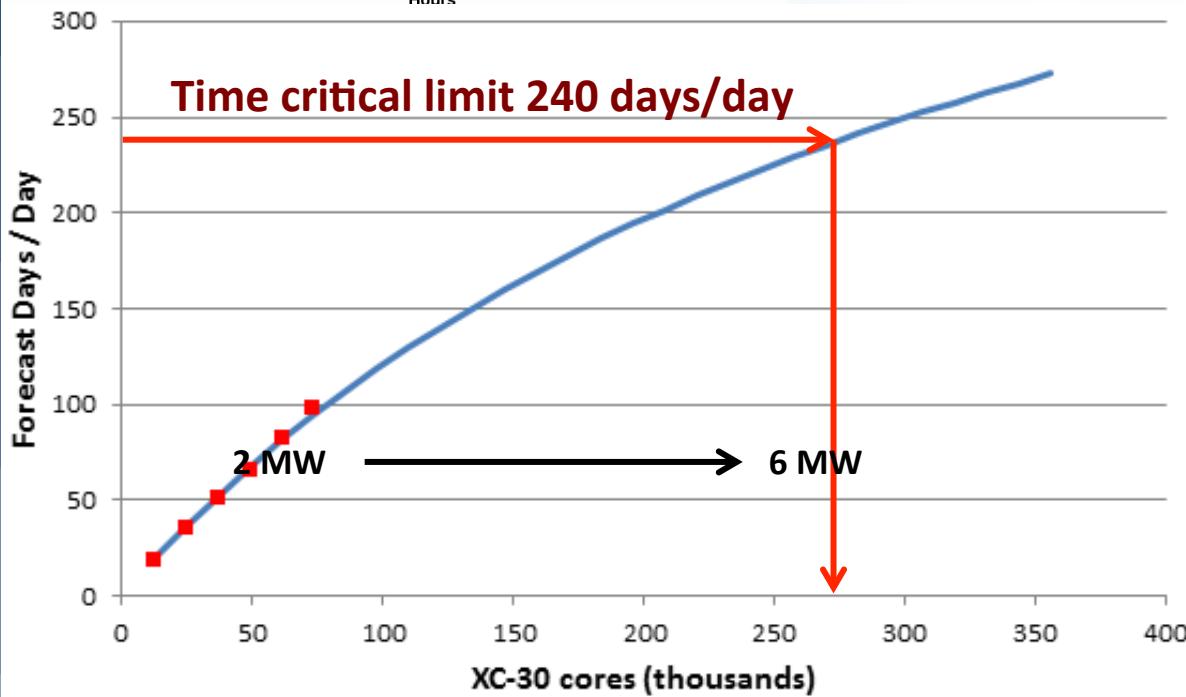
Resolution	Grid size	Grid Points	Field Size (in memory)
T319	62.5 km	204 k	1.6 MB
T511	39 km	524 k	4 MB
T799	25 km	1.2 M	9.6 MB
T1279	16 km	2.1 M	16.8 MB
T2047	10 km	8.4 M	67.2 MB
T3999	5 km	20 M	160 MB
T7999	2.5 km	80 M	640 MB

As memory per core diminishes (think GPU's and Intel Phi) ...
... this may have serious implications on the data processing software!

Power Challenge



- 25% of HPC used for operations
- 10% of operations for HRES forecast
- 2.5% of HPC



- 2.5 km* forecast on Cray XC-30 (84,000 cores)
= 2 MW
- 240 days/day require 270,000 cores
= 6 MW
- $\times 10^{**} = 60 \text{ MW ?}$

* planned for 2025
** scaling from HRES to full HPC

Platform Uncertainty Challenge

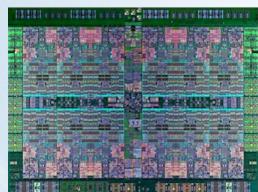
- How to keep up with the **changing** hardware landscape?
- How to **best** use the new heterogeneous systems?
 - CPU's
 - GPU's
 - Accelerators
- With very limited human resources...

Hardware Landscape

Proliferation of HPC Architectures



Intel Xeon



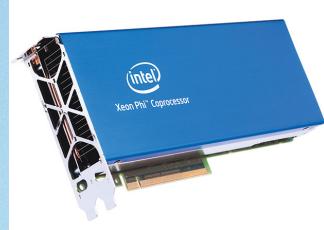
IBM Power8



NVidia Tesla



AMD FirePro



Intel Xeon Phi



NEC SX-ACE Aurora



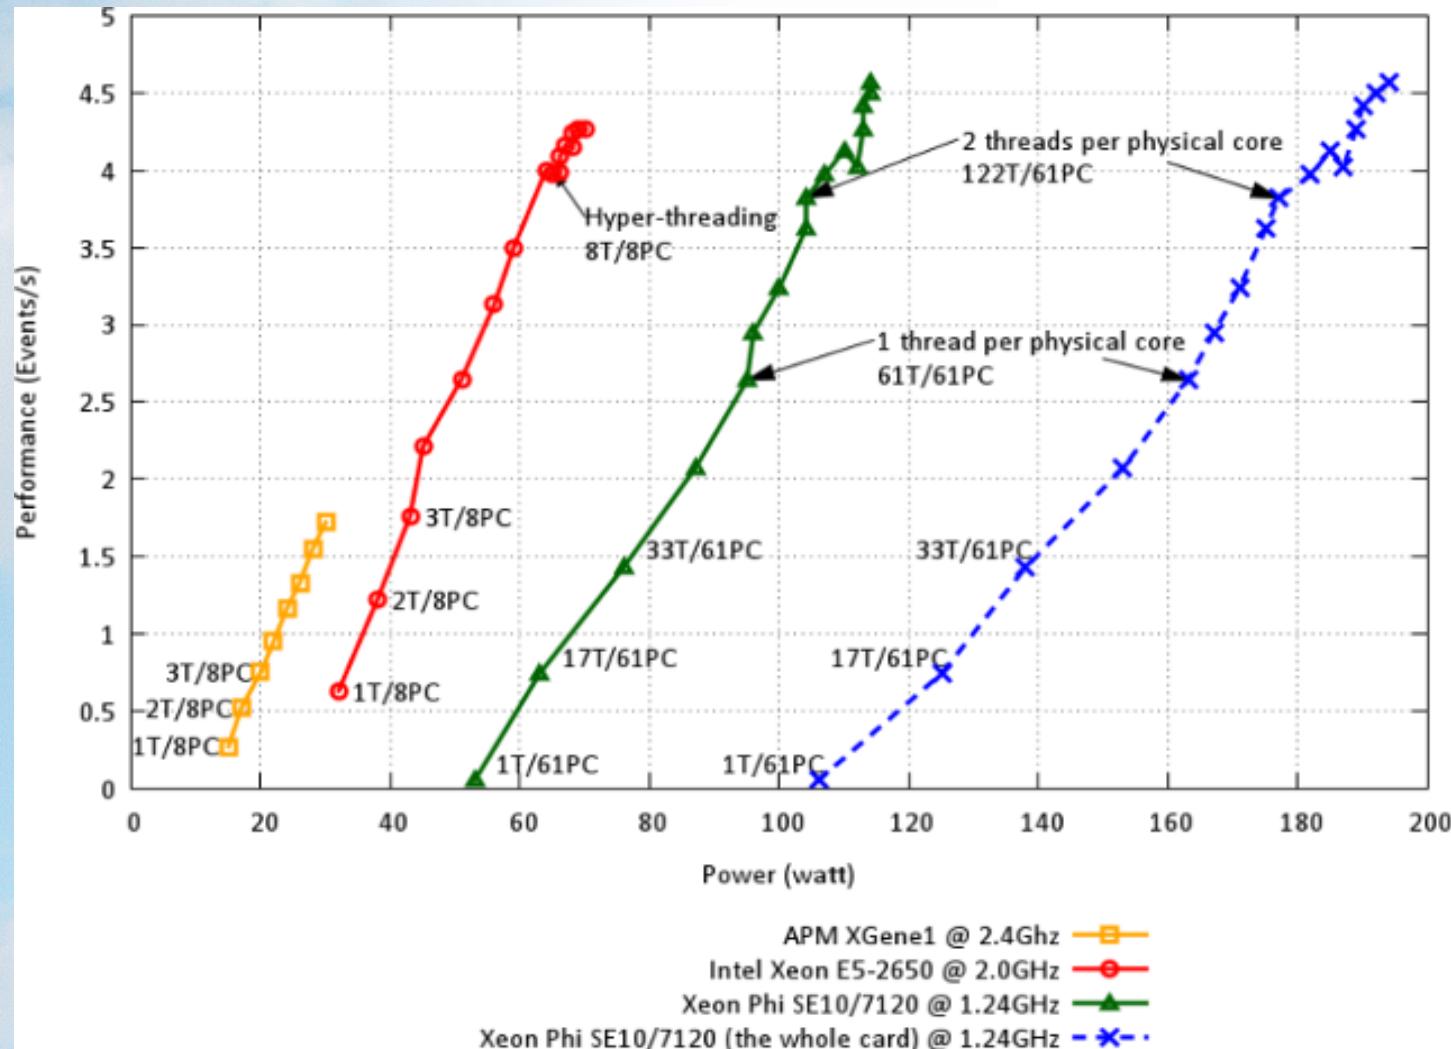
ThunderX ARM



X-Gene ARM

Power Matters

"Heterogeneous High Throughput Scientific Computing with APM X-Gene and Intel Xeon Phi", Adburachmanov et. al., 2014



Hypothetical Solution

- Need flexibility?
- Need performance?
- Code easier to maintain?
- A modern language?

We have a solution...

C++



But ...



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

36

HPC Platform Restrictions

- **POSIX**
- Only supported languages:
 - **Fortran**
 - **C99**
 - Recently **C++ (but not C++11)**
 - Shell (Ksh) + Perl (**but not Python**)
- Most of our scientists know
 - **Fortran**
 - Python (maybe)
 - Basic shell scripting

C++ Support on HPC

- Support for C++ in HPC platforms is ... **patchy!**
- Some HPC compilers aren't fully C++98 and choke on heavily template code:
 - Eigen library for linear algebra
 - Parts of Boost
 - serialization
 - Proto
- Not to mention **C++11!**

Change on the way?

Community & Vendors start talking about C++

Scientists are reluctant to use C++

Possible explanations (my view):

- Fortran is what they know
- Fortran is very fast for numeric computations
- C++ was (initially) perceived as slow:
 - hidden constructors
 - misuse of classes
 - **restrict** keyword
- Mostly corrected now (e.g. work of T. Veldhuizen)
- But C++ is still perceived as ...
 - complex (templates are a common complaint)
 - less portable (than C or Fortran)

Big Data Challenge

“Big Data is **high volume**, **high velocity**, and/or **high variety** information assets that require new forms of processing to enable enhanced decision making, insight discovery and process optimization.”

“3D Data Management: Controlling Data Volume, Velocity and Variety”, D. Laney, Gartner, 2001

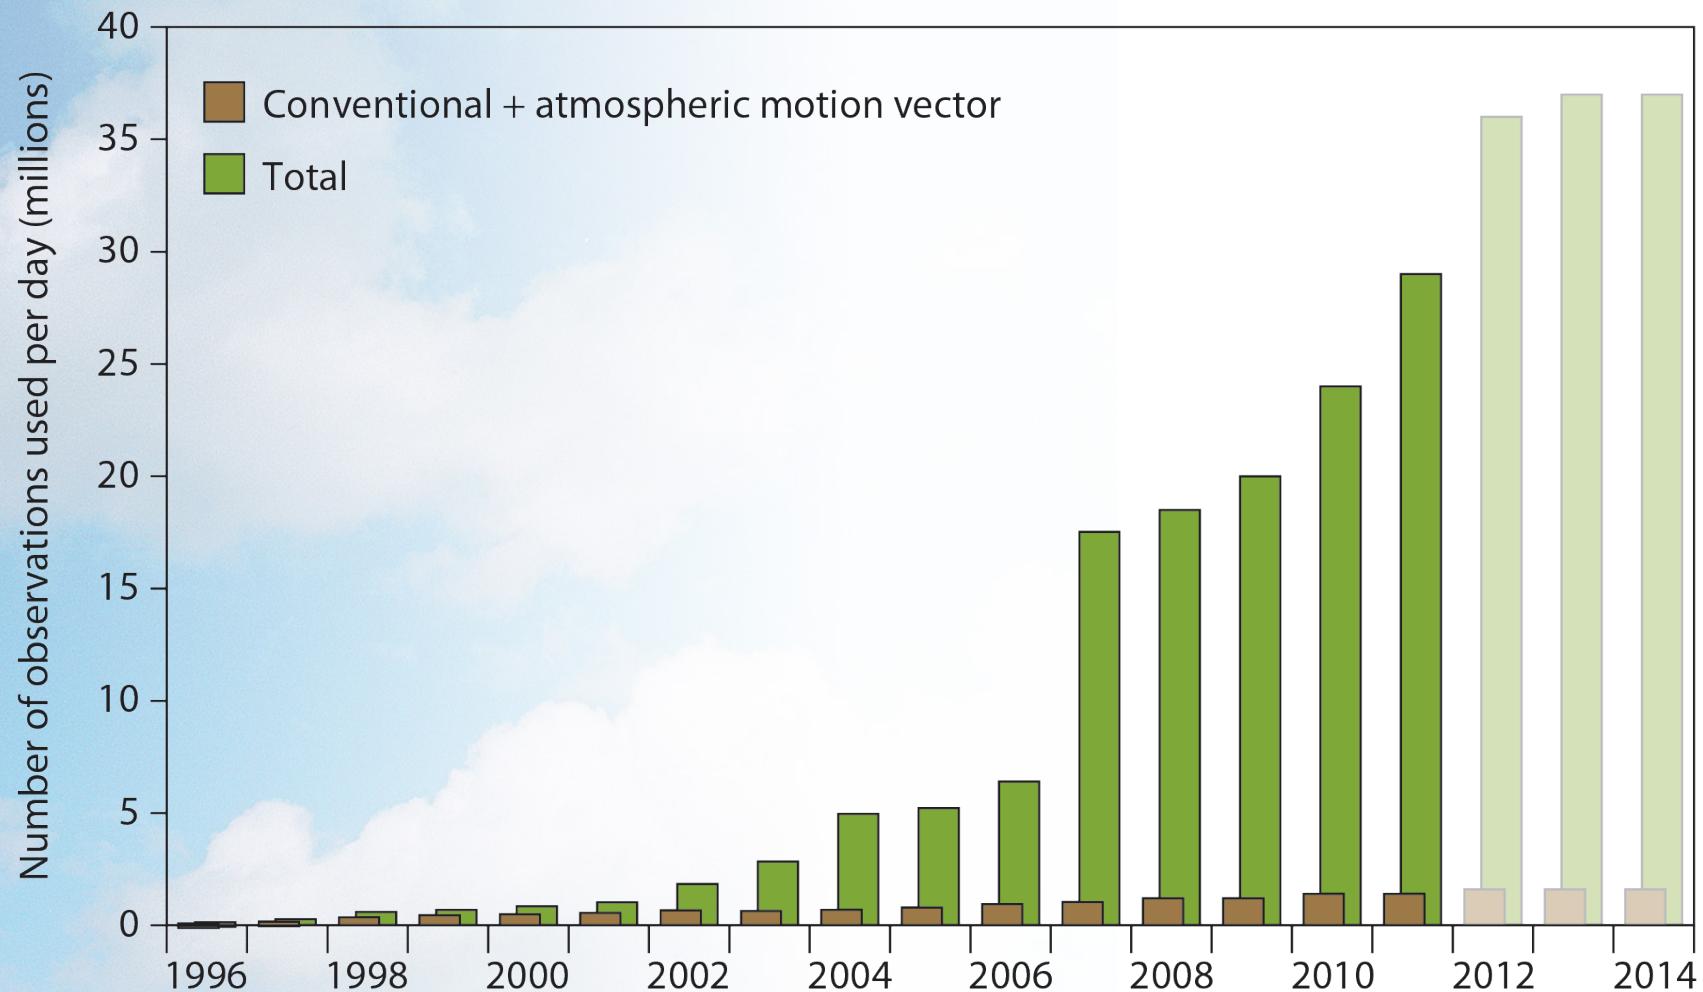
The 3 V's of Big Data



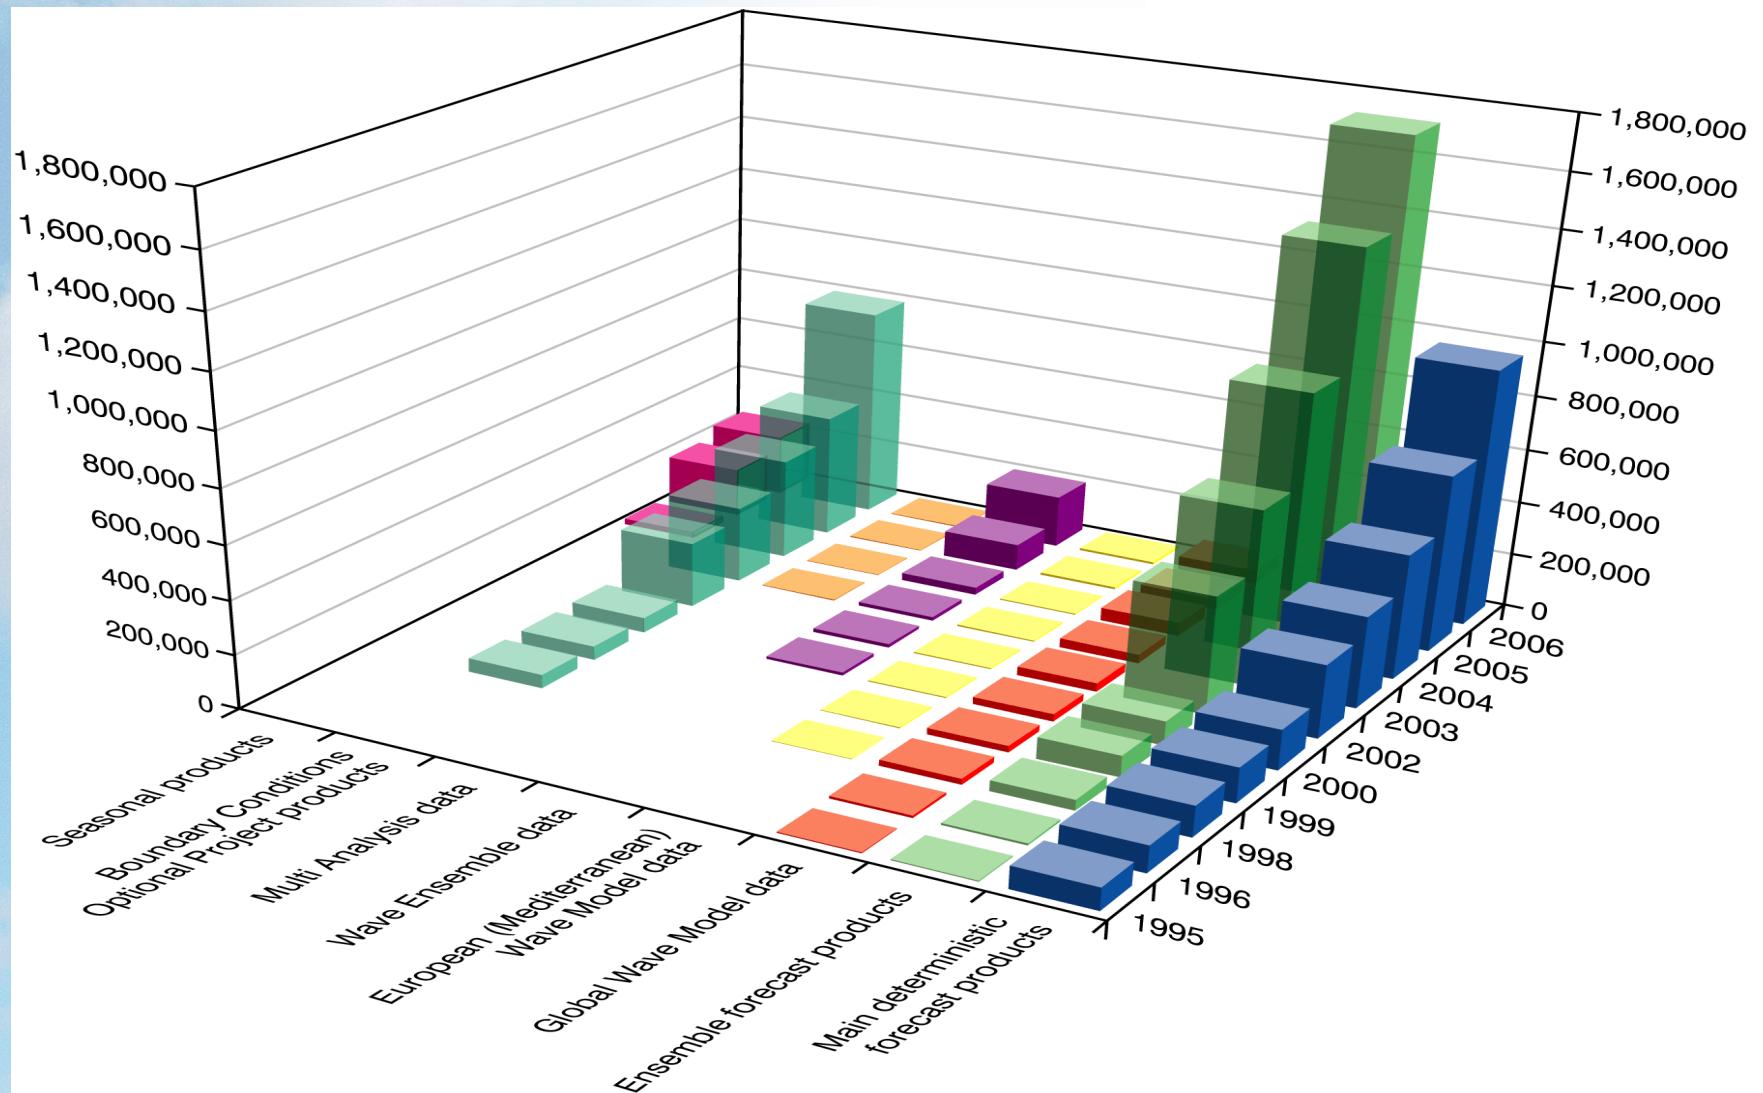
EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

V is for Volume: Observations

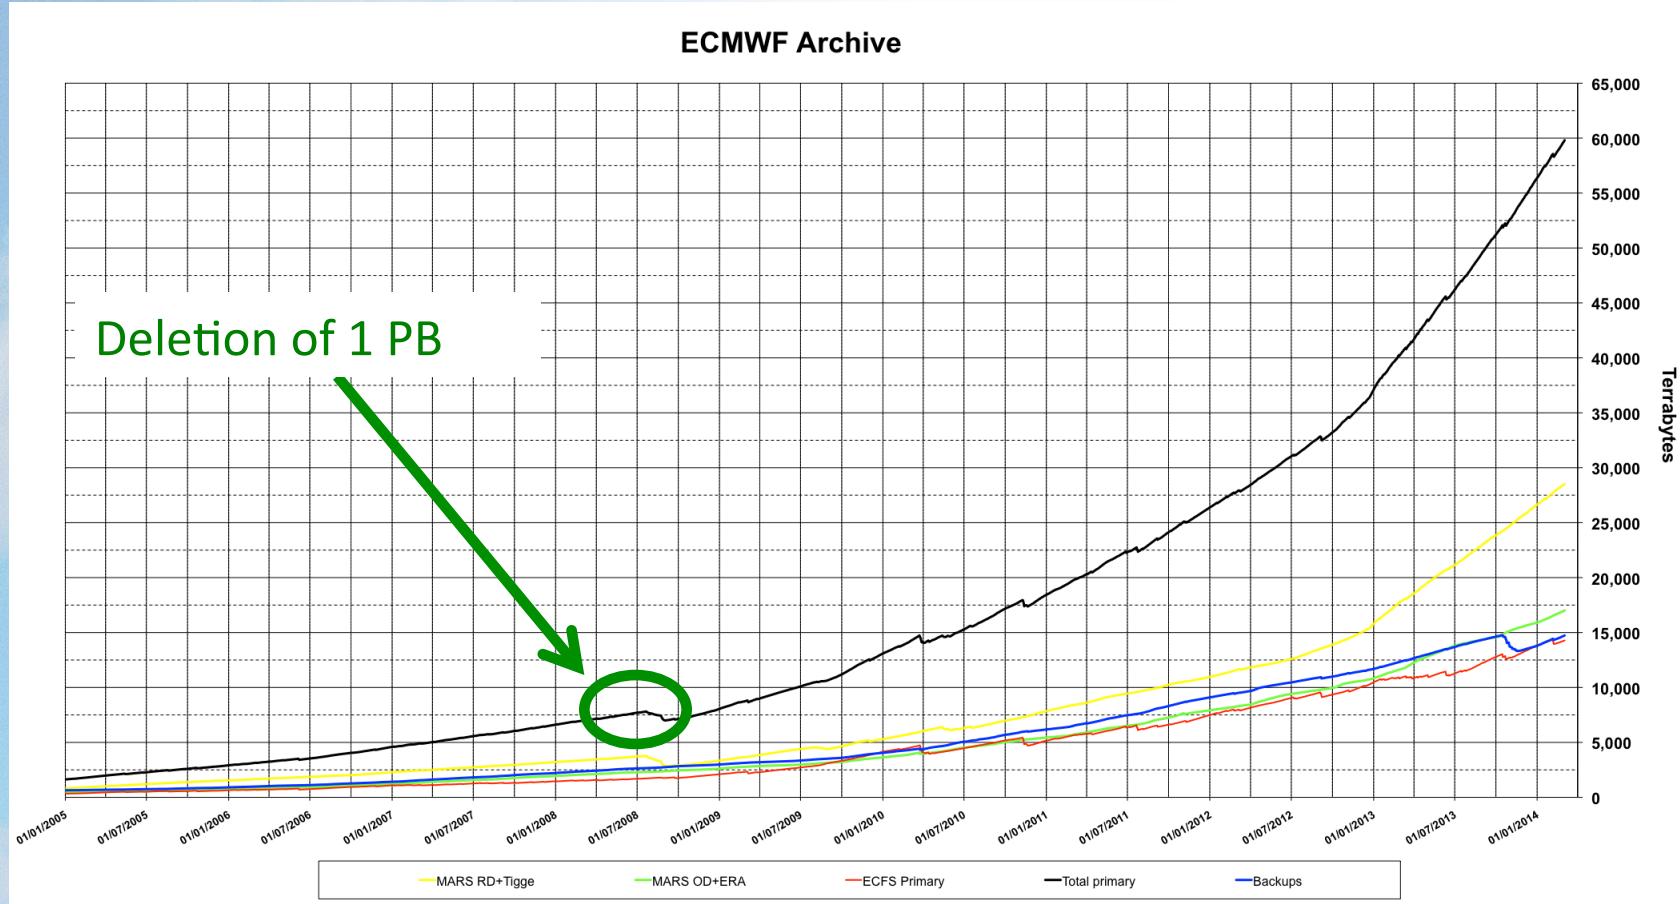
Increase of satellite data usage



V is for Volume: Products



V is for Volume: Archive



V is for Velocity

- ECMWF's archive grows exponentially:

$$V = V_0(1 + r)^t$$

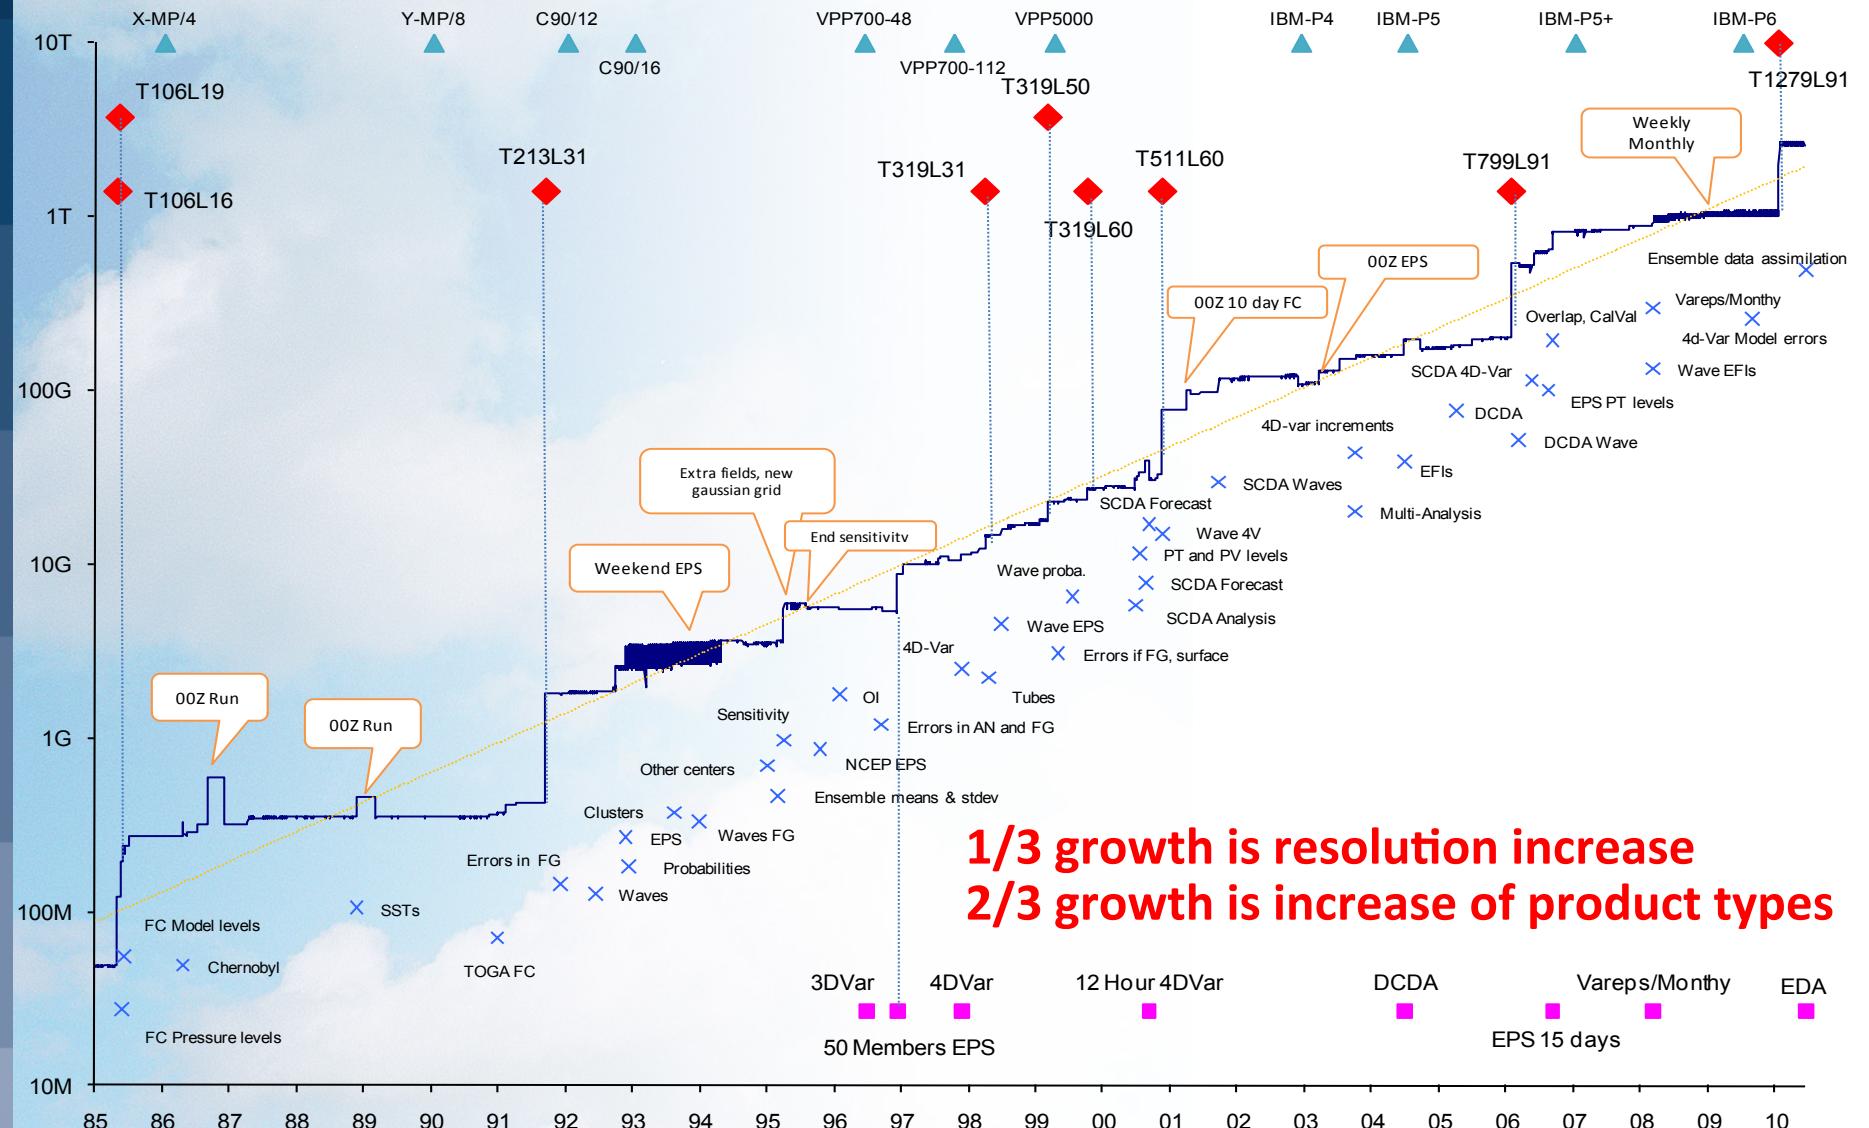
Initial volume *Time*
Volume of the archive *Rate of growth*

- r is around 0.5, which is a 50% increase per year
- The rate of added data also grows exponentially at the **same rate!**

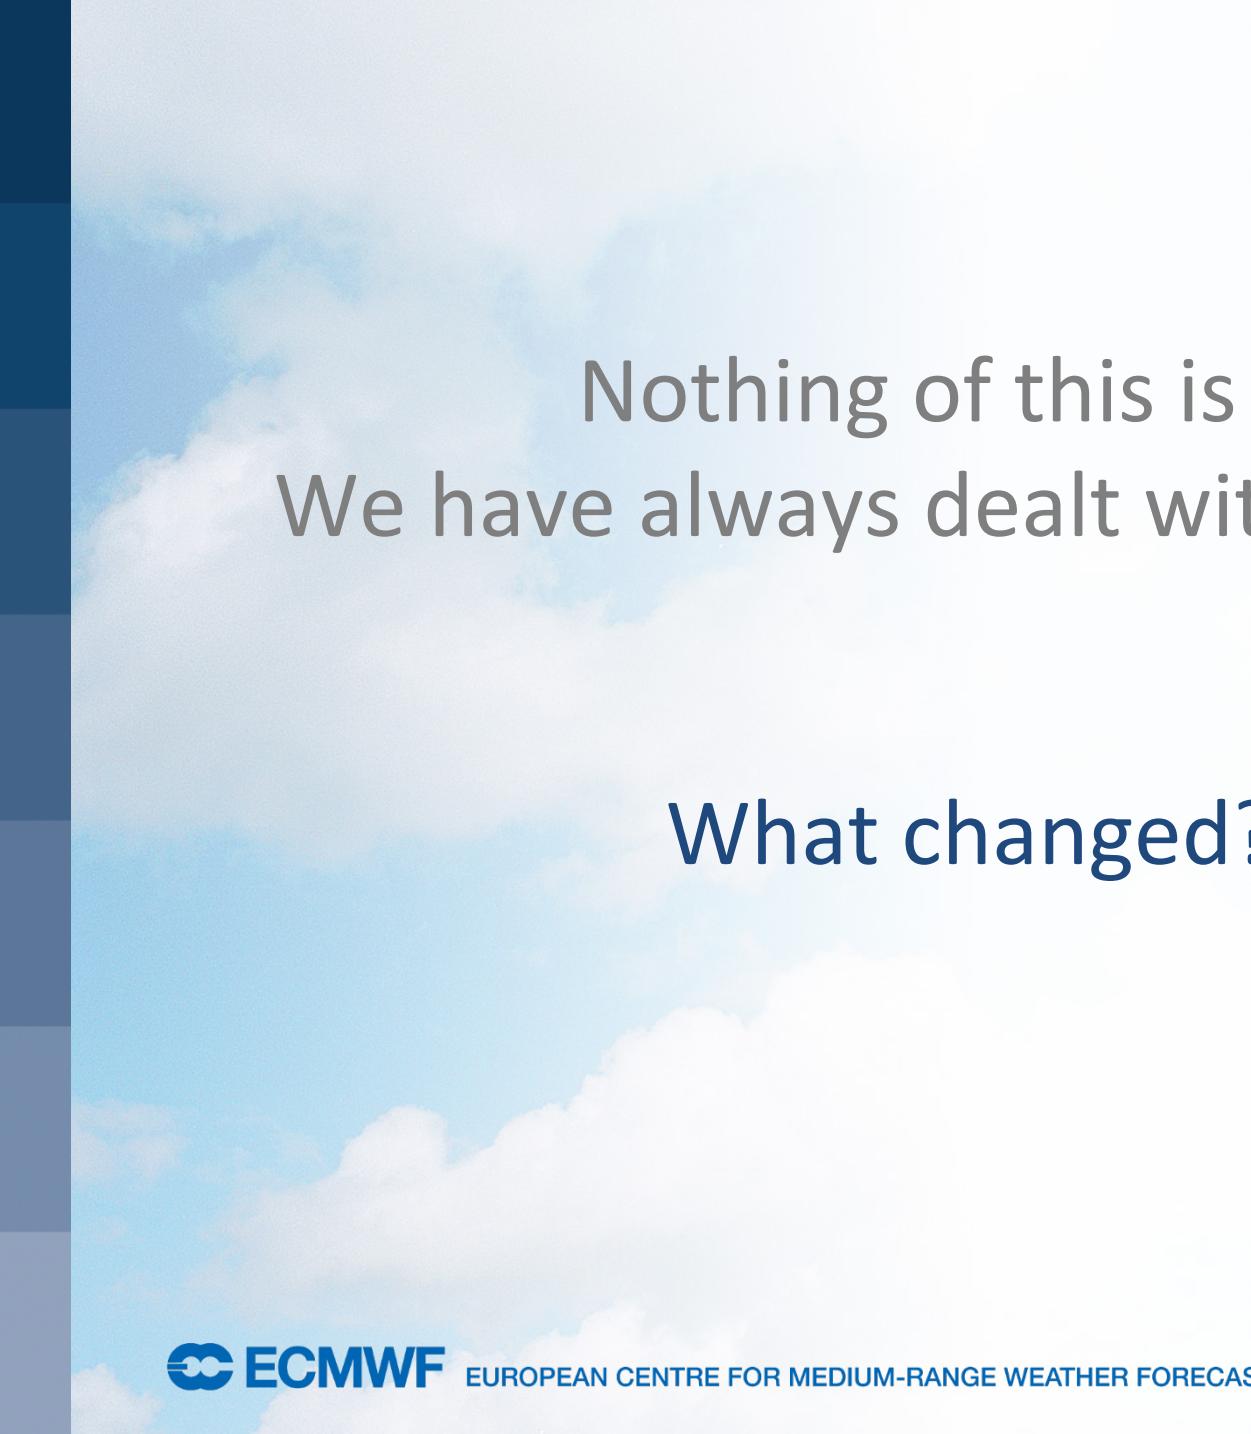
$$\frac{\partial V_0(1 + r)^t}{\partial t} = V_0 \log(1 + r)(1 + r)^t = A_0(1 + r)^t$$

- In 1995, the size of the archive was increasing at a rate of 14 TB/year.
- In 2015, the size of the archive increases at a rate of 100 TB/day

V is for Variety



**1/3 growth is resolution increase
2/3 growth is increase of product types**



Nothing of this is new
We have always dealt with Big Data...

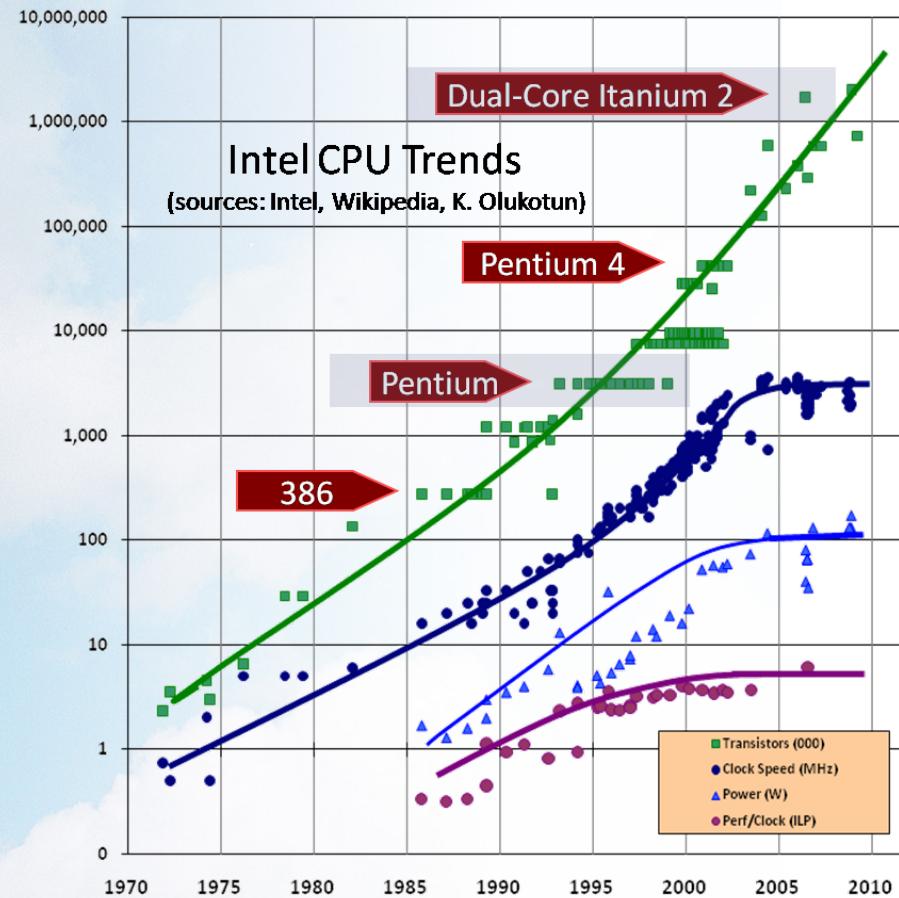
What changed?



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

CPU Power Growth

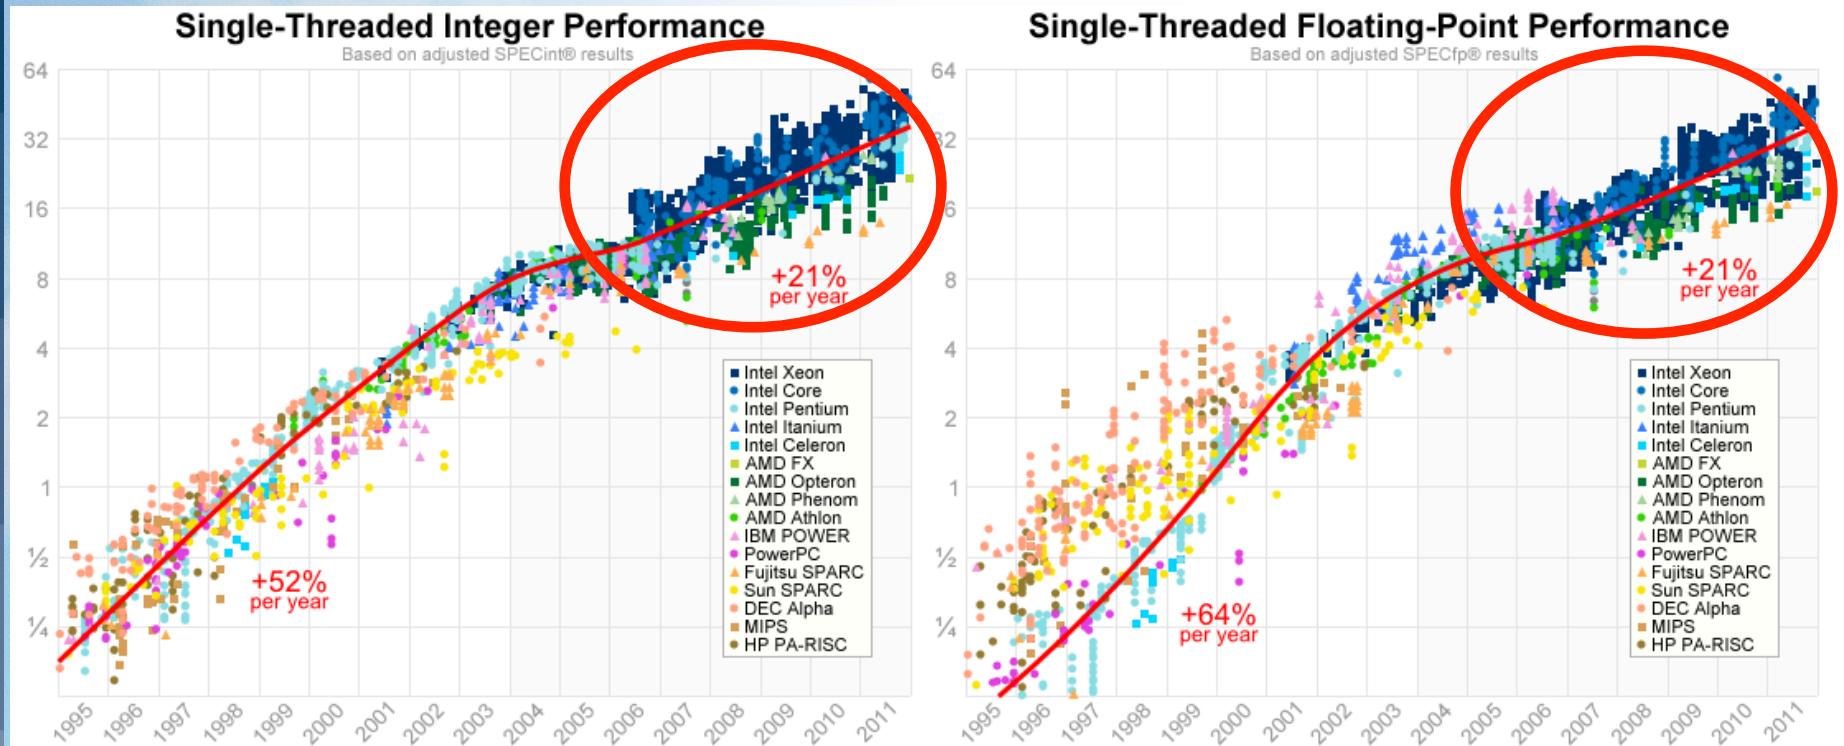
"The Free Lunch is Over". H. Sutter, Dr. Dobb's Journal, 30(3), March 2005



But what about “real” performance?

CPU Performance Growth (single-threaded)

"A Look Back at Single-Threaded CPU Performance", J. Pershing Feb 2012

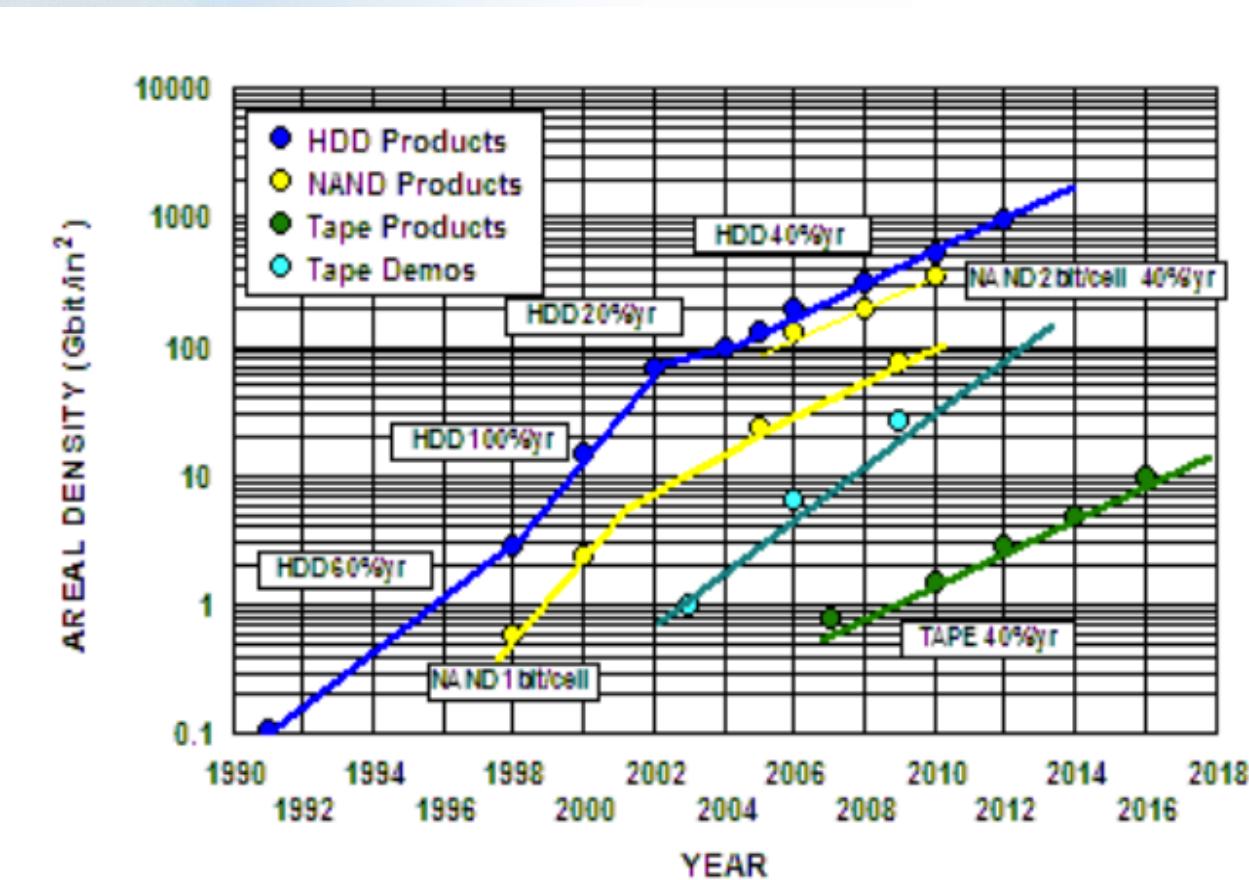


More registers, vector units, branch prediction ...
but also harder to achieve!

Deeper memory hierarchy, out-of-order execution ...

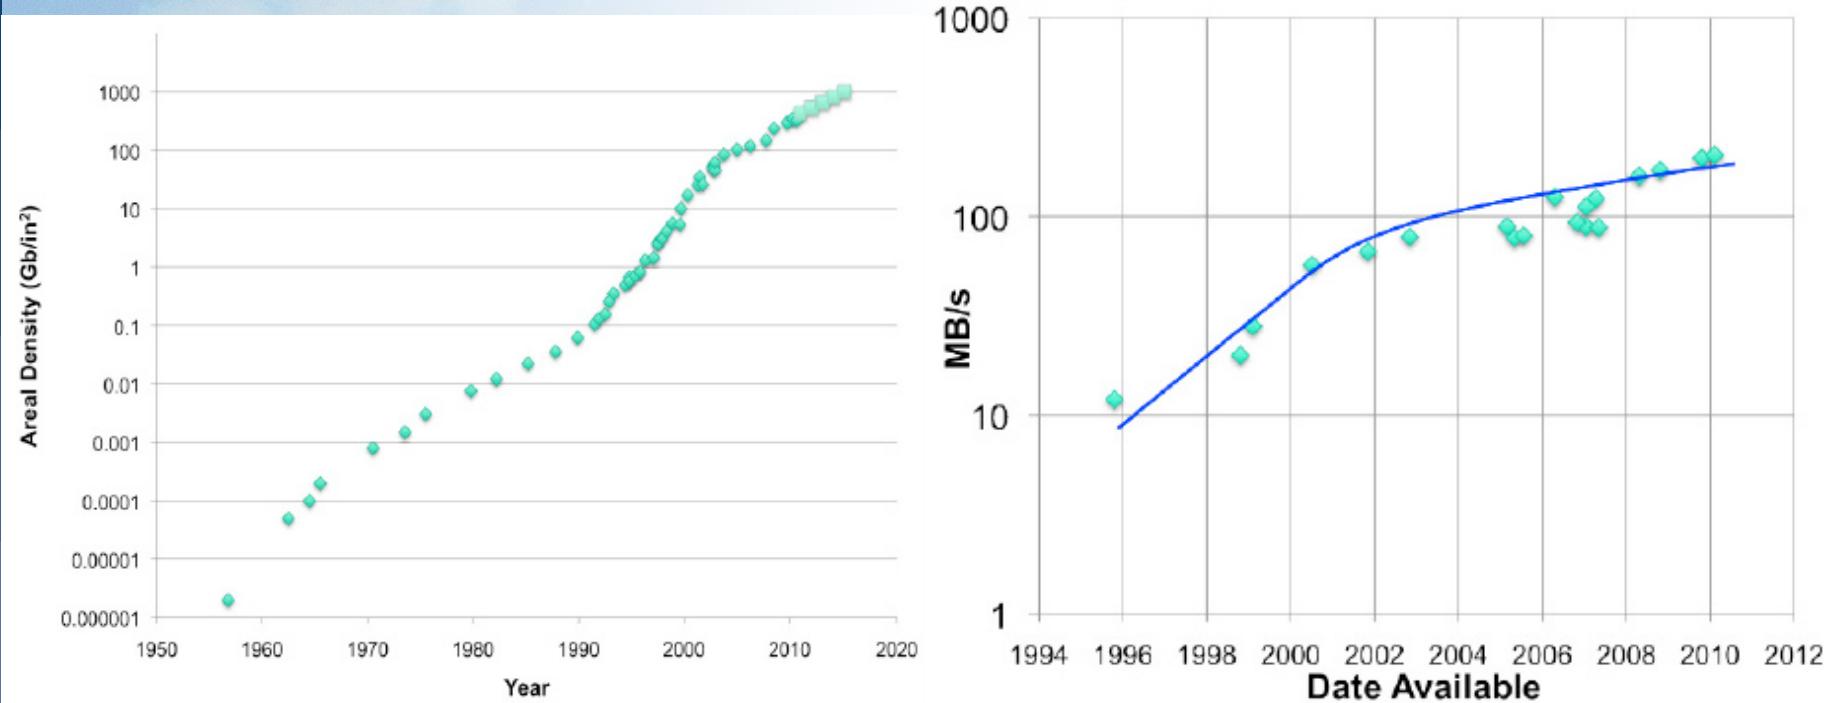
Storage Density Growth Multiple Technologies

*"Tape based magnetic recording: technology landscape comparisons with hard disk drive and flash roadmaps",
R. Fontana et al, IBM Research Division, 2011*



HDD Storage Growth

"GPFS Scans 10 Billion Files in 43 minutes". R. Freitas, et al. IBM Research Division, 2011



Volume is linearly proportional to area density

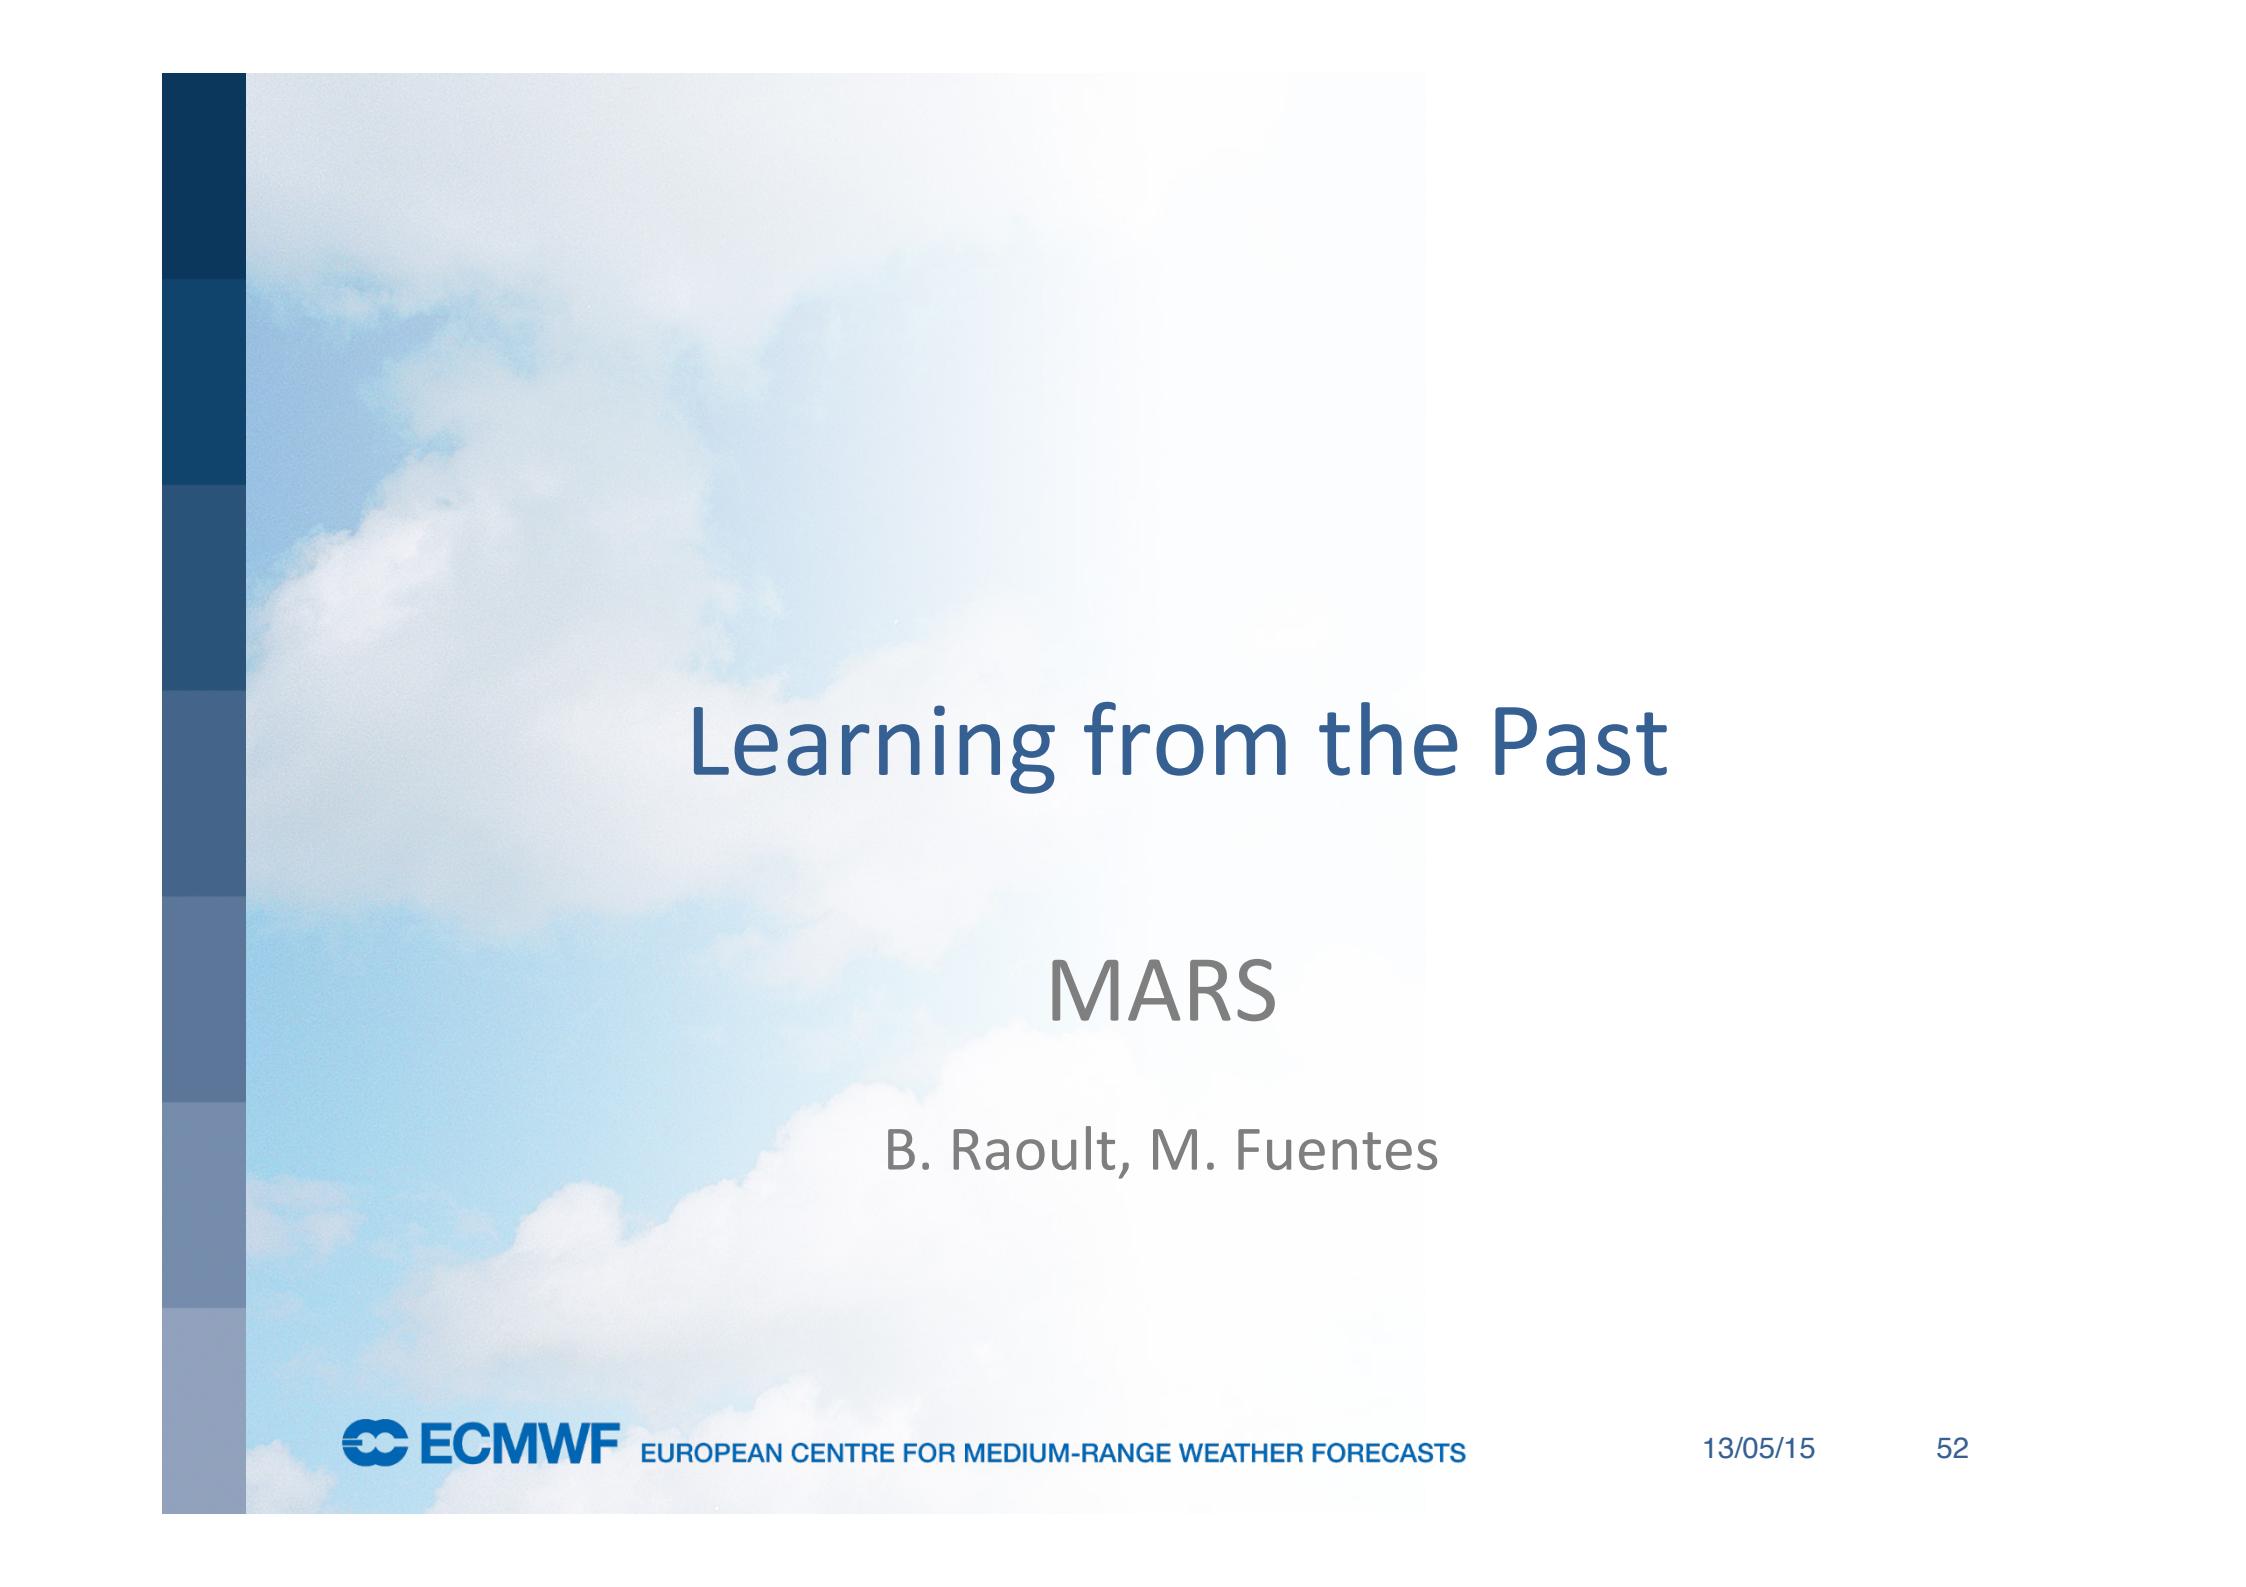
Recently follows 25-40% CAGR...

... but transaction rate hasn't kept up!

This means that we may have the capacity, but maybe not the bandwidth ...

What does it imply?

- “**No Free Lunch**” → Improve our software
 - NWP models have been parallel and concurrent for many decades
- Optimize **Data Movement**
- Develop new **Algorithms** that explore...
 - (More) Concurrent computations
 - Data locality
 - Computational intensity (CPU usage/MB transferred)
- Software must cope with changes – **Flexibility**
 - Best use of new hardware
 - Unknown future for parallel platforms



Learning from the Past

MARS

B. Raoult, M. Fuentes



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

52

Meteorological Archival and Retrieval System

A managed archive, **not a file system**

- Users not aware of data location
- Express access in meteorological terms

Data is kept **forever**:

- As data is accumulated, datasets become more useful
- Deleting old data in an exponentially growing archive is meaningless

Consists of 3 layers:

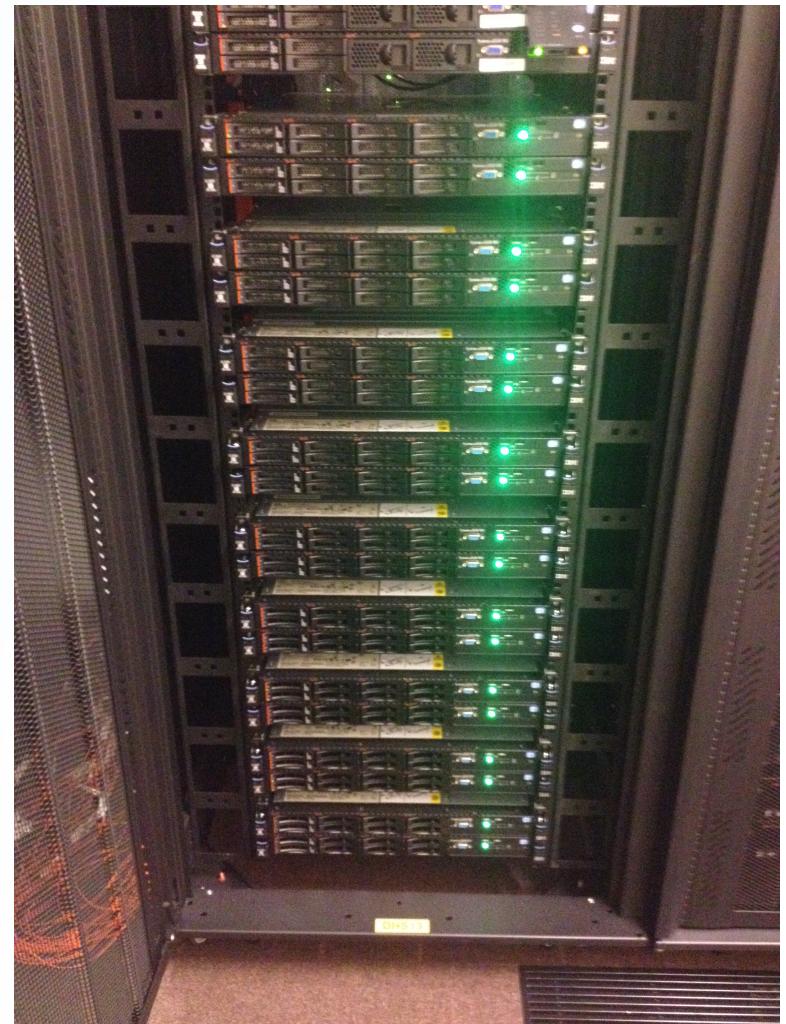
- Cache in the HPC (~80% hit ratio)
- Servers HDD cache (~80% hit ratio)
- HPSS Tape system



Meteorological Archival and Retrieval System

Fully distributed

- 15 servers for metadata and data movers
- 80 PB primary archive
- 1.5 PB of disk cache (2.5%)
- 110 billion fields in 8.5 million files (2014)
- 200 million objects (2014)
- 100 TB added daily
- 100 TB retrieved per day
- 2 million requests per day
- Fully written in C++
- Since 1995
- Index is Object Oriented DB
- Persisted in filesystem

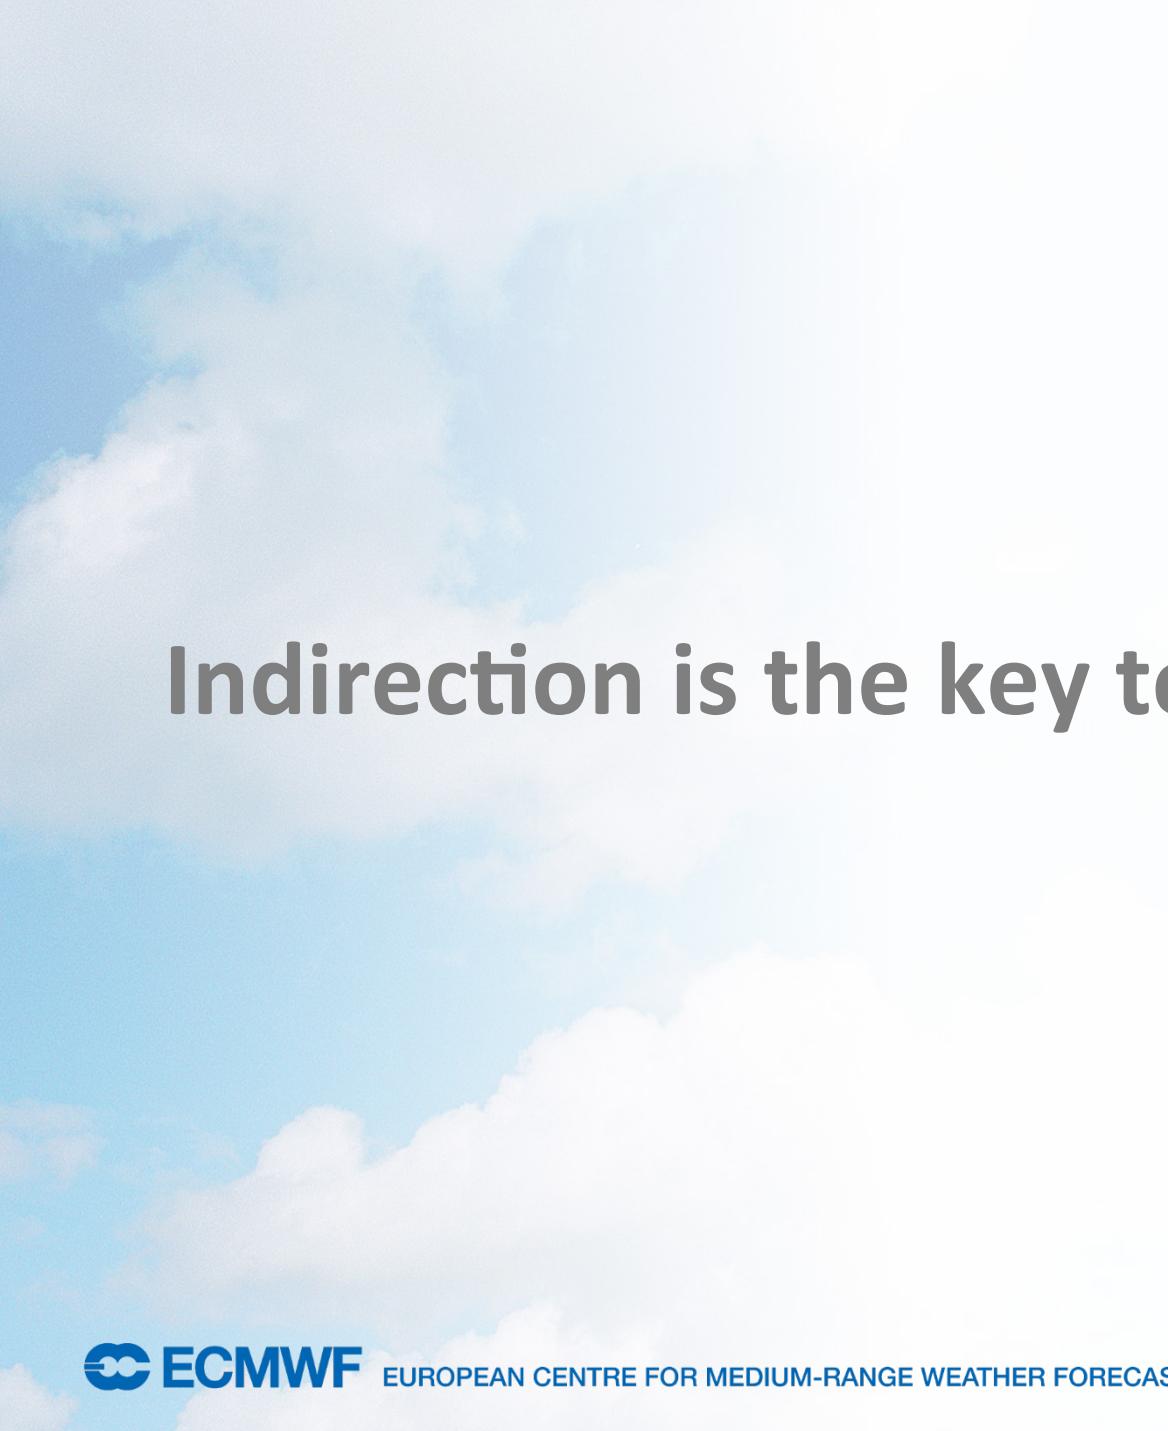


MARS 2011 Migration

- From monolithic to distributed
- From AIX to Linux
- From Big Endian to Little Endian
- Metadata indexing 40PB
- Job 2 years in the planning
- Degraded service for a couple of hours
- No one noticed...

A meteorological language

- retrieve,
date = 20110101/to/20110131,
parameter = temperature/geopotential,
type = forecast,
step = 12/to/240/by/12,
levtype = pressure levels,
levels = 1000/850/500/200,
grid = 2/2,
area = -10/20/10/0
- This request:
 - represents $31 \times 2 \times 20 \times 4 = 4960$ fields
 - defines an interpolation followed by cropping



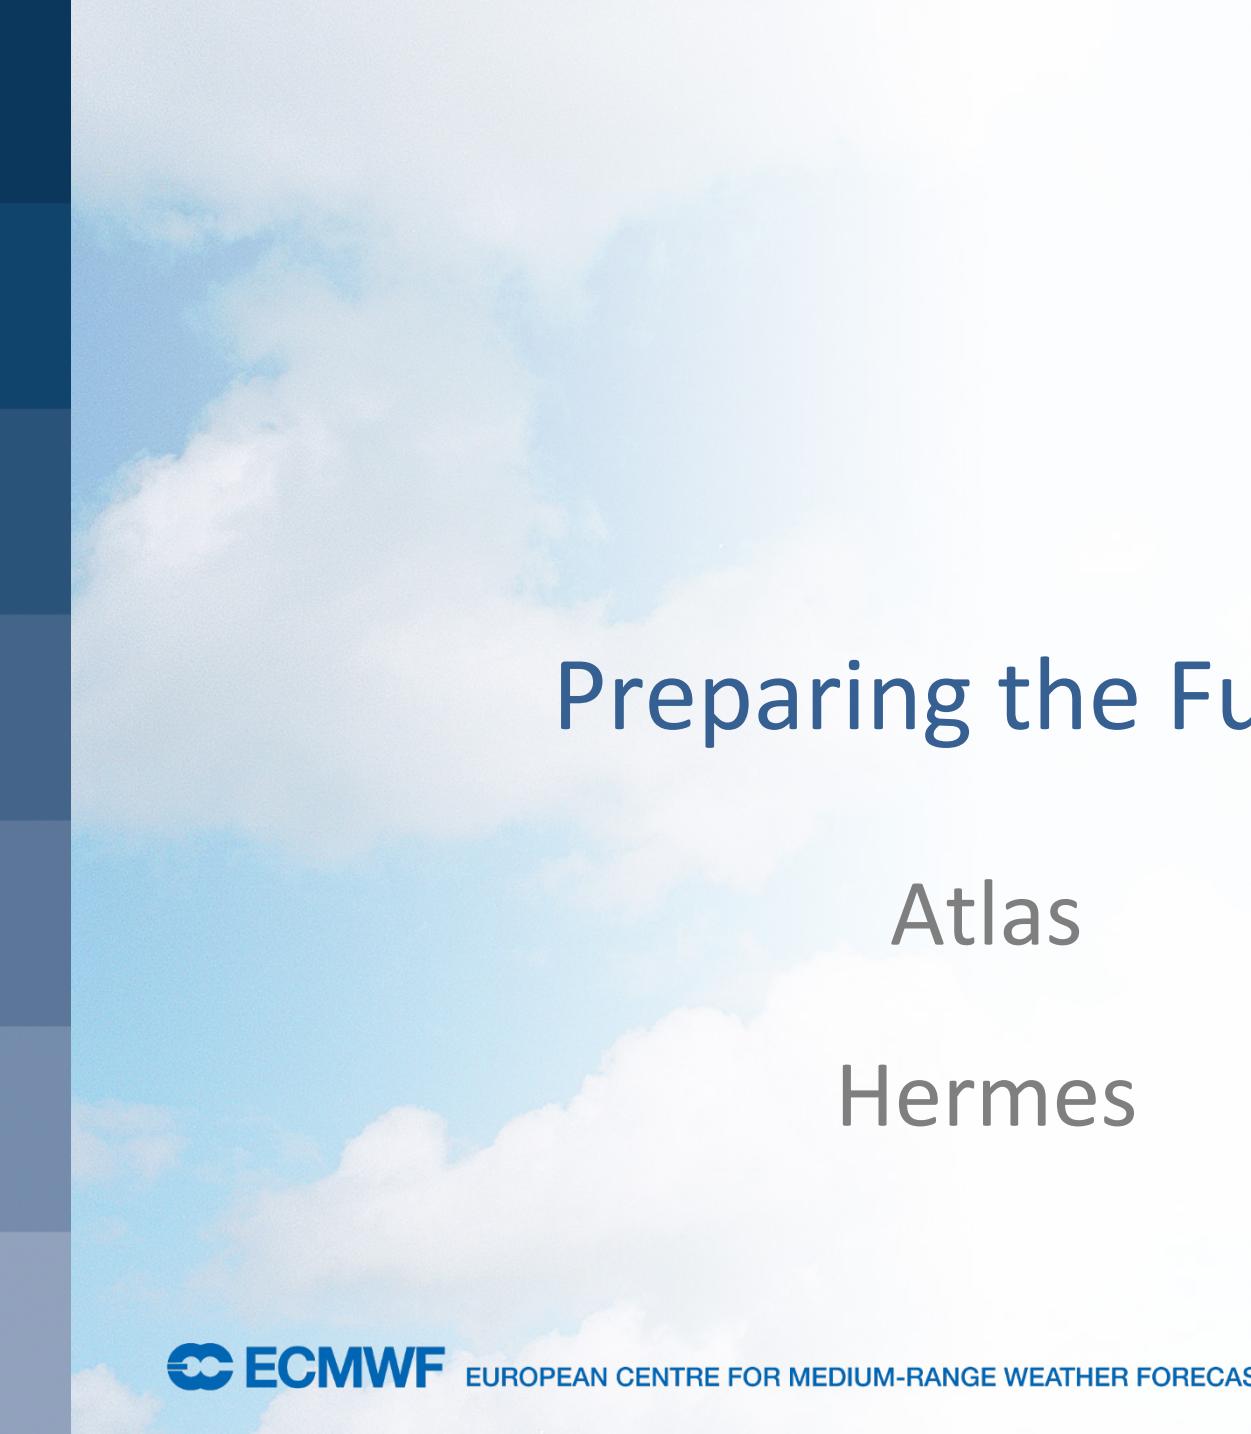
Indirection is the key to scalability



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

57



Preparing the Future

Atlas

Hermes



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

58

Aims

- Big data challenge (scalability)
- Adaptable solution (flexibility)
- Code evolution (maintainability)

Lets use C++ (98)

Atlas

W. Deconinck, T. Quintino

Challenges: New Algorithms & Platform Uncertainty



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

60

Current IFS Model

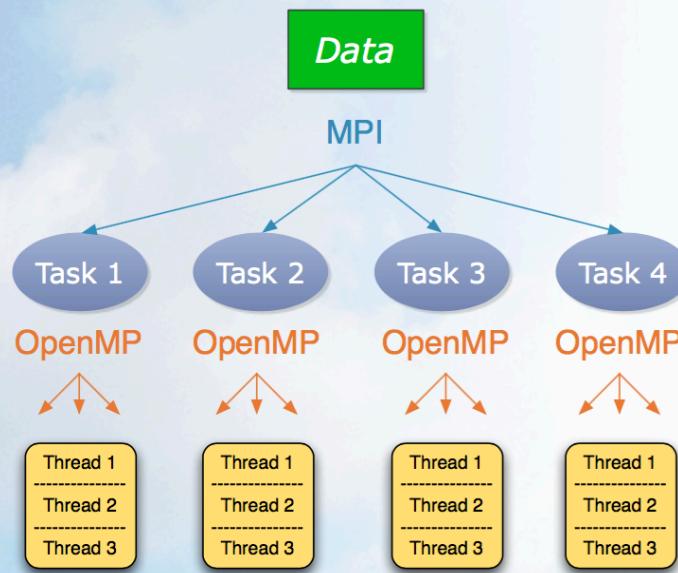
- 2M lines of Fortran code
- 9000 routines / modules
- Hybrid MPI + OpenMP
- Fixed data-structure
- Well organized but...
- ... lots of global data.
- Impractical to rewrite
- Fortunately most global operations are located in the “dynamical core”.

Atlas

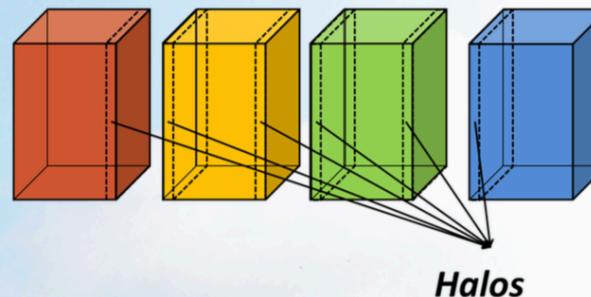
- Framework for parallel, dynamic data structures
- Supporting multiple types of grids
- Fully written in **C++**
- Providing Fortran 2003 interfaces
- Basis to develop **scalable** dynamical core

Atlas capabilities

Parallelisation



Local computations in every subdomain

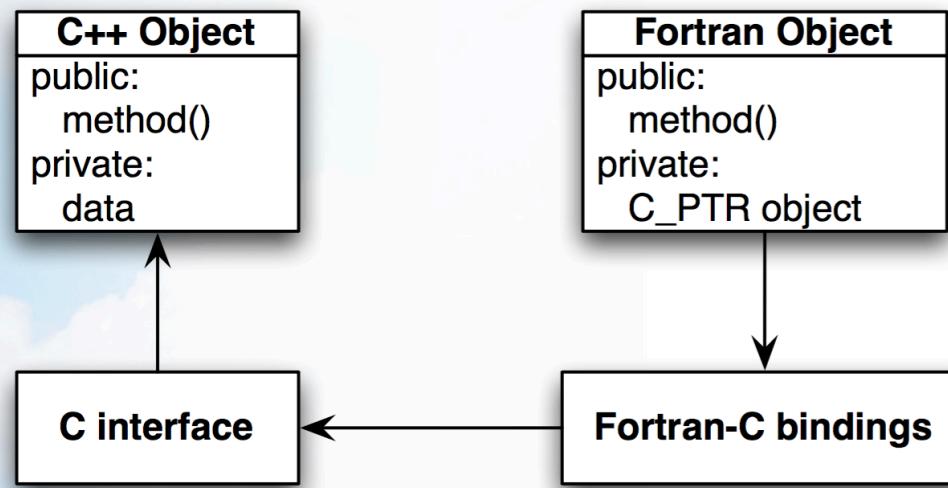


Same Equal-Area Domain decomposition
(T255 – 96 partitions)



So what about Fortran?

- Keep Fortran for numeric algorithms
- C++ for ...
 - All data & memory management
 - Message passing
- For every C++ class -> Fortran Derived Type
- Use Fortran ISO C Bindings (F2003)



C++ contains implementation

```
// C++ implementation
class Field {
public:
    Field(int size) { data_.resize(size); }
    double* data() { return &data_[0]; }
private:
    vector<double> data_;
}

// C interface – to be called from Fortran
Field* atlas_new_Field (int size) {
    return new Field( size );
}
void atlas_Field_access_data(Field* this,
                             double& data,
                             int& s){
    data = this->data();
    size = this->size();
}
```



WARNING: Fortran code ahead!



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

66

Fortran/C interface

```
! Fortran Derived type , contains no data , only methods F03
TYPE :: Field_type
PRIVATE
TYPE(C_PTR) :: object = C_NULL_PTR
CONTAINS
    PROCEDURE, PUBLIC :: access_data => Field__access_data
END TYPE

! Method calls C function, which uses C++ implementation
SUBROUTINE Field__access_data(this,data)
    CLASS(Field_type), intent(in) :: this
    TYPE(C_PTR) :: cptr_data
    REAL(JPRB), POINTER :: data
    INTEGER :: size
    ! This is the C function
    CALL atlas_Field__access_data(this%object ,cptr_data ,size)
    ! Converts the C pointer to a fortran pointer to array of doubles
    CALL C_F_POINTER(cptr_data ,data ,(/size/))
END SUBROUTINE
```

Fortran Program

```
PROGRAM main

  TYPE(Field_type) :: field
  REAL(JPRB), POINTER :: field_data(:)
  INTEGER :: j

  ! We can ask the field for access to its data
  CALL field%access_data(field_data)

  ! Now we can use the data as any other Fortran array
  DO j=1,size(field_data)
    field_data(j) = 0.
  ENDDO

END PROGRAM
```

- Furthermore, we provide another higher level API to hide Fortran 2003.

Hermes

T. Quintino, B. Raoult, F. Rathgeber

Challenges: Big Data & Platform Uncertainty



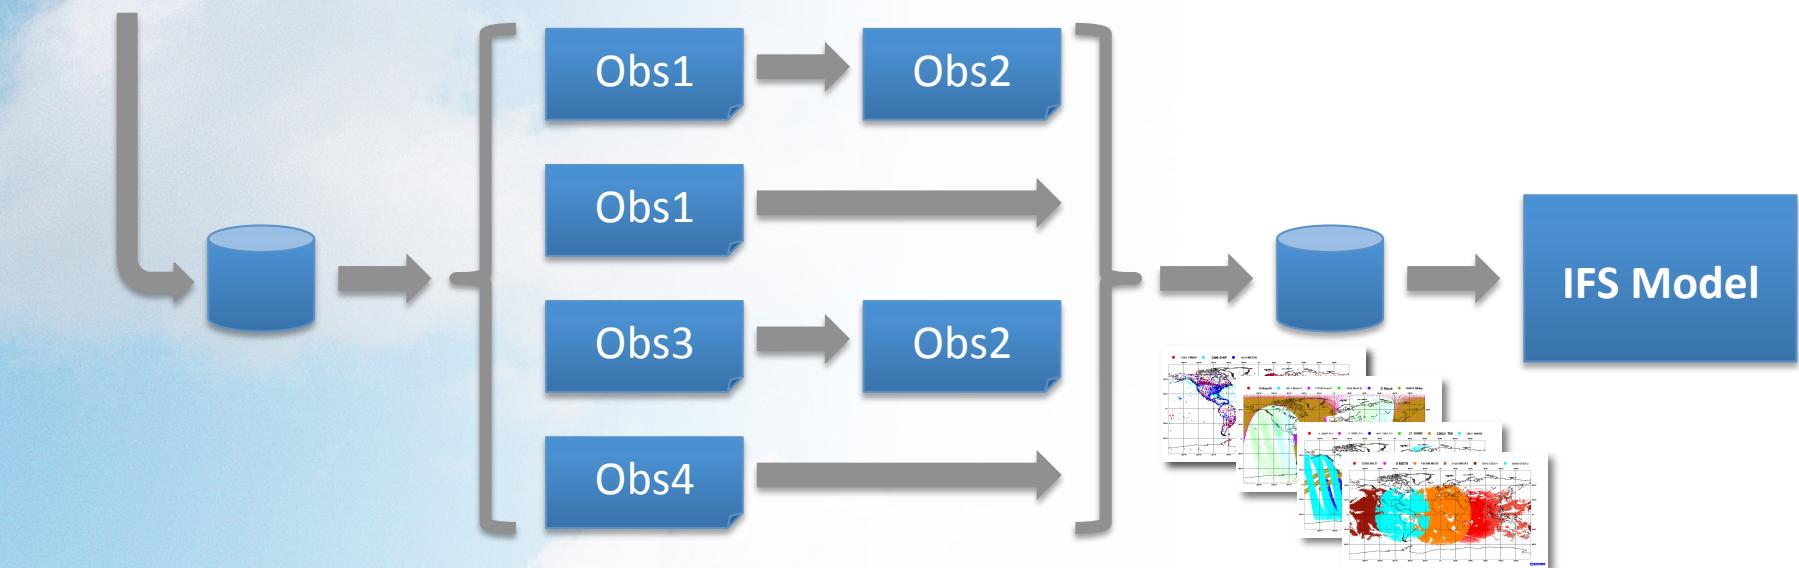
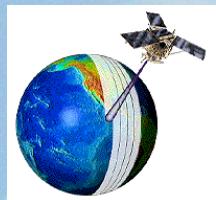
EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

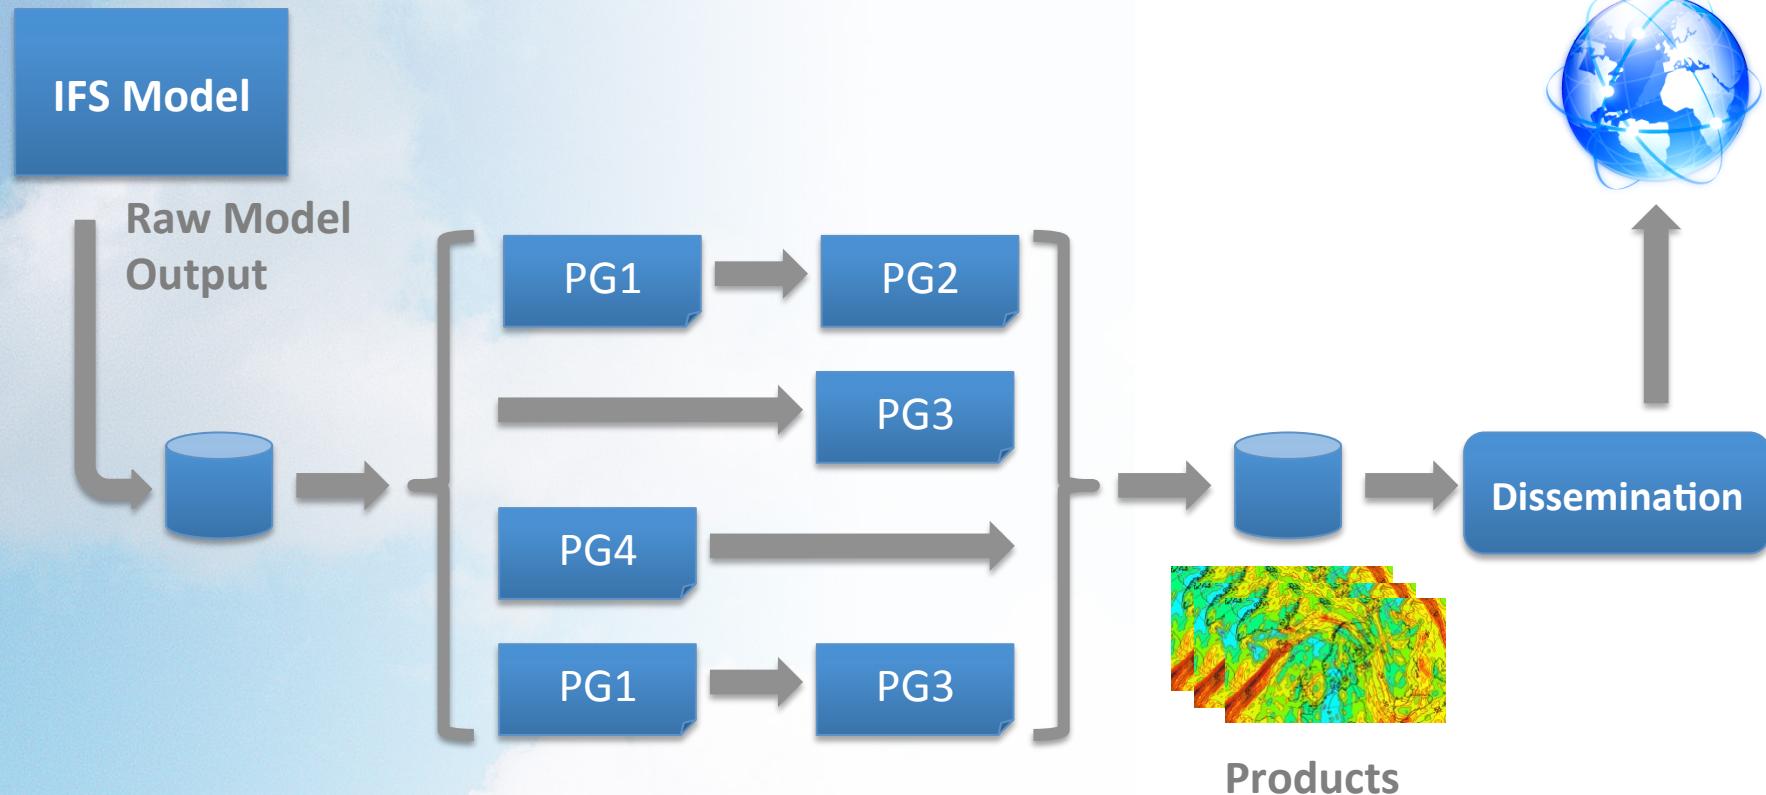
69

A look at the data-chain

Observations



A look at the data-chain



Observations & Fields: Some Similarities

- Both processed in pipelines
- Both are parallel...
 - Concurrent (Threads / OpenMP)
 - Distributed (MPI)

How does it **differ** from an atmospheric Model?

- Parallel, but little or no synchronisation
- No time evolution
- “Embarassingly parallel” (nearly)

Possible: *Automatic Parallelisation*

Requirements

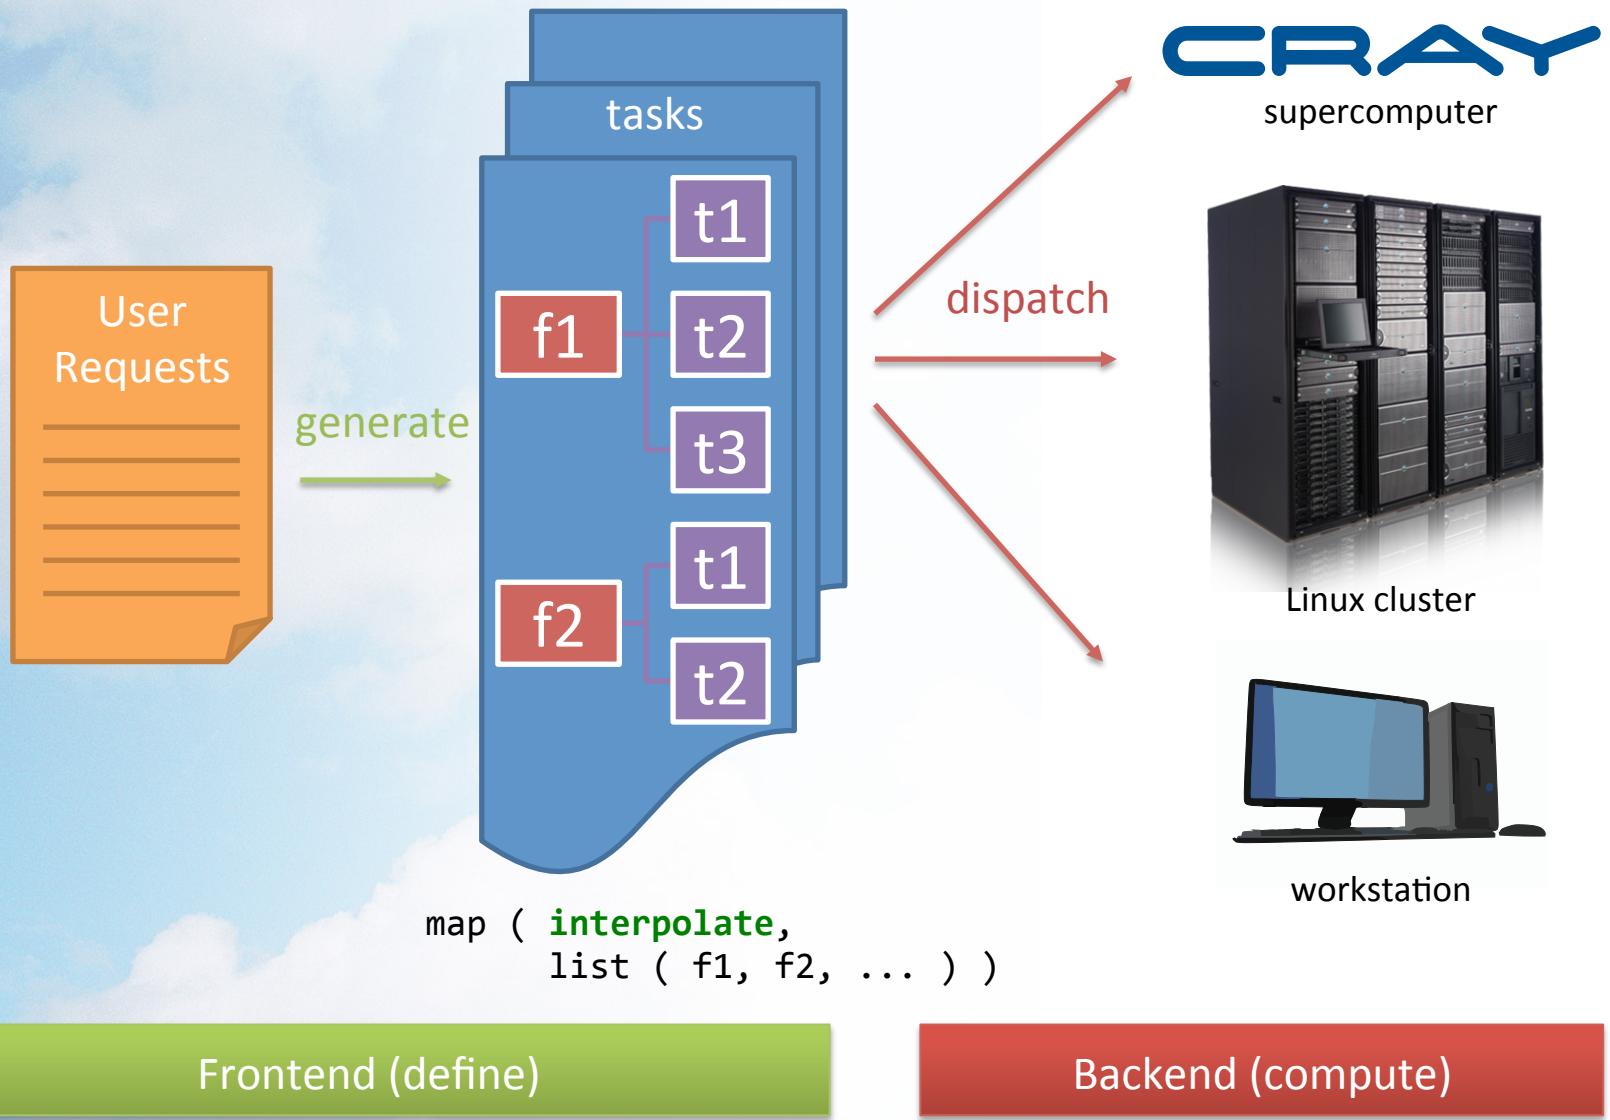
- Way to define computations
- Efficient parallel computation
- Way to accept data and deliver results
- Flexibility for future architectures
- Extendable and understandable by scientists

Solution

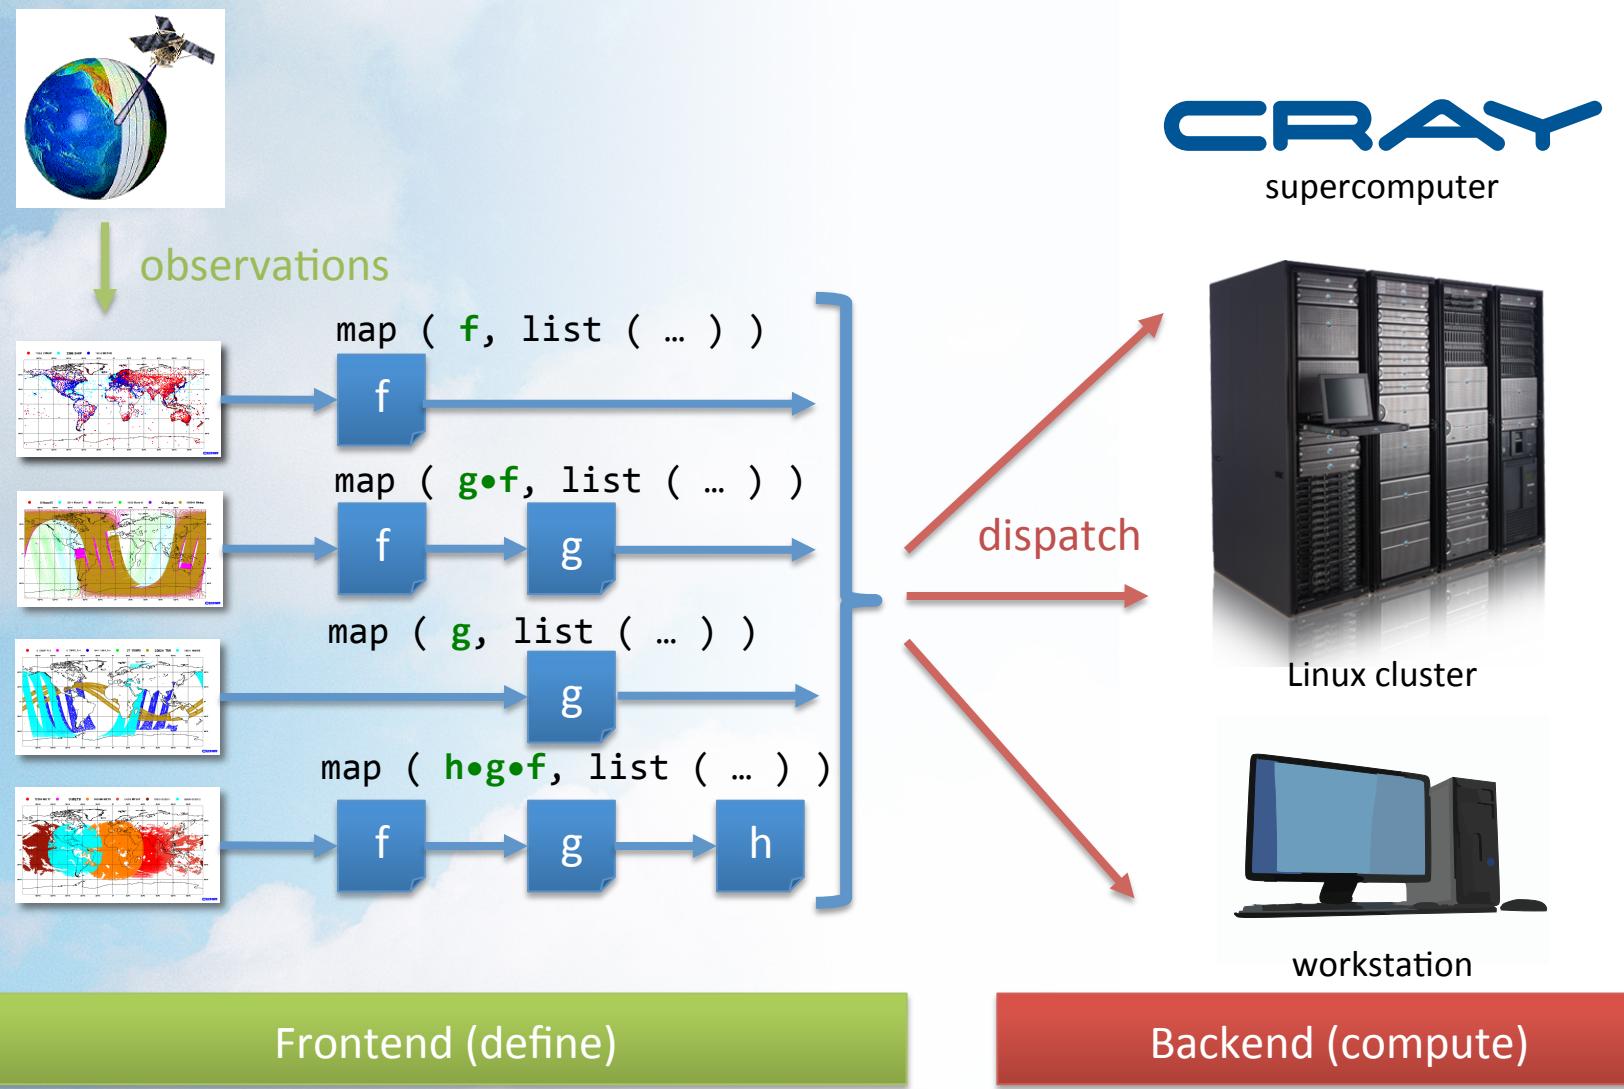
Indirection is key to Scalability

Decouple **what** to compute
from **how** and **where** to compute

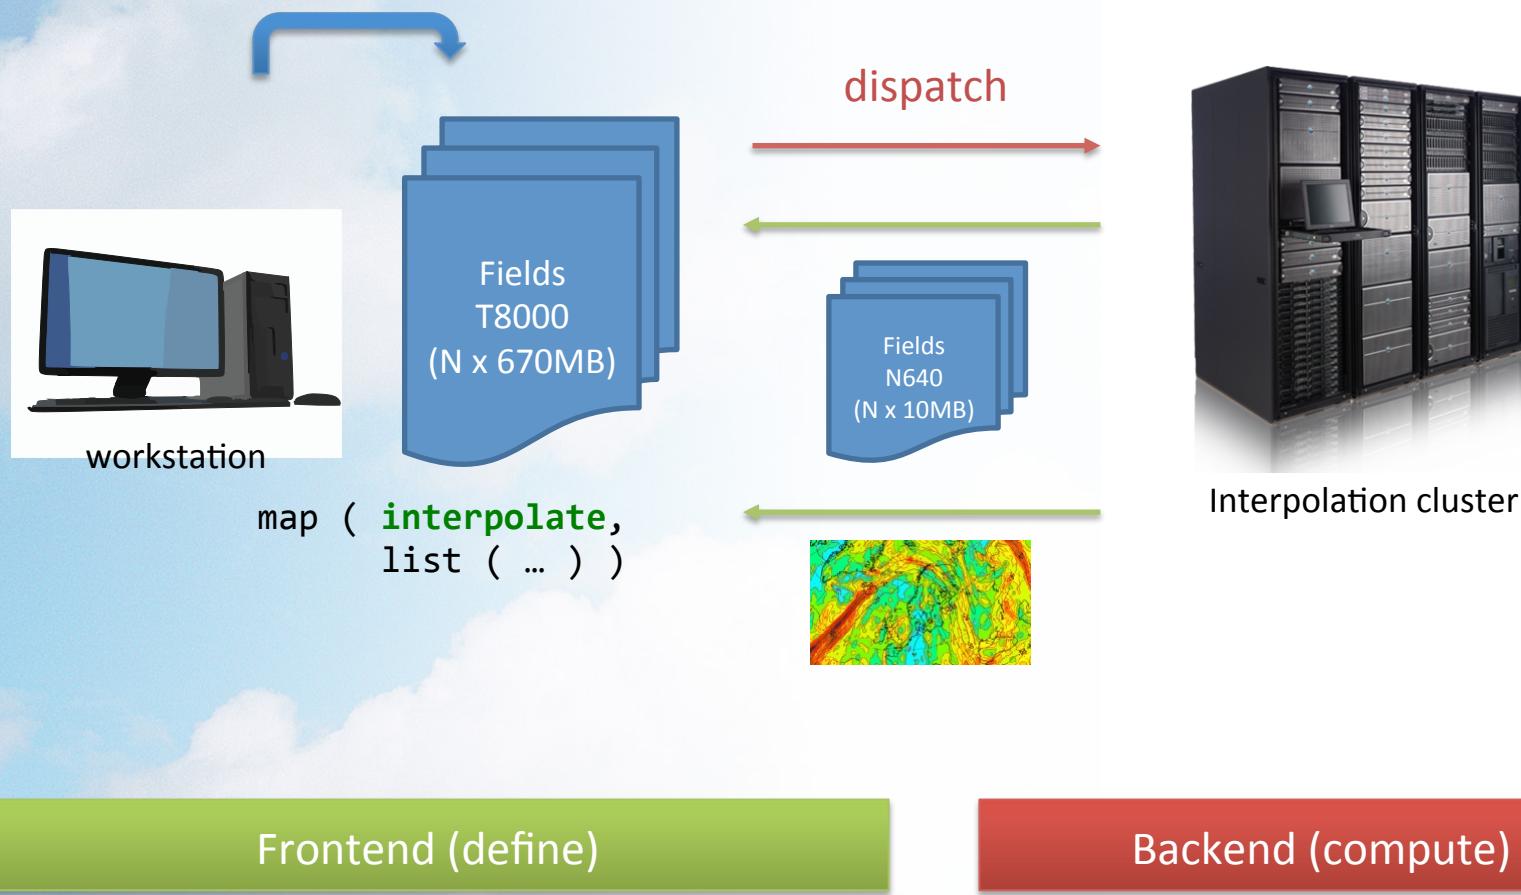
Use Case: Product Generation



Use Case: Observation Filters



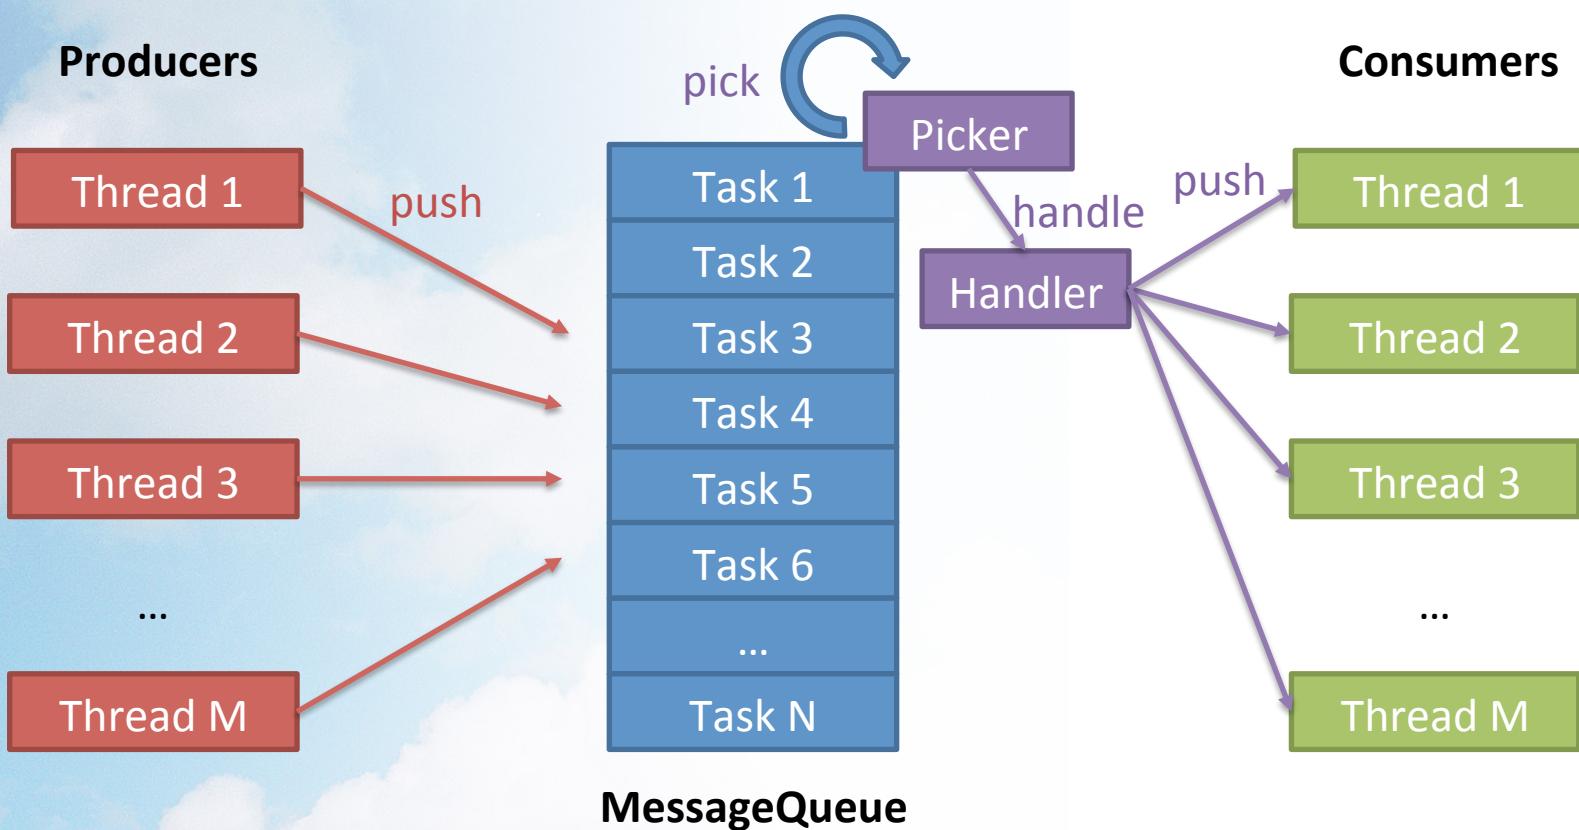
Use Case: Workstation Interpolation & Visualisation



Hermes

- DSL to define computations
- Frontend for sending requests and data
- Broker to dispatch work to PU's
- Multiple computational back-ends:
 - CPU
 - CUDA (Linear Algebra)
 - Intel MKL (Linear Algebra)
 - *Intel TBB ??*
 - *HPX ??*

Dispatcher



Expression Engine

- Prototype DSEL
- Embedded in C++
- *Plug-ins in Fortran (not yet)*
- Computations as expression trees
- Hierarchy of expressions:
 - terminals (value semantics)
 - functions (pure operations)

Expression Engine

- Inspired by functional language Haskell
- Expressions immutable
- **Functions pure**
- No variables, only constants
- **Lazy** evaluation
- Optimiser: find most efficient path to evaluate expression
 - Pattern matching (e.g. linear combination)
 - Collapse expression tree into terminal
 - Detect common sub-expressions
 - *Automatic parallelisation*

Expression

```
typedef std::shared_ptr<Expression> ExpPtr;
typedef std::vector< ExpPtr > args_t;

class Expression {
public:
    ...
    Expression(const args_t& args);
    size_t arity() const { return args_.size(); }
    ...
    ExpPtr self();
    ExpPtr eval() const { Scope s; return ExpPtr(optimise()->evaluate(s)) };
    ExpPtr eval(Scope& s) const { return ExpPtr(optimise()->evaluate(s)) };
    ...
    // used to bind lambda parameters
    virtual ExpPtr resolve(Scope&) const { return self(); }
    virtual ExpPtr optimise(size_t depth) const { return self(); }
    template<typename T> shared_ptr<T> as() { ... }
private:
    virtual ExpPtr evaluate(Scope&) const = 0;
    virtual ExpPtr cloneWith(args_t&) const = 0;
    virtual std::string signature() const = 0;
};
```

Values

```
class Value : public Expression {  
protected:  
    Value( const args_t& args ) : Expression(args) {}  
private:  
    virtual ExpPtr evaluate(Scope&) const { return self(); }  
};  
  
class Real : public Value {  
public:  
    typedef real_t value_t;  
    ...  
    Real(ExpPtr e);  
    Real(const real_t& v);  
    ...  
    value_t value() const { return v_; }  
    static value_t extract (const ExpPtr& e) { return e->as<Real>()->value(); }  
    ...  
private:  
    virtual std::string signature() const { return "r"; }  
};
```

Bool, Integer, Real, String, List, Vector, Matrix

Functions

```
class Function: public Expression {  
public:  
  
    typedef ExpPtr (*func_t) (Scope&, const args_t&);  
    typedef std::map< key_t, func_t > dispatcher_t;  
  
    virtual ExpPtr evaluate(Scope&) const;  
    virtual ExpPtr optimise(size_t depth) const;  
  
    virtual std::string signature() const { return signatureArguments(args()); }  
  
protected:  
    Function( const args_t& args ) : Expression(args) {}  
    // returns something like "func(r,v,r,v)"  
    std::string signatureArguments(const args_t& args) const {...}  
  
private:  
    virtual ExpPtr evaluate(Scope&) const { return self(); }  
};
```

Functions (2)

```
ExpPtr Function::evaluate(Scope& ctx) const {

    // create temporary args
    args_t tmp = args();
    for(size_t i = 0; i < tmp.size(); ++i) {
        tmp[i] = args(i, ctx, true)->self(); // evaluate the arguments
    }

    std::string sig = signatureArguments(tmp); // new signature

    dispatcher_t& d = dispatcher(); // find dispatching function
    dispatcher_t::iterator itr = d.find(sig);
    if(itr == d.end()) { ... fail ... }

    return (*itr).second)( ctx, tmp ); // dispatch
}

ExpPtr Function::optimise(size_t depth) const {
    // revert to generic optimizer if Expression doesn't optimize itself
    return Optimiser::apply(self(), depth);
}
```

Functions (3)

```
class Map : public Function {  
public:  
    Map(ExpPtr func, ExpPtr lst );  
  
private:  
    ...  
    ExpPtr Map::evaluate(Scope& ctx) const {  
  
        ExpPtr f = args(0, ctx, false); // get the function, don't evaluate  
        ExpPtr ls = args(1, ctx, true); // get the list and evaluate if necessary  
  
        const List::value_t& list = List::extract(ctx, ls);  
  
        List::value_t res;  
        res.reserve(list.size());  
        for( size_t i = 0; i < list.size(); ++i ) {  
            ExpPtr e = list[i]->resolve(ctx)->eval(ctx); // evaluate each element  
            ExpPtr v = f->eval(e);  
            res.push_back( v );  
        }  
        return ExpPtr(new List(res));  
    }  
    ...  
};
```

Functions (4)

Map
Reduce
Merge
Take
ZipWith
IfElse
Call
Lambda

Filter
Bind
UnaryOperators
• neg,sqrt,abs
BinaryOperators
• mul,add,sub,div
• mod,min,max
Predicates

Examples

```
auto a = real(2.);  
auto b = real(4.);  
auto x = vector({1.,2.,3.});  
auto y = vector(3, 5.);  
  
auto c = add(a,b);  
  
std::cout << *c << std::endl; // Add(Real(2), Real(4))  
  
auto e = add(mul(c, x) , mul(b, y));  
  
c = c->eval();  
  
std::cout << *c << std::endl; // Real(6)  
  
// Add(Mul(Add(Real(2), Real(4)), Vector(1, 2, 3)), Mul(Real(4), Vector(5, 5, 5)))  
std::cout << *e << std::endl;  
  
auto opt = e->optimise(true);  
  
std::cout << opt->signature() << std::endl; // optimised to "Linear(r,v,r,v)"  
  
auto f1 = reduce(mul(), list(a, x, x))->eval(); // Vector(2, 8, 18)
```

Examples

class **Xpr** provides value semantics and convinience operators

```
// lambdas

Xpr twice = call(lambda("x", Xpr(2.0) * Xpr("x"))); // define 2*x

Xpr l = list(a,b,c);

Xpr r = map(twice, l);

std::cout << r() << std::endl; // List(Real(4), Real(8), Real(12))

// currying

Xpr F = lambda("i", call( lambda("j", Xpr("i") * Xpr("j")))); // define i*j

Xpr H = call(F, Xpr( 2. )); // curry to i*2

Xpr G = H( 3. );           // curry to 3*2

std::cout << G() << std::endl; // Real(6)
```

Examples

```
// Define the Y-combinator

Xpr Y = lambda("f", call(Xpr("f"), Xpr("f")));

// lookup table for first 10 elements of series
Xpr table = list( Xpr(0.), Xpr( 1.), Xpr( 1.), Xpr( 2.), Xpr( 3.), Xpr(5.),
                  Xpr(8.), Xpr(13.), Xpr(21.), Xpr(34.), Xpr(55.) );

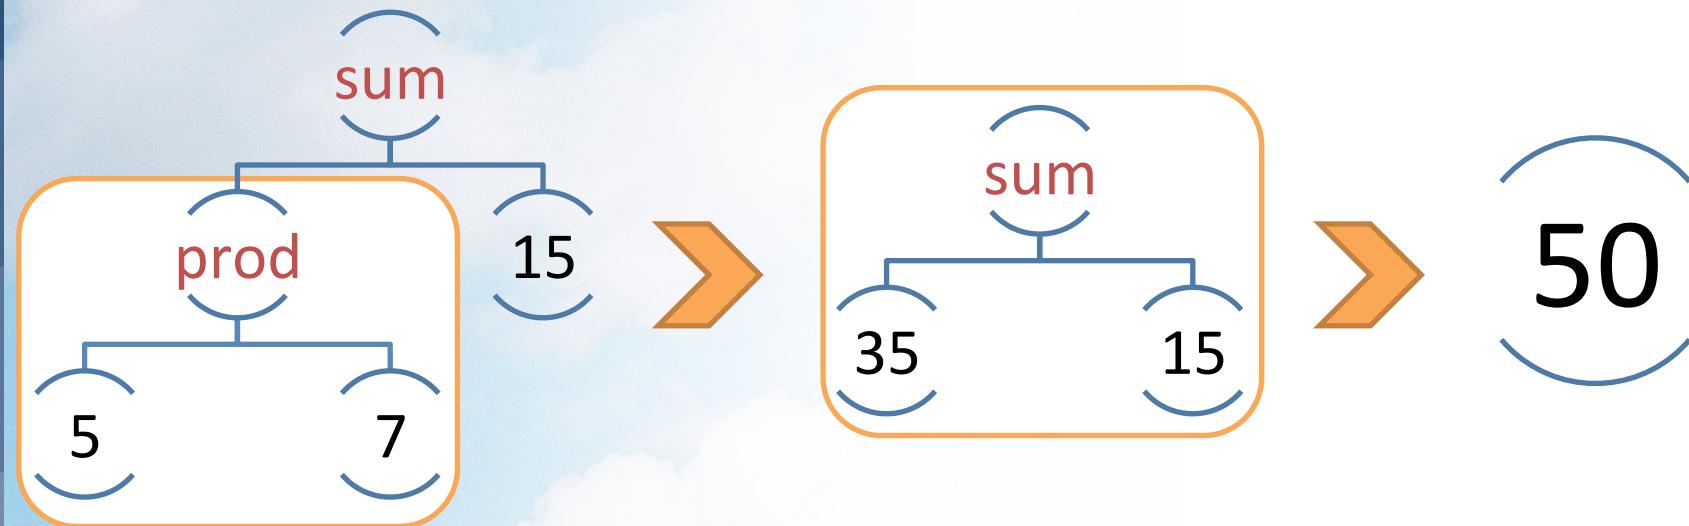
// non-recursive fibonacci function
Xpr fib = lambda("f",
                  call(lambda("n",
                            ifelse(
                                Xpr("n") > Xpr(10.),
                                Xpr(call(Xpr("f"),Xpr("f"), Xpr("n")-Xpr(1.0))) +
                                Xpr(call(Xpr("f"),Xpr("f"), Xpr("n")-Xpr(2.0))),
                                take( Xpr("n"), table )
                            )
                        )
                );
);

Xpr fibonnaci = call(Y, fib); // applying the Y-combinator

std::cout << fibonnaci(20.) << std::endl; // Real(6765)
```

Arithmetic Expression

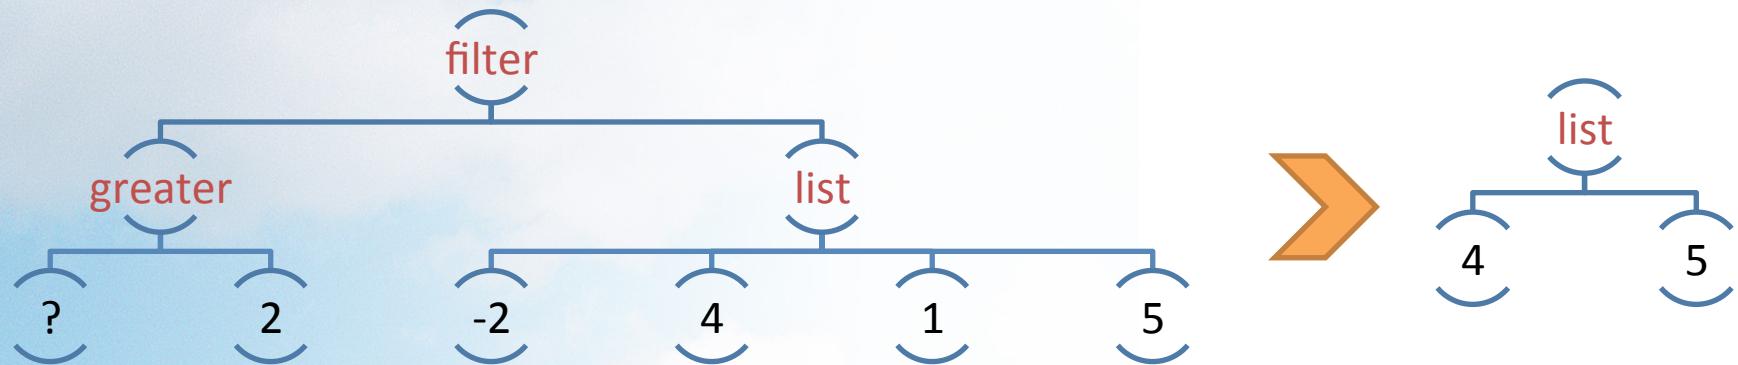
$5 * 7 + 15 == \text{add}(\text{prod}(5, 7), 15)$



constant folding and propagation

Filter Expression

```
filter ( greater ( ?, 2 ), list ( -2, 4, 1, 5 ) )
```



expression elimination

Future work: Map fusion

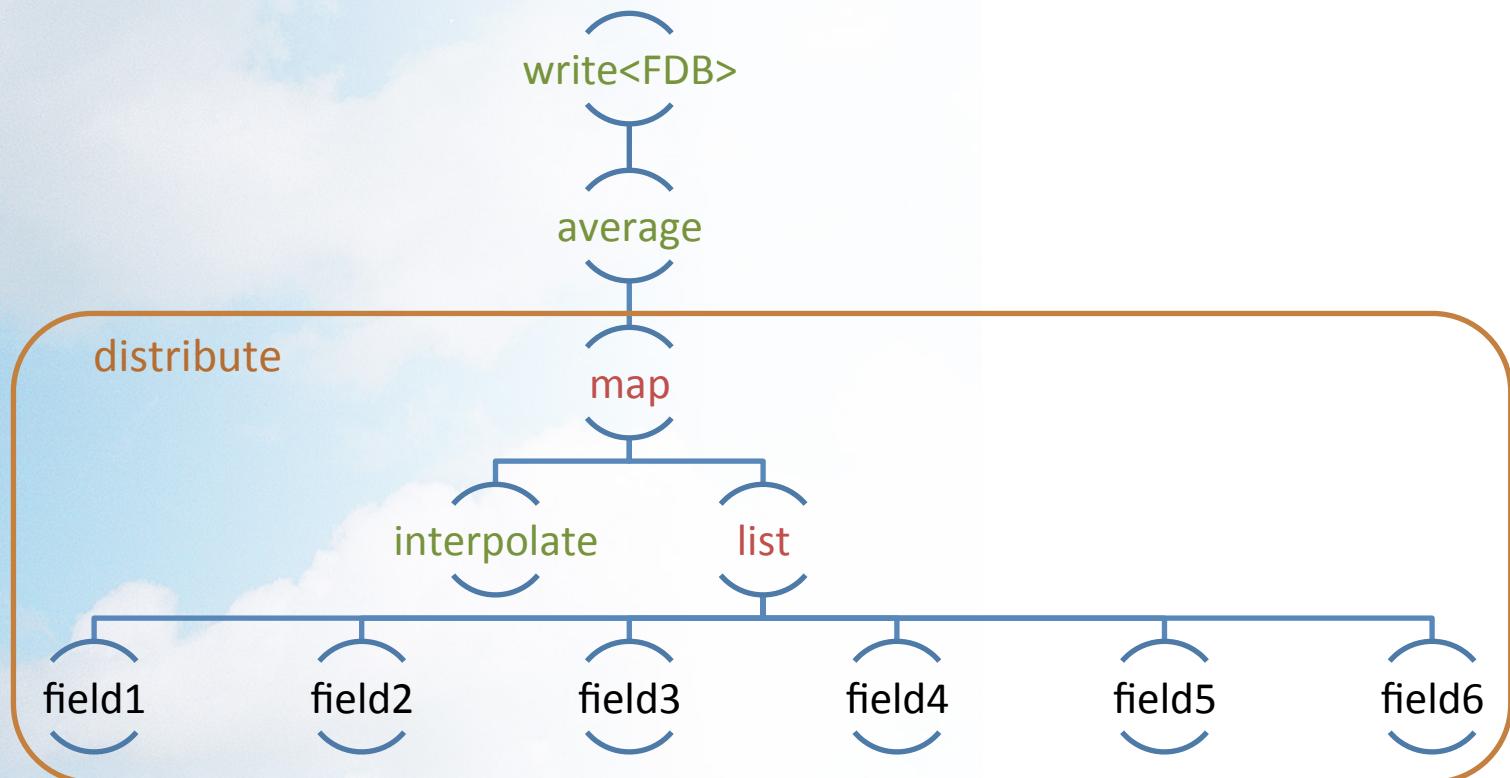
map (h, map (g, map (f, list (...))))



Future work

Distributed interpolation of fields

```
write<FDB> ( average ( map ( interpolate, list ( field1,  
field2,  
...,  
field6 ) ) ) )
```

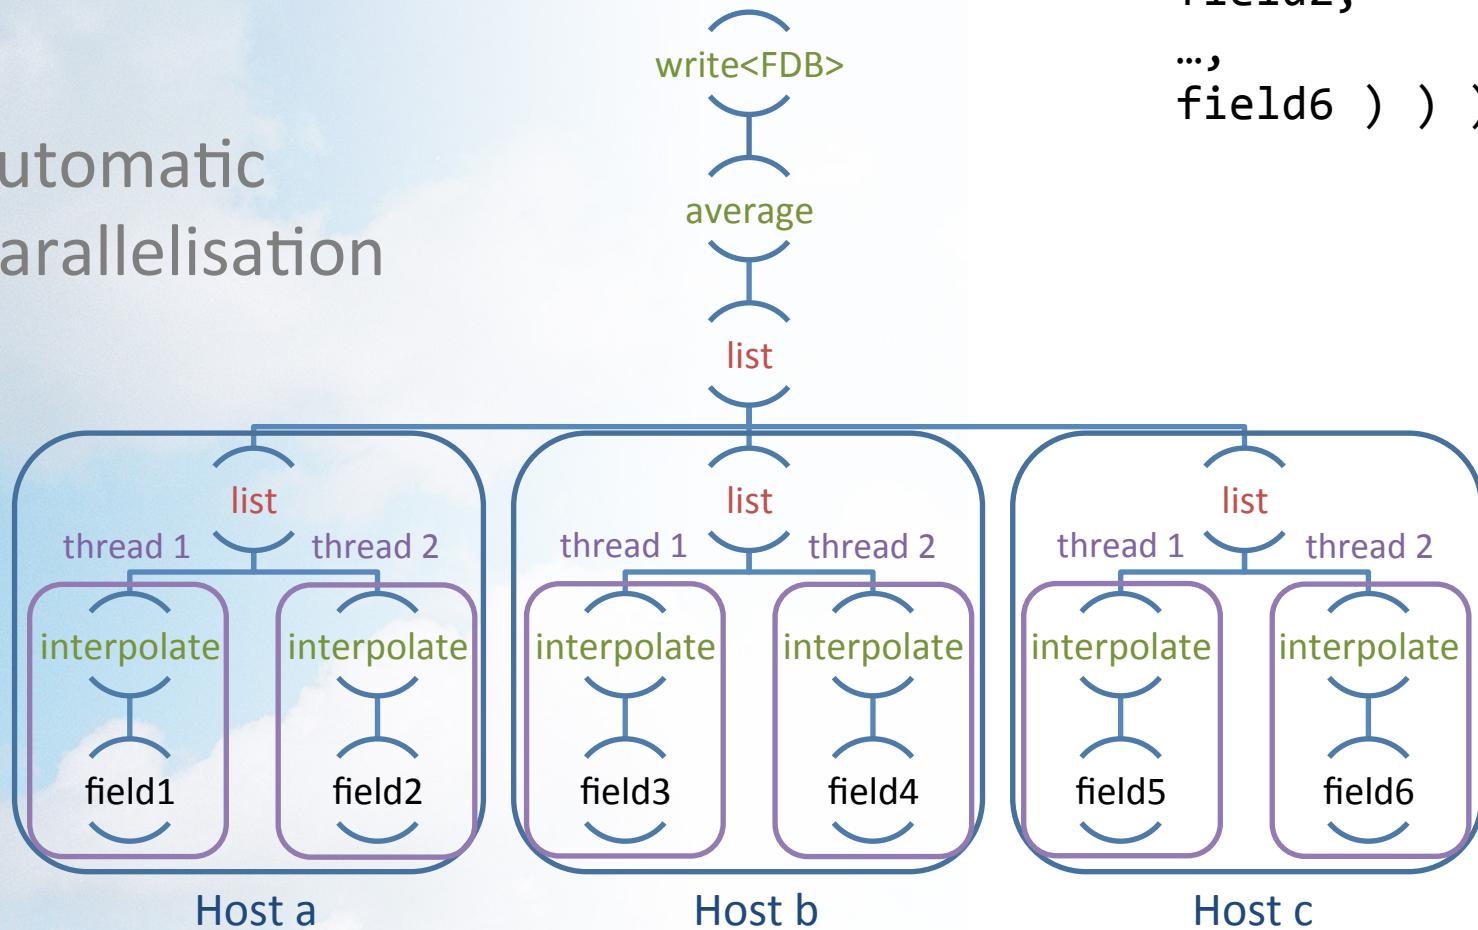


Future work

Distributed interpolation of fields

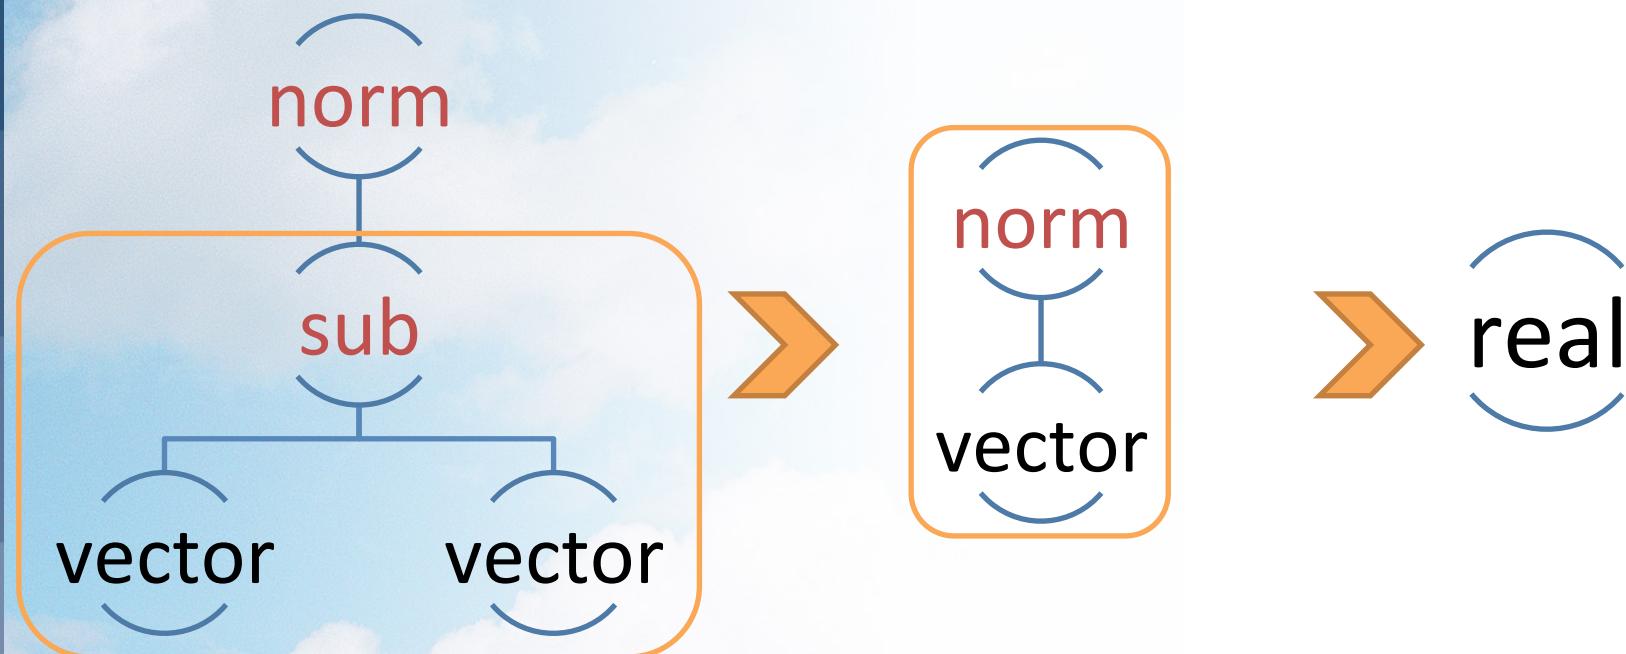
```
write<FDB> ( average ( map ( interpolate, list ( field1,  
field2,  
...,  
field6 ) ) ) )
```

automatic
parallelisation



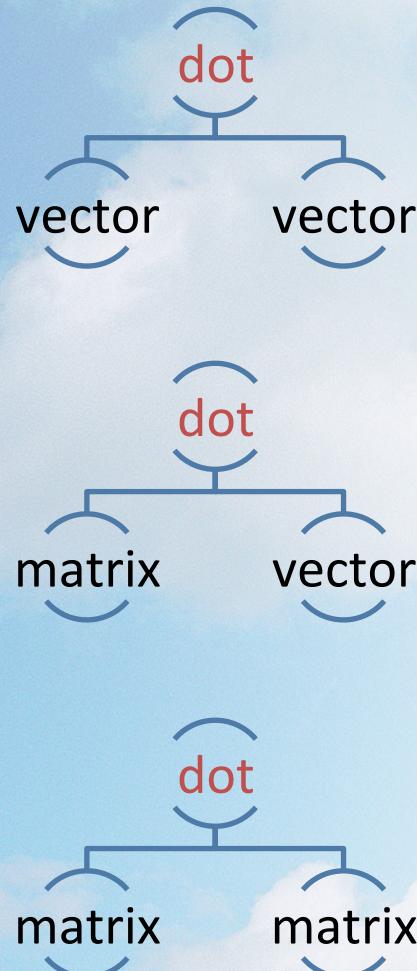
Linear Algebra Expressions

norm (sub (vector, vector))



norm of difference of two vectors

Dispatch to Compute Engines



Expression
signatures

Dot(v,v)
Dot(m,v)
Dot(m,m)

dispatch

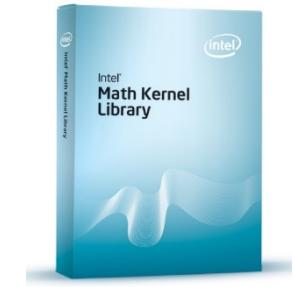
GPU engine

cublasDdot
cublasDgemv
cublasDgemm



MKL engine

cblas_ddot
cblas_dgemv
cblas_dgemm



Generic engine

ddot
dgemv
dgemm

**BLAS/
LAPACK**

Expressions: client code

```
const size_t ncols = 1024;
const size_t nrows = 1024;

auto init = [] (size_t, size_t j=0) {
    return (real_t)std::rand() / RAND_MAX;
};

XprClient c;

// L2 norm of difference of 2 vectors
c.send( norm( sub( vector(ncols, init), vector(ncols, init) ) ) );

// Dot product of vectors
c.send( dot( vector(ncols, init), vector(ncols, init) ) );

// Dot product of matrix and vector
c.send( dot( matrix(nrows, ncols, init), vector(ncols, init) ) );

// Dot product of matrix and matrix
c.send( dot( matrix(nrows, ncols, init), matrix(ncols, nrows, init) ) );
```

Generic Engine implementation

```
ExpPtr compute_dot_generic(Scope&, const args_t &p)
{
    Vector::value_t v1 = Vector::extract(p[0]);
    Vector::value_t v2 = Vector::extract(p[1]);
    ASSERT( v1.size() == v2.size() );
    Vector::elemt_t r;
    for( size_t i = 0; i < v1.size(); ++i )
        r += v1[i] * v2[i];
    return real(r);
}

ExpPtr compute_gemv_generic(Scope&, const args_t &p);
ExpPtr compute_gemm_generic(Scope&, const args_t &p);

// Register generic functions in dispatch table
struct DotRegister {
    DotRegister() {
        Function::dispatcher()["xpr::Dot(v,v)"] = &compute_dot_generic;
        Function::dispatcher()["xpr::Dot(m,v)"] = &compute_gemv_generic;
        Function::dispatcher()["xpr::Dot(m,m)"] = &compute_gemm_generic;
    }
};
Static DotRegister dotRegister;
```

MKL Engine implementation

```
ExpPtr compute_dot_MKL(Scope&, const args_t &p)
{
    const Vector::value_t& v1 = Vector::extract(p[0]);
    const Vector::value_t& v2 = Vector::extract(p[1]);
    ASSERT( v1.size() == v2.size() );
    Vector::elemt_t r;
    return real( cblas_ddot( v1.size(), v1.data(), 1, v2.data(), 1 ) );
}

ExpPtr compute_gemv_MKL(Scope&, const args_t &p)
ExpPtr compute_gemm_MKL(Scope&, const args_t &p)

// Register MKL functions in dispatch table
void registerMKL() {
    Function::dispatcher()["xpr::Dot(v,v)"] = &compute_dot_MKL;
    Function::dispatcher()["xpr::Dot(m,v)"] = &compute_gemv_MKL;
    Function::dispatcher()["xpr::Dot(m,m)"] = &compute_gemm_MKL;
}
```

GPU Engine implementation

```
ExpPtr compute_dot_CUDA(Scope&, const args_t &p)
{
    const Vector::value_t& x = Vector::extract(p[0]);
    const Vector::value_t& y = Vector::extract(p[1]);
    ASSERT( x.size() == y.size() );
    const size_t size = x.size()*sizeof(real_t);
    Vector::elem_t r;

    real_t* d_x; // device memory vector x
    real_t* d_y; // device memory vector y
    cublasHandle_t handle;

    CALL_CUDA( cudaMalloc((void**) &d_x, size) );
    CALL_CUDA( cudaMalloc((void**) &d_y, size) );
    CALL_CUDA( cudaMemcpy(d_x, x.data(), size, cudaMemcpyHostToDevice) );
    CALL_CUDA( cudaMemcpy(d_y, y.data(), size, cudaMemcpyHostToDevice) );

    CALL_CUBLAS( cublasCreate(&handle) );
    CALL_CUBLAS( cublasDdot(handle, x.size(), d_x, 1, d_y, 1, &r) );
    CALL_CUBLAS( cublasDestroy(handle) );

    CALL_CUDA( cudaFree(d_x) );
    CALL_CUDA( cudaFree(d_y) );

    return real( r );
}
```

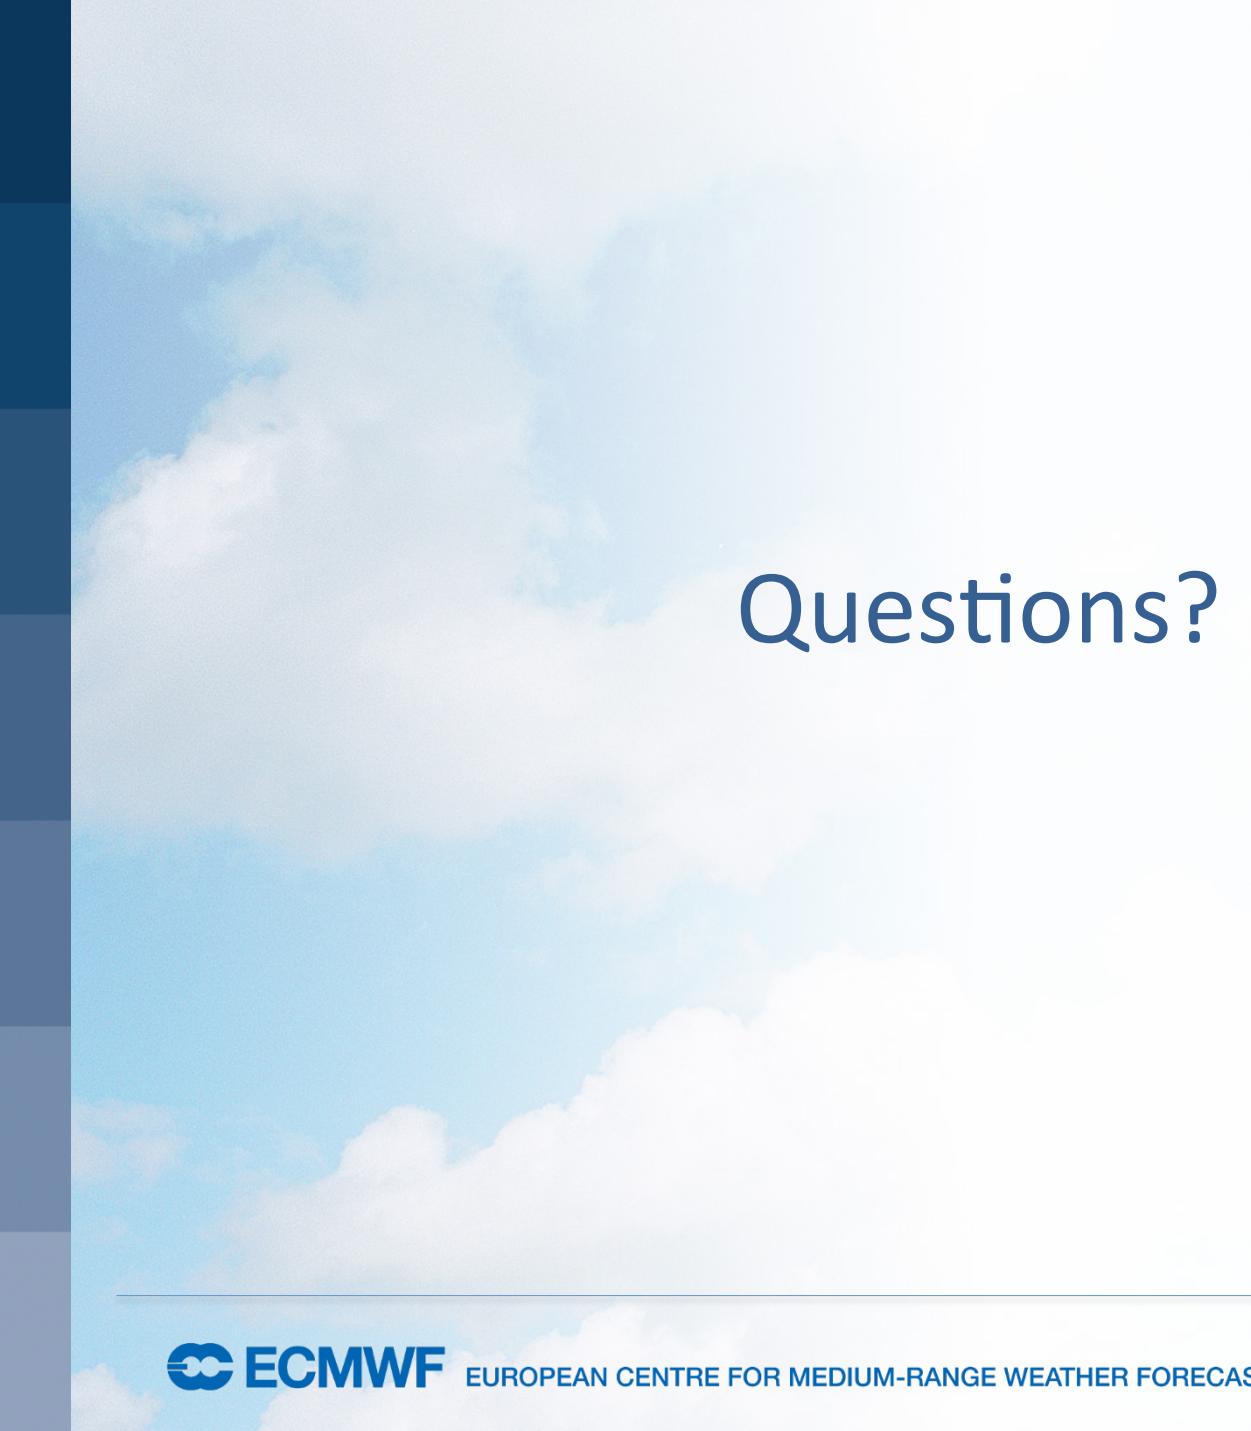
Live Demo



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

10
2



Questions?



EUROPEAN CENTRE FOR MEDIUM-RANGE WEATHER FORECASTS

13/05/15

10
3